



subtractive mixture models

representation, learning & inference

antonio vergari (he/him)

 @tetraduzione

18th Feb 2025 - Flatiron Institute NYC

april

april-tools.github.io

april

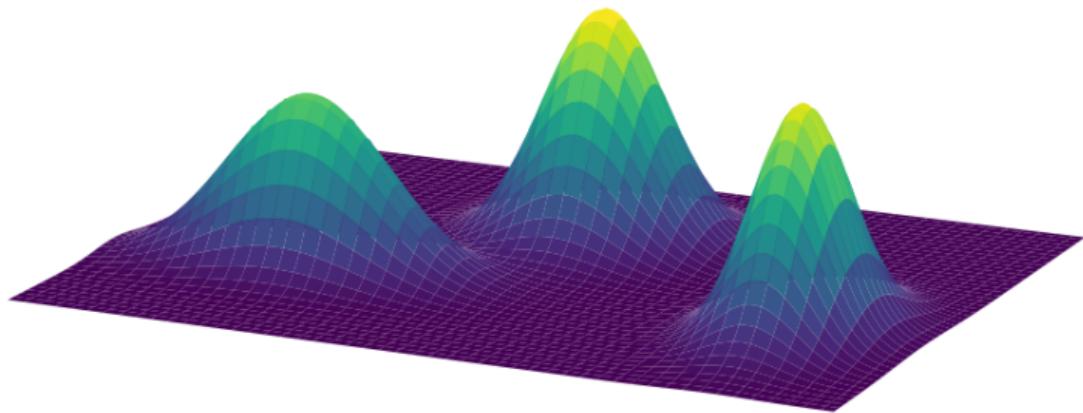
*about
probabilities
integrals &
logic*

april

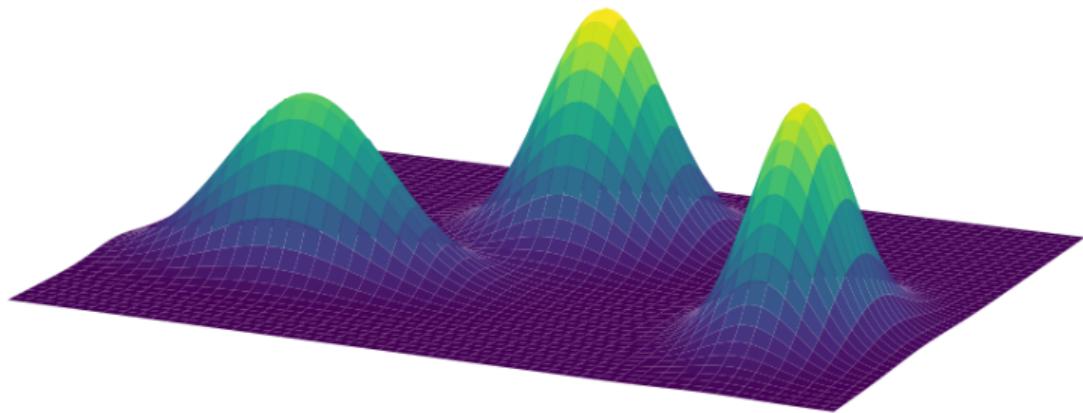
*autonomous &
provably
reliable
intelligent
learners*

april

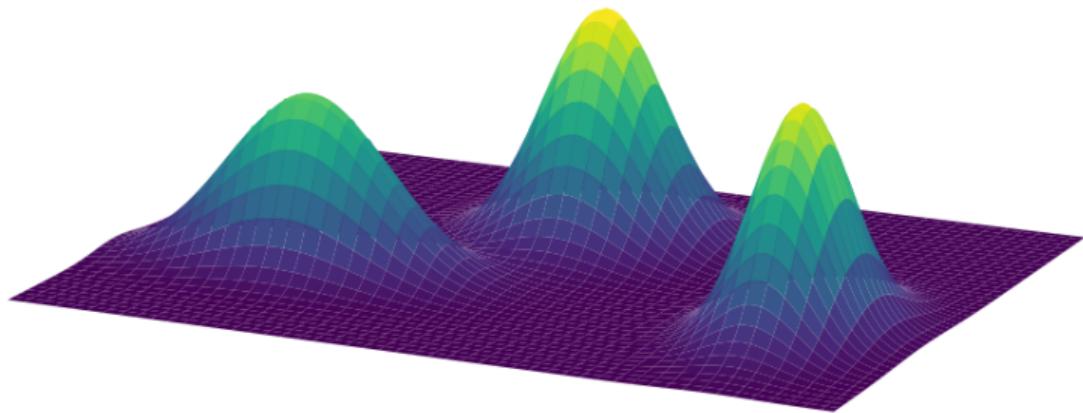
*april is
probably a
recursive
identifier of a
lab*



who knows mixture models?



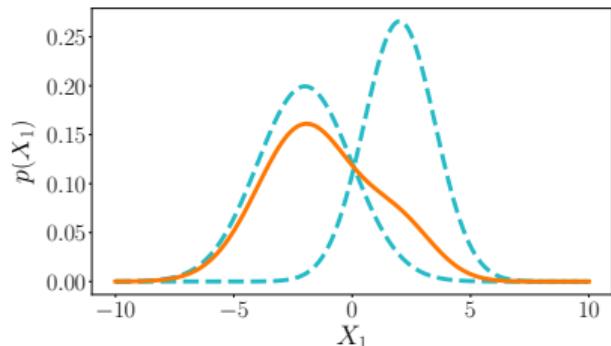
who loves mixture models?



$$c(\mathbf{X}) = \sum_{i=1}^K w_i c_i(\mathbf{X}), \quad \text{with} \quad w_i \geq 0, \quad \sum_{i=1}^K w_i = 1$$

GMMs

as computational graphs

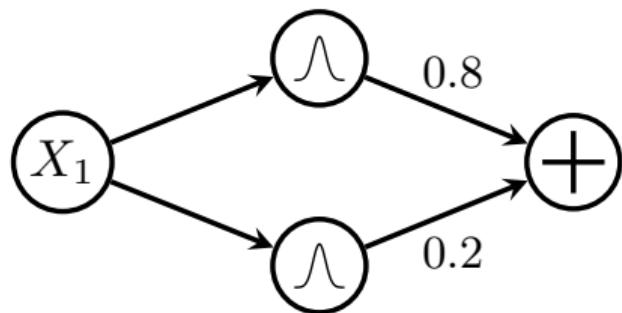


$$p(X) = w_1 \cdot p_1(X_1) + w_2 \cdot p_2(X_1)$$

⇒ translating inference to data structures...

GMMs

as computational graphs

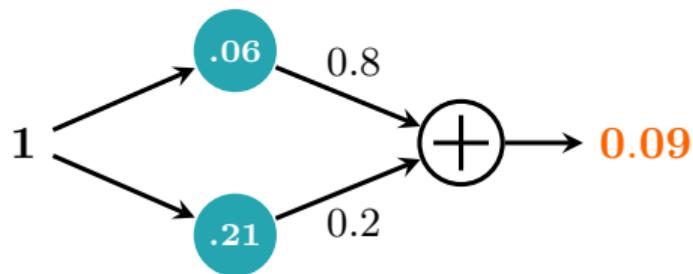


$$p(X_1) = 0.2 \cdot p_1(X_1) + 0.8 \cdot p_2(X_1)$$

⇒ ...e.g., as a weighted sum unit over Gaussian input distributions

GMMS

as computational graphs

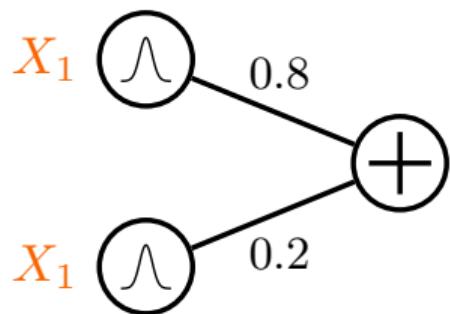


$$p(X = 1) = 0.2 \cdot p_1(X_1 = 1) + 0.8 \cdot p_2(X_1 = 1)$$

⇒ inference = feedforward evaluation

GMMs

as computational graphs

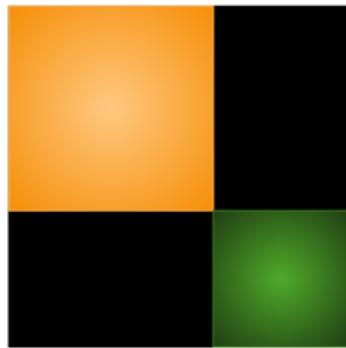
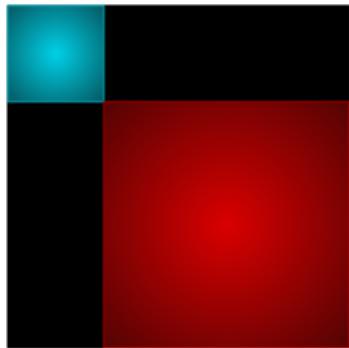


A simplified notation:

- ⇒ **scopes** attached to inputs
- ⇒ edge directions omitted

GMMS

as computational graphs

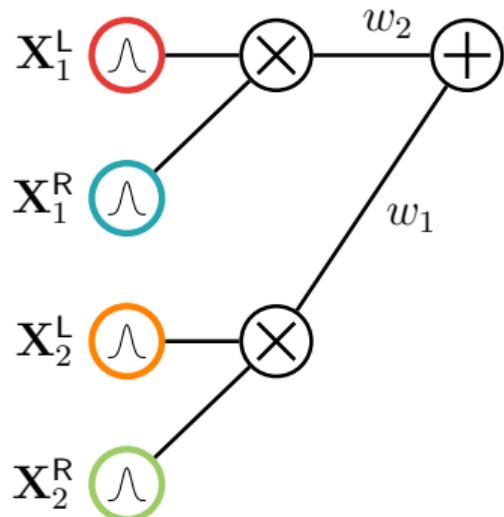


$$p(\mathbf{X}) = w_1 \cdot p_1(\mathbf{X}_1^L) \cdot p_1(\mathbf{X}_1^R) + \\ w_2 \cdot p_2(\mathbf{X}_2^L) \cdot p_2(\mathbf{X}_2^R)$$

⇒ local factorizations...

GMMs

as computational graphs



$$p(\mathbf{X}) = w_1 \cdot p_1(\mathbf{X}_1^L) \cdot p_1(\mathbf{X}_1^R) + \\ w_2 \cdot p_2(\mathbf{X}_2^L) \cdot p_2(\mathbf{X}_2^R)$$

⇒ ...are product units

probabilistic circuits (PCs)

a grammar for tractable computational graphs

I. A simple tractable function is a circuit

⇒ e.g., a multivariate Gaussian or
orthonormal polynomial



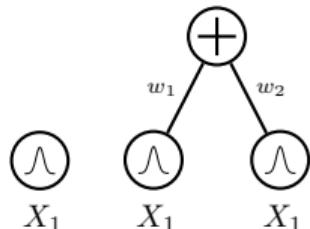
X_1

probabilistic circuits (PCs)

a grammar for tractable computational graphs

I. A simple tractable function is a circuit

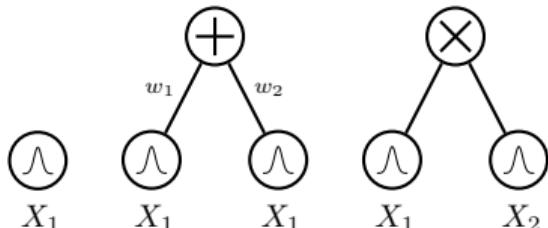
II. A weighted combination of circuits is a circuit



probabilistic circuits (PCs)

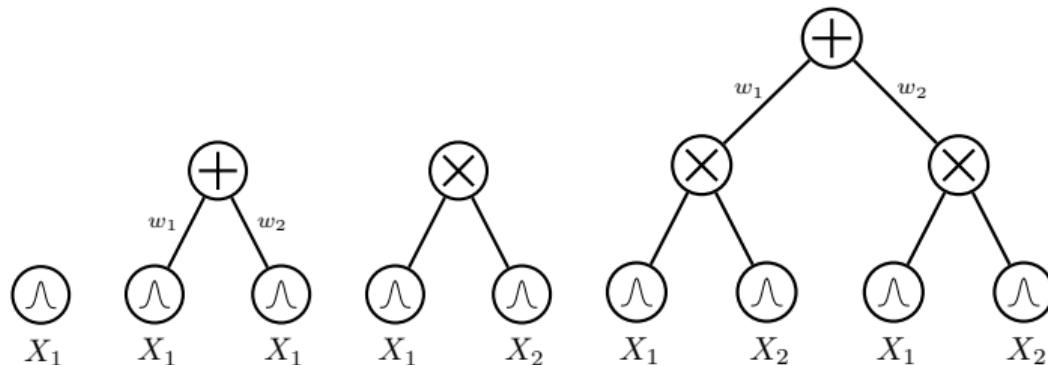
a grammar for tractable computational graphs

- I. A simple tractable function is a circuit
- II. A weighted combination of circuits is a circuit
- III. A product of circuits is a circuit



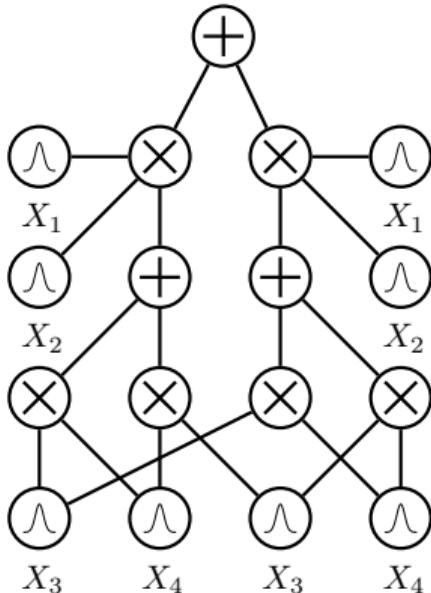
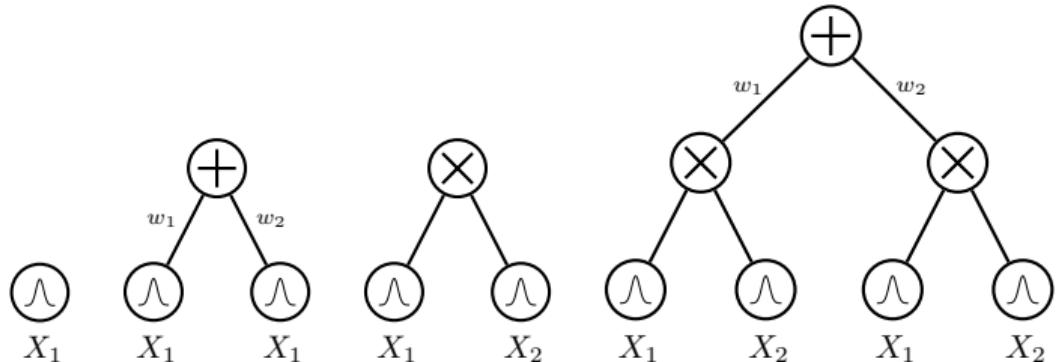
probabilistic circuits (PCs)

a grammar for tractable computational graphs

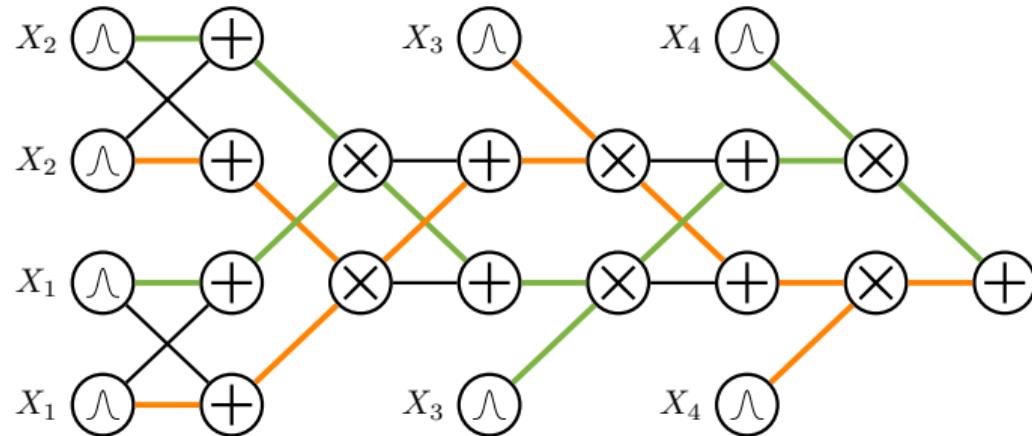


probabilistic circuits (PCs)

a grammar for tractable computational graphs

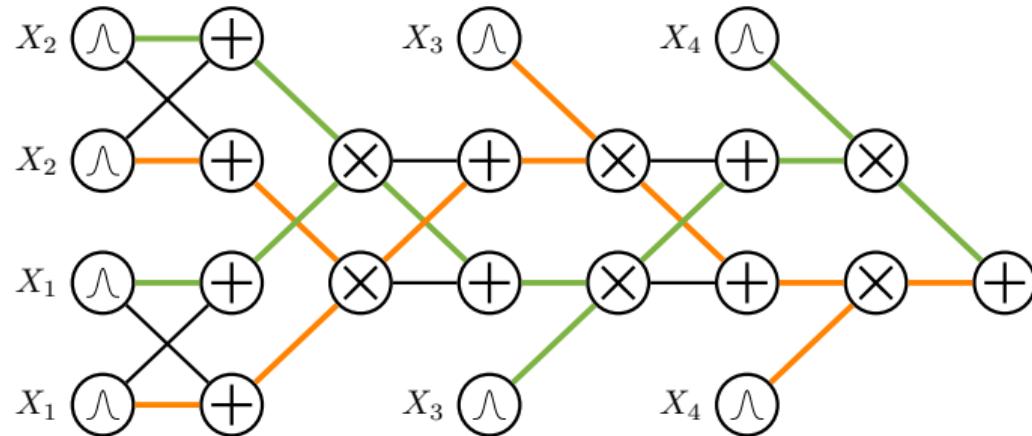


deep mixtures



$$p(\mathbf{x}) = \sum_{\mathcal{T}} \left(\prod_{w_j \in \mathbf{w}_{\mathcal{T}}} w_j \right) \prod_{l \in \text{leaves}(\mathcal{T})} p_l(\mathbf{x})$$

deep mixtures



an exponential number of mixture components!

...why PCs?

1. A grammar for tractable models

One formalism to represent many probabilistic models

⇒ #HMMs #Trees #XGBoost, Tensor Networks, ...

...why PCs?

1. A grammar for tractable models

One formalism to represent many probabilistic models

→ #HMMs #Trees #XGBoost, Tensor Networks, ...

2. Tractability == Structural Properties!!!

Exact computations of reasoning tasks are certified by guaranteeing certain structural properties. #marginals #expectations #MAP, #product ...

...why PCs?

1. A grammar for tractable models

One formalism to represent many probabilistic models

→ #HMMs #Trees #XGBoost, Tensor Networks, ...

2. Tractability == Structural Properties!!!

Exact computations of reasoning tasks are certified by guaranteeing certain structural properties. #marginals #expectations #MAP, #product ...

3. Reliable neuro-symbolic AI

logical constraints as circuits, multiplied to probabilistic circuits

circuits
(and variants)
everywhere

Semantic Probabilistic Layers for Neuro-Symbolic Learning

Kareem Ahmed

CS Department
UCLA

ahmedk@cs.ucla.edu

Stefano Teso

CIMeC and DISI
University of Trento

stefano.teso@unitn.it

Kai-Wei Chang

CS Department
UCLA

kwchang@cs.ucla.edu

Guy Van den Broeck

CS Department
UCLA

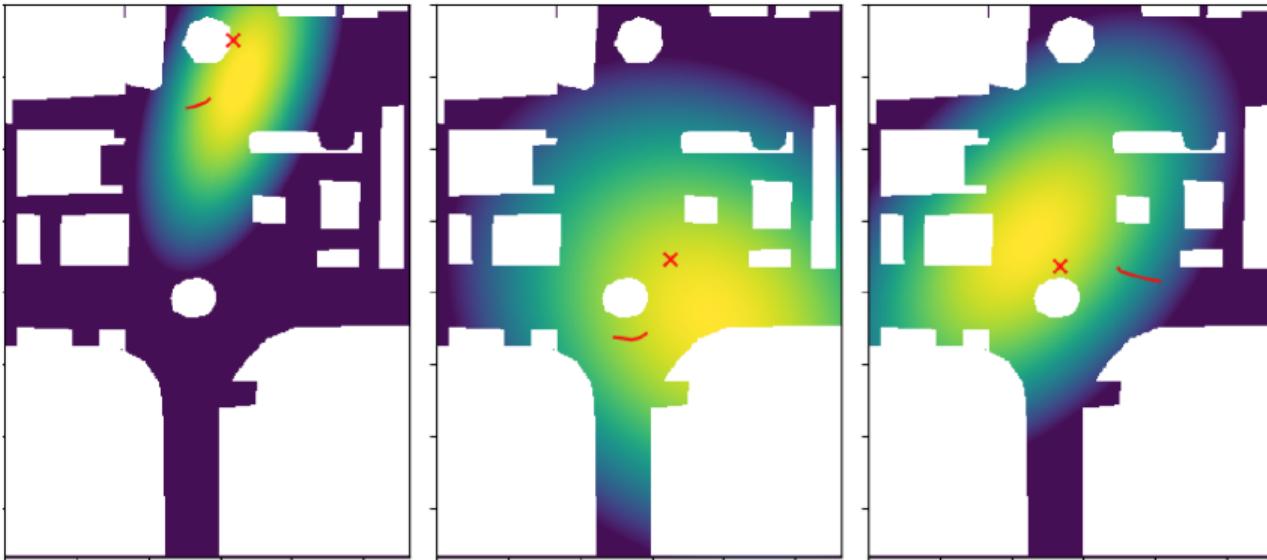
guyvdb@cs.ucla.edu

Antonio Vergari

School of Informatics
University of Edinburgh

avergari@ed.ac.uk

enforce constraints in neural networks at NeurIPS 2022



$$p(\mathbf{y} \mid \mathbf{x}) = \mathbf{q}_{\Theta}(\mathbf{y} \mid g(\mathbf{z})) \cdot \mathbf{c}_K(\mathbf{x}, \mathbf{y}) / \mathcal{Z}(\mathbf{x})$$



Ground Truth



ResNet-18



Semantic Loss

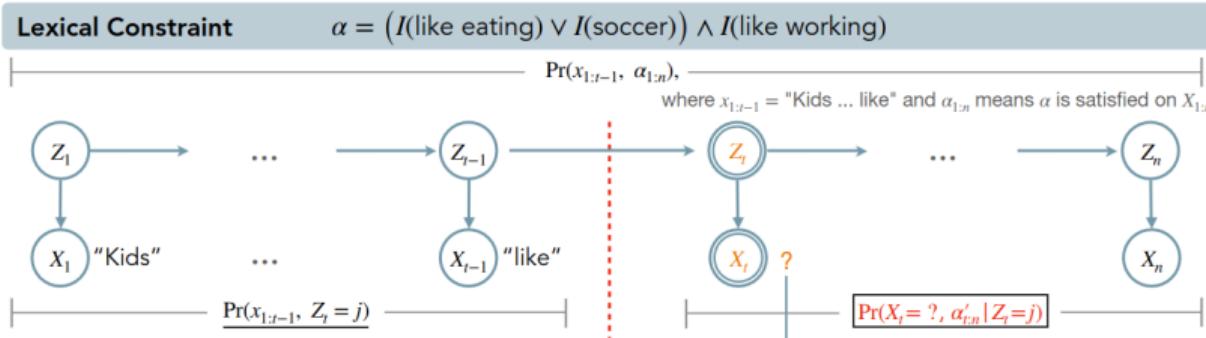


circuits

predictions guarantee a logical constraint 100% of the time!

Tractable Control for Autoregressive Language Generation

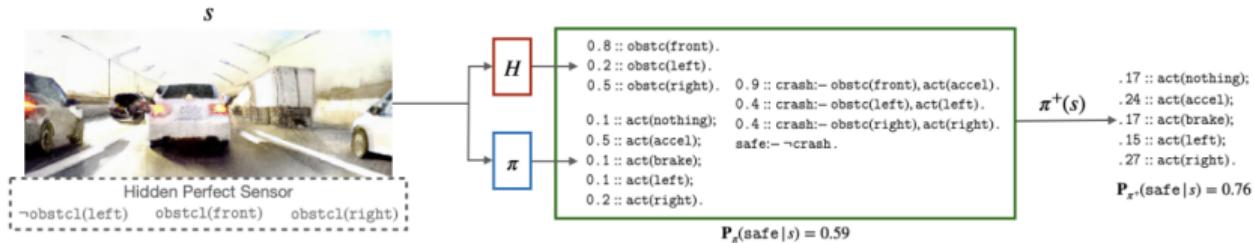
Honghua Zhang^{* 1} Meihua Dang^{* 1} Nanyun Peng¹ Guy Van den Broeck¹



constrained text generation with LLMs (ICML 2023)

Safe Reinforcement Learning via Probabilistic Logic Shields

Wen-Chi Yang¹, Giuseppe Marra¹, Gavin Rens and Luc De Raedt^{1,2}



reliable reinforcement learning (AAAI 23)

How to Turn Your Knowledge Graph Embeddings into Generative Models

Lorenzo Loconte
University of Edinburgh, UK
l.loconte@sms.ed.ac.uk

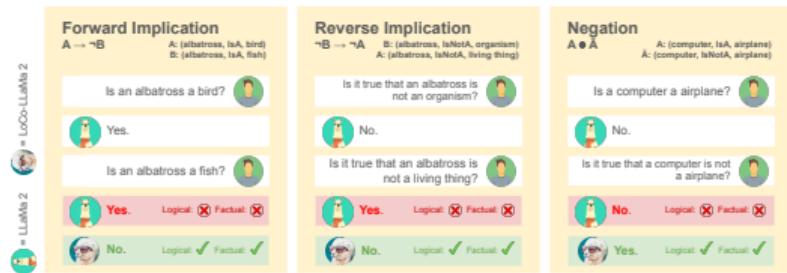
Nicola Di Mauro
University of Bari, Italy
nicola.dimauro@uniba.it

Robert Peharz
TU Graz, Austria
robert.peharz@tugraz.at

Antonio Vergari
University of Edinburgh, UK
avergari@ed.ac.uk

*enforce constraints in knowledge graph embeddings
oral at NeurIPS 2023*

Logically Consistent Language Models via Neuro-Symbolic Integration

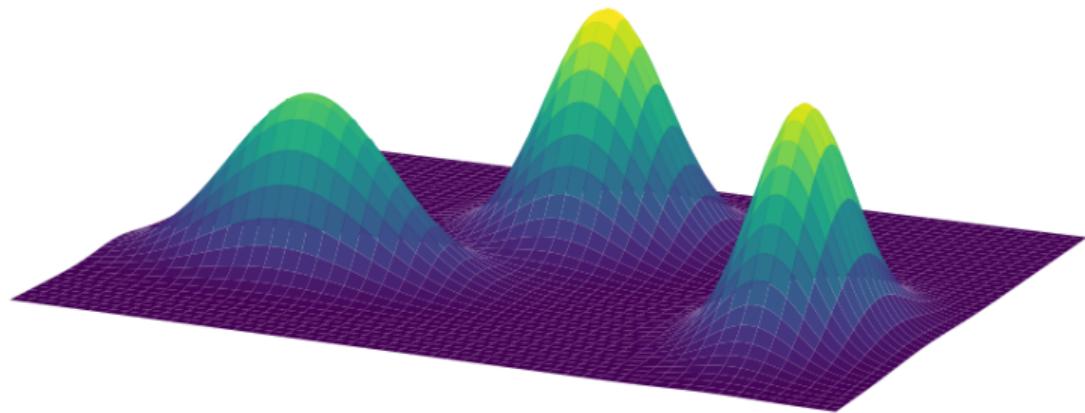


improving logical (self-)consistency in LLMs at ICLR 2025



learning & reasoning with circuits in pytorch

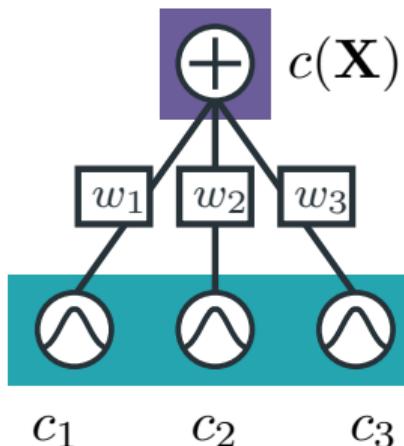
github.com/april-tools/cirkit



$$c(\mathbf{X}) = \sum_{i=1}^K w_i c_i(\mathbf{X}), \quad \text{with } w_i \geq 0, \quad \sum_{i=1}^K w_i = 1$$

additive MMs

are so cool!



easily represented as shallow PCs

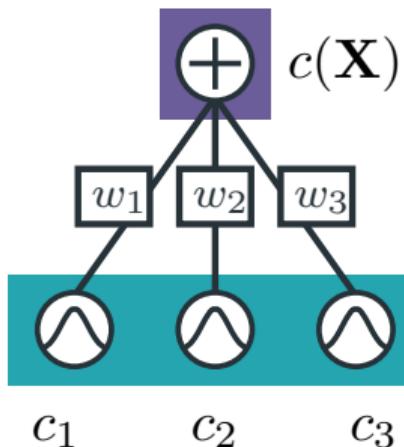
these are **monotonic** PCs

if marginals/conditionals are tractable for the components, they are tractable for the MM

universal approximators...

additive MMs

are so cool!



easily represented as shallow PCs

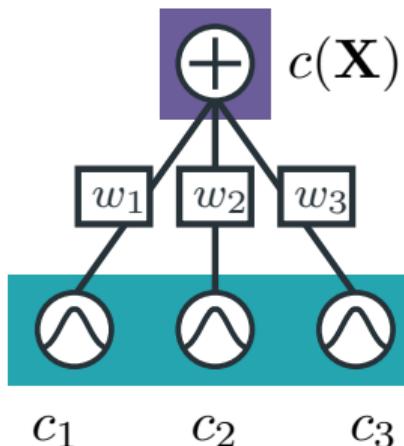
these are **monotonic** PCs

if marginals/conditionals are tractable for the components, they are tractable for the MM

universal approximators...

additive MMs

are so cool!



easily represented as shallow PCs

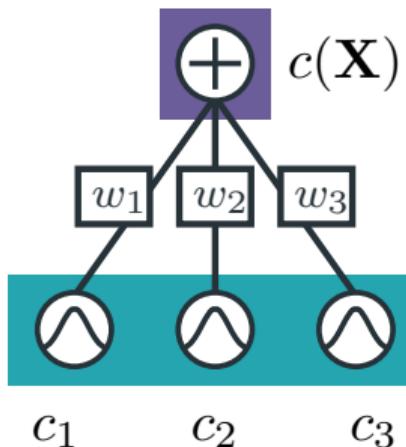
these are **monotonic** PCs

if marginals/conditionals are tractable for the components, they are tractable for the MM

universal approximators...

additive MMs

are so cool!



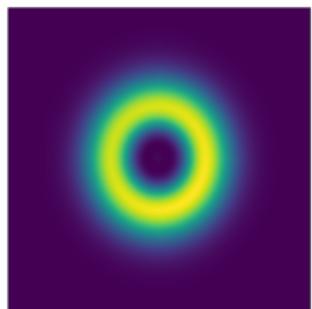
easily represented as shallow PCs

these are **monotonic** PCs

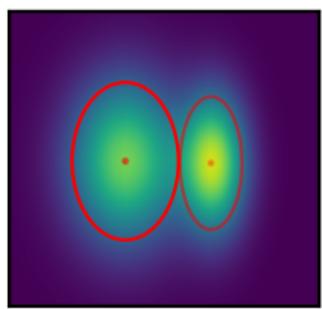
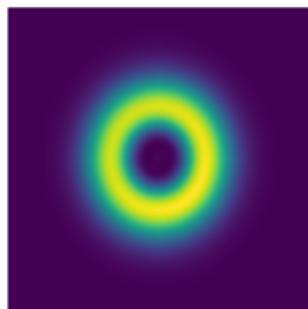
if marginals/conditionals are tractable for the components, they are tractable for the MM

universal approximators...

however...

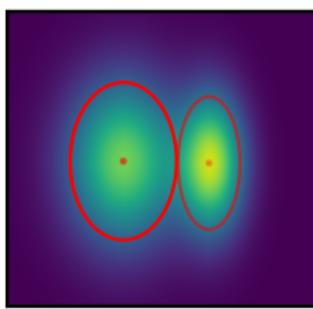
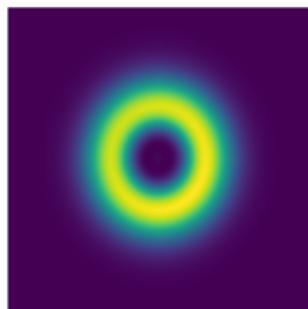


however...

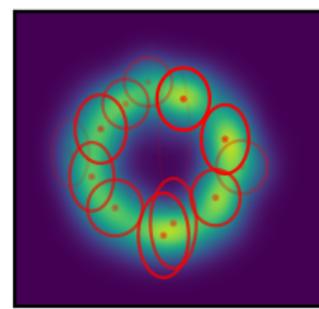


GMM ($K = 2$)

however...

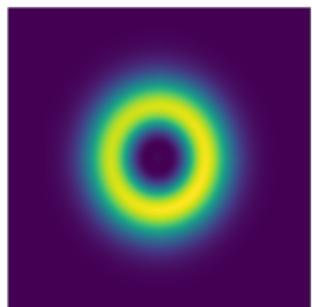


GMM ($K = 2$)

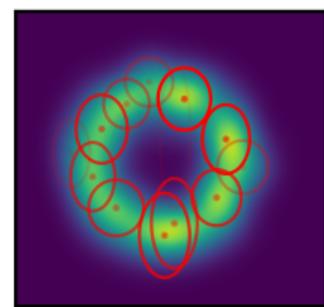


GMM ($K = 16$)

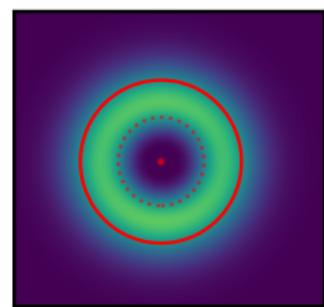
however...



GMM ($K = 2$)



GMM ($K = 16$)

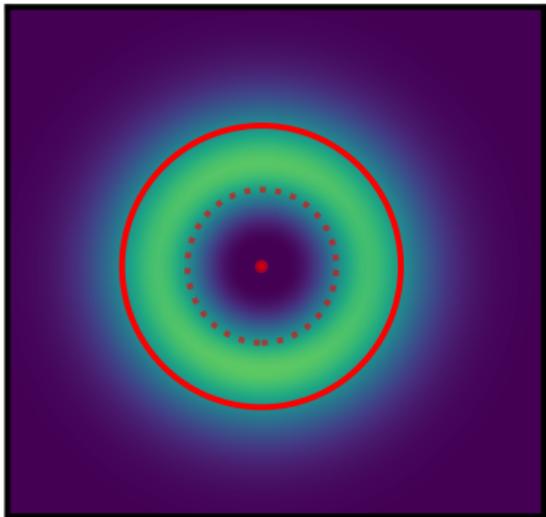


nGMM² ($K = 2$)

spoiler

**shallow mixtures
with negative parameters
can be *exponentially more compact* than
deep ones with positive parameters.**

subtractive MMs



also called negative/signed/**subtractive** MMs

⇒ or **non-monotonic** circuits,...

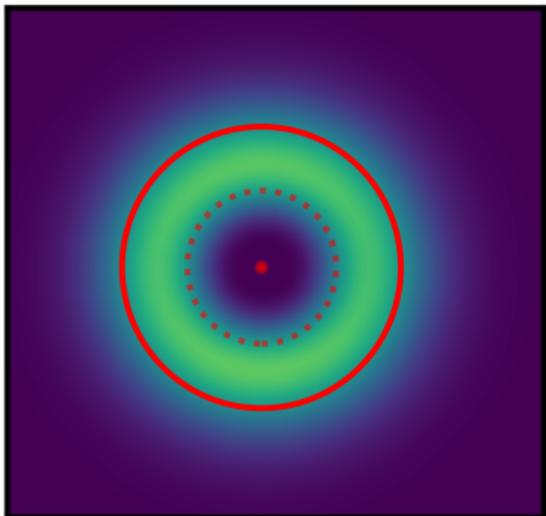
issue: how to preserve non-negative outputs?

well understood for simple parametric forms

e.g., Weibulls, Gaussians

⇒ constraints on variance, mean

subtractive MMs



also called negative/signed/**subtractive** MMs

⇒ or **non-monotonic** circuits,...

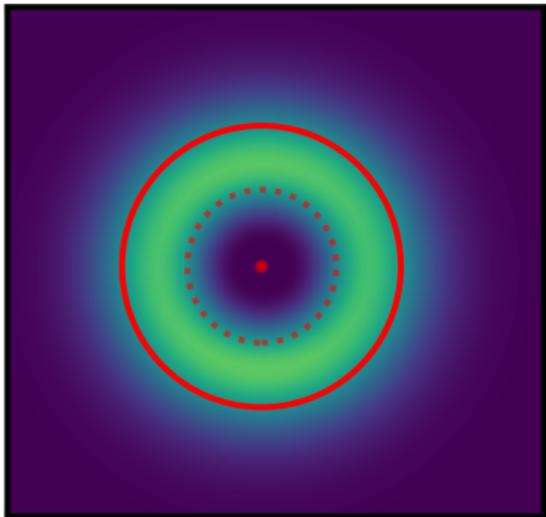
issue: how to preserve non-negative outputs?

well understood for simple parametric forms

e.g., Weibulls, Gaussians

⇒ constraints on variance, mean

subtractive MMs



also called negative/signed/**subtractive** MMs

⇒ or **non-monotonic** circuits,...

issue: how to preserve non-negative outputs?

well understood for simple parametric forms

e.g., Weibulls, Gaussians

⇒ constraints on variance, mean

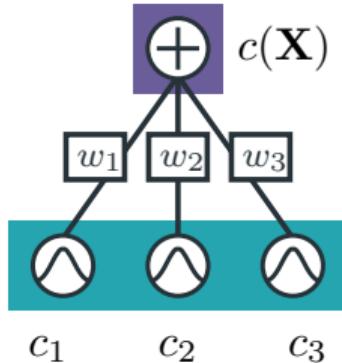
tl;dr

***“Understand when and how
we can use negative parameters
in deep subtractive mixture models”***

tl;dr

***“Understand when and how
we can use negative parameters
in deep **non-monotonic squared circuits**”***

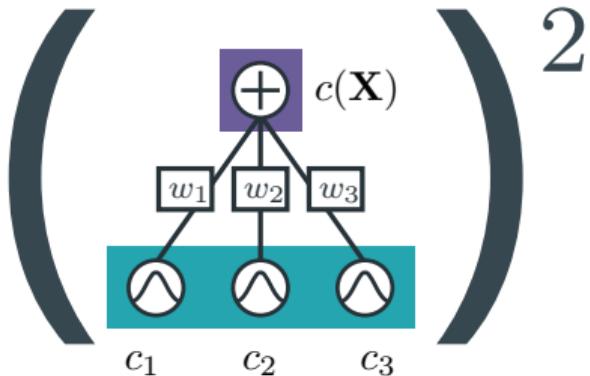
subtractive MMs as circuits



a **non-monotonic** smooth and (structured) decomposable circuit
⇒ possibly with negative outputs

$$c(\mathbf{X}) = \sum_{i=1}^K w_i c_i(\mathbf{X}), \quad w_i \in \mathbb{R},$$

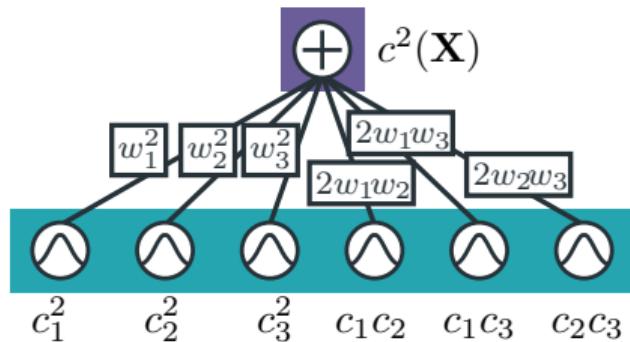
squaring shallow MMs



$$c^2(\mathbf{X}) = \left(\sum_{i=1}^K w_i c_i(\mathbf{X}) \right)^2$$

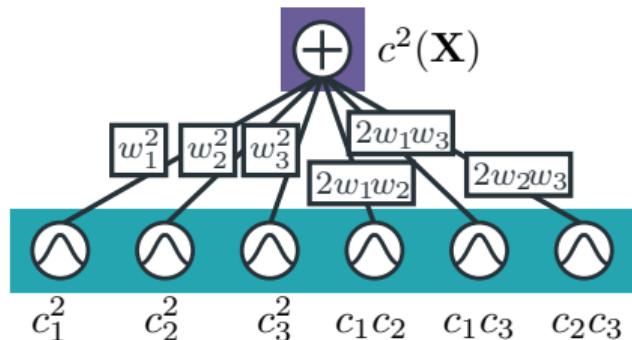
⇒ ensure non-negative output

squaring shallow MMs



$$\begin{aligned} c^2(\mathbf{X}) &= \left(\sum_{i=1}^K w_i c_i(\mathbf{X}) \right)^2 \\ &= \sum_{i=1}^K \sum_{j=1}^K w_i w_j c_i(\mathbf{X}) c_j(\mathbf{X}) \end{aligned}$$

squaring shallow MMs

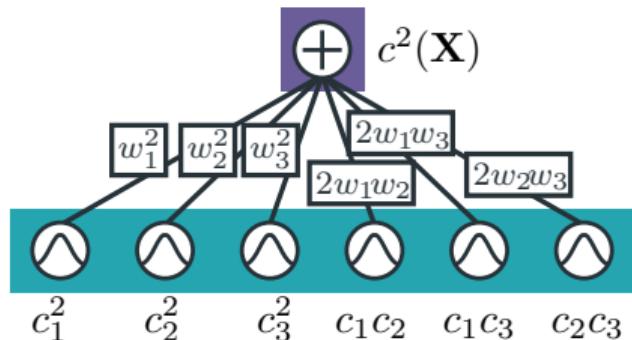


$$\begin{aligned} c^2(\mathbf{X}) &= \left(\sum_{i=1}^K w_i c_i(\mathbf{X}) \right)^2 \\ &= \sum_{i=1}^K \sum_{j=1}^K w_i w_j c_i(\mathbf{X}) c_j(\mathbf{X}) \end{aligned}$$

still a smooth and (str) decomposable PC with $\mathcal{O}(K^2)$ components!

\Rightarrow but still $\mathcal{O}(K)$ parameters

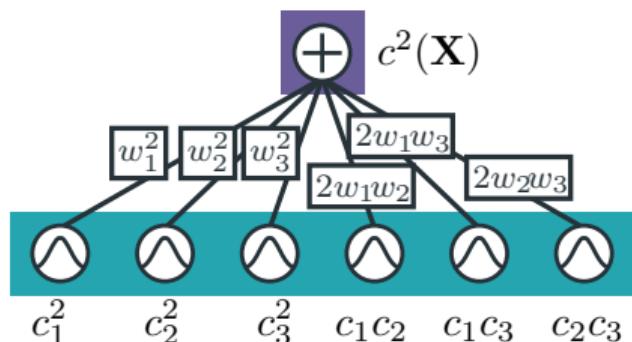
squaring shallow MMs



$$\begin{aligned} c^2(\mathbf{X}) &= \left(\sum_{i=1}^K w_i c_i(\mathbf{X}) \right)^2 \\ &= \sum_{i=1}^K \sum_{j=1}^K w_i w_j c_i(\mathbf{X}) c_j(\mathbf{X}) \end{aligned}$$

to **renormalize**, we have to compute $\sum_i \sum_j w_i w_j \int c_i(\mathbf{x}) c_j(\mathbf{x}) d\mathbf{x}$

squaring shallow MMs



$$\begin{aligned} c^2(\mathbf{X}) &= \left(\sum_{i=1}^K w_i c_i(\mathbf{X}) \right)^2 \\ &= \sum_{i=1}^K \sum_{j=1}^K w_i w_j c_i(\mathbf{X}) c_j(\mathbf{X}) \end{aligned}$$

to **renormalize**, we have to compute $\sum_i \sum_j w_i w_j \int c_i(\mathbf{x}) c_j(\mathbf{x}) d\mathbf{x}$
⇒ or we pick c_i, c_j to be **orthonormal**...!

EigenVI: score-based variational inference with orthogonal function expansions

Diana Cai
Flatiron Institute
dcai@flatironinstitute.org

Chirag Modi
Flatiron Institute
cmodi@flatironinstitute.org

Charles C. Margossian
Flatiron Institute
cmargossian@flatironinstitute.org

Robert M. Gower
Flatiron Institute
rgower@flatironinstitute.org

David M. Blei
Columbia University
david.blei@columbia.edu

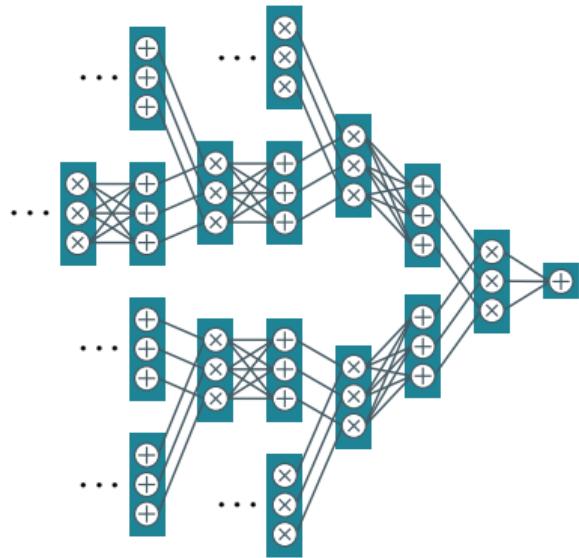
Lawrence K. Saul
Flatiron Institute
lsaul@flatironinstitute.org

orthonormal squared mixtures for VI

wait...

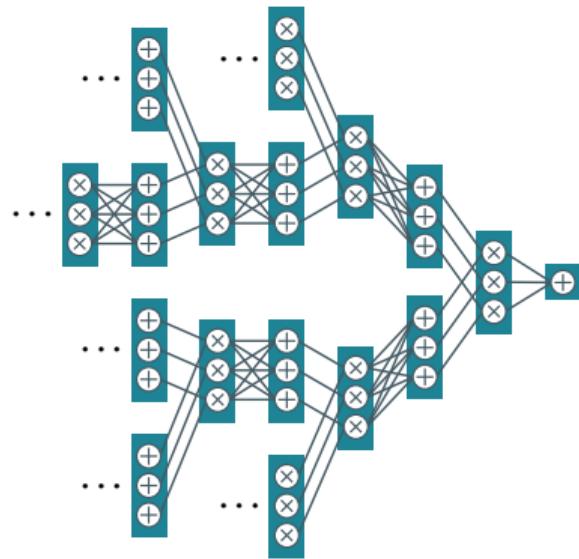
***“do negative parameters
really boost expressiveness?
and...always?”***

theorem



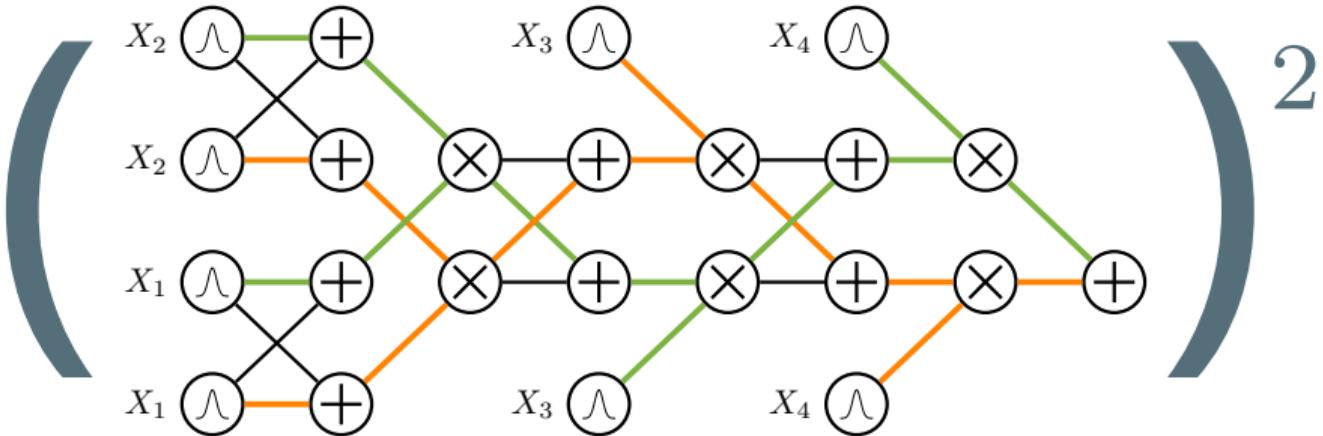
$\exists p$ requiring exponentially large
monotonic circuits...

theorem



$$\left(\begin{array}{c} \text{...} \\ \text{...} \\ \text{...} \end{array} \right)^2$$

**...but compact
squared non-monotonic circuits**



how to efficiently square (and *renormalize*) a deep PC?

compositional inference I



```
1 from cirkit.symbolic.functional import integrate, multiply  
2  
3 #  
4 # create a deep circuit  
5 c = build_symbolic_circuit('quad-tree-4')  
6  
7 #  
8 # compute the partition function of c^2  
9 def renormalize(c):  
10     c2 = multiply(c, c)  
11     return integrate(c2)
```

probabilistic circuits (PCs)

the unit-wise definition

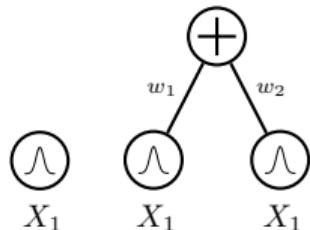
I. A simple tractable function is a circuit

$$\bigcirc \wedge \\ X_1$$

probabilistic circuits (PCs)

the unit-wise definition

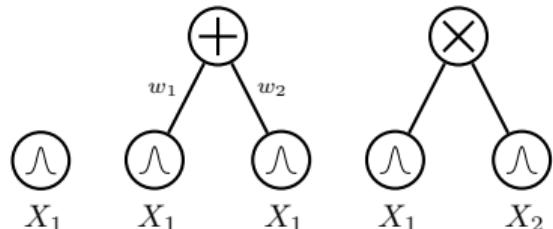
- I. A simple tractable function is a circuit
- II. A weighted combination of circuits is a circuit



probabilistic circuits (PCs)

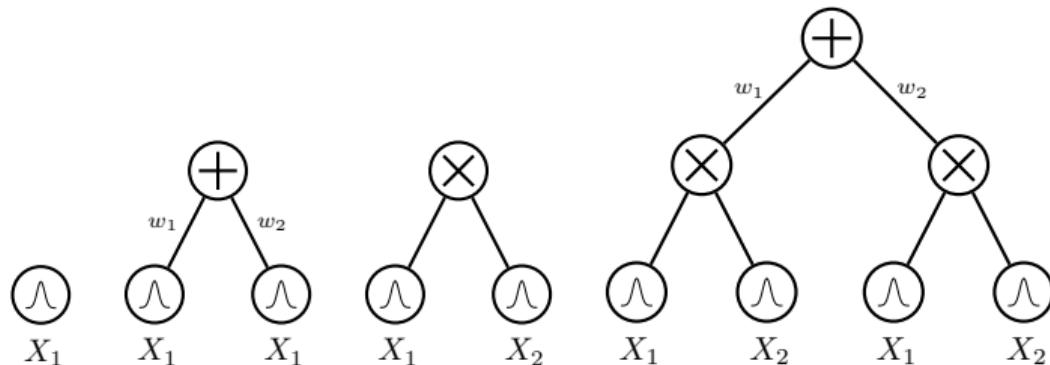
the unit-wise definition

- I. A simple tractable function is a circuit
- II. A weighted combination of circuits is a circuit
- III. A product of circuits is a circuit



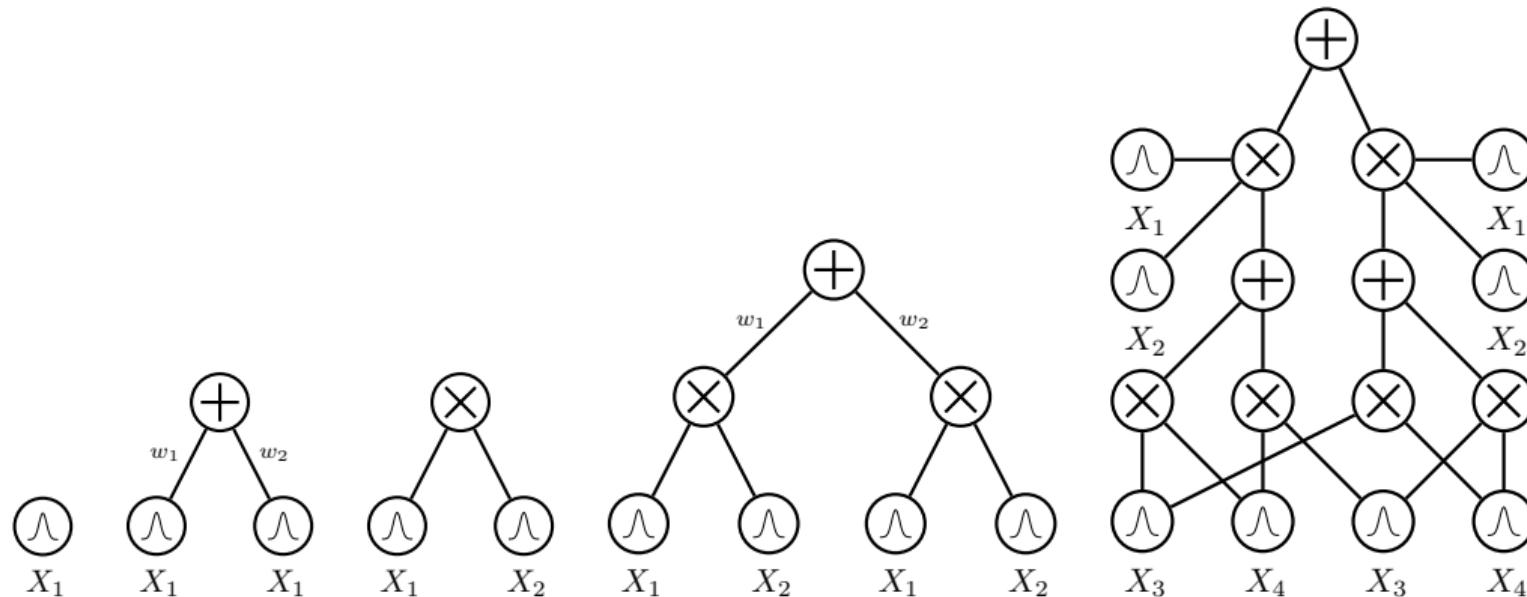
probabilistic circuits (PCs)

the unit-wise definition



probabilistic circuits (PCs)

the unit-wise definition



probabilistic circuits (PCs)

a tensorized definition

- I. A set of tractable functions is a circuit layer



probabilistic circuits (PCs)

a tensorized definition

- I. A set of tractable functions is a circuit layer
- II. A linear projection of a layer is a circuit layer

$$c(\mathbf{x}) = \mathbf{W}l(\mathbf{x})$$



probabilistic circuits (PCs)

a tensorized definition

- I. A set of tractable functions is a circuit layer
- II. A linear projection of a layer is a circuit layer

$$c(\mathbf{x}) = \mathbf{W}l(\mathbf{x})$$



probabilistic circuits (PCs)

a tensorized definition

- I. A set of tractable functions is a circuit layer
- II. A linear projection of a layer is a circuit layer
- III. The product of two layers is a circuit layer

$$c(\mathbf{x}) = \mathbf{l}(\mathbf{x}) \odot \mathbf{r}(\mathbf{x}) \quad // \text{ Hadamard}$$



probabilistic circuits (PCs)

a tensorized definition

- I. A set of tractable functions is a circuit layer
- II. A linear projection of a layer is a circuit layer
- III. The product of two layers is a circuit layer

$$c(\mathbf{x}) = \mathbf{l}(\mathbf{x}) \odot \mathbf{r}(\mathbf{x}) \quad // \text{ Hadamard}$$

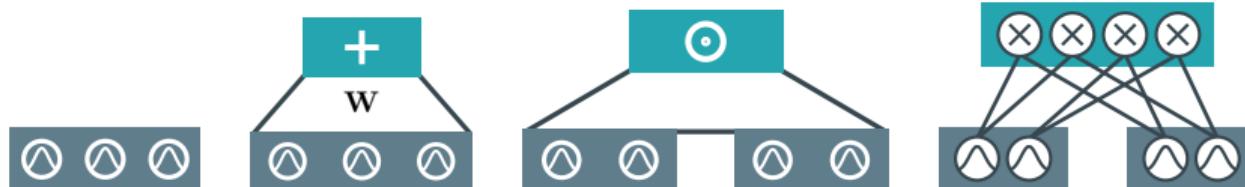


probabilistic circuits (PCs)

a tensorized definition

- I. A set of tractable functions is a circuit layer
- II. A linear projection of a layer is a circuit layer
- III. The product of two layers is a circuit layer

$$c(\mathbf{x}) = \text{vec}(\mathbf{l}(\mathbf{x})\mathbf{r}(\mathbf{x})^\top) \quad // \text{ Kronecker}$$

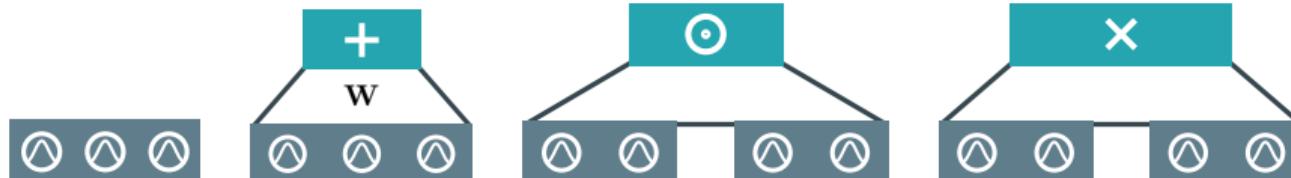


probabilistic circuits (PCs)

a tensorized definition

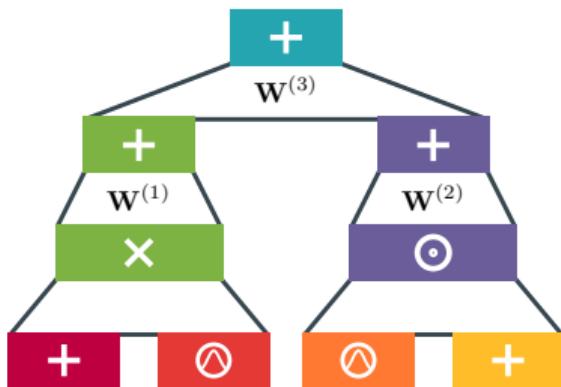
- I. A set of tractable functions is a circuit layer
- II. A linear projection of a layer is a circuit layer
- III. The product of two layers is a circuit layer

$$c(\mathbf{x}) = \text{vec}(\mathbf{l}(\mathbf{x})\mathbf{r}(\mathbf{x})^\top) \quad // \text{ Kronecker}$$



probabilistic circuits (PCs)

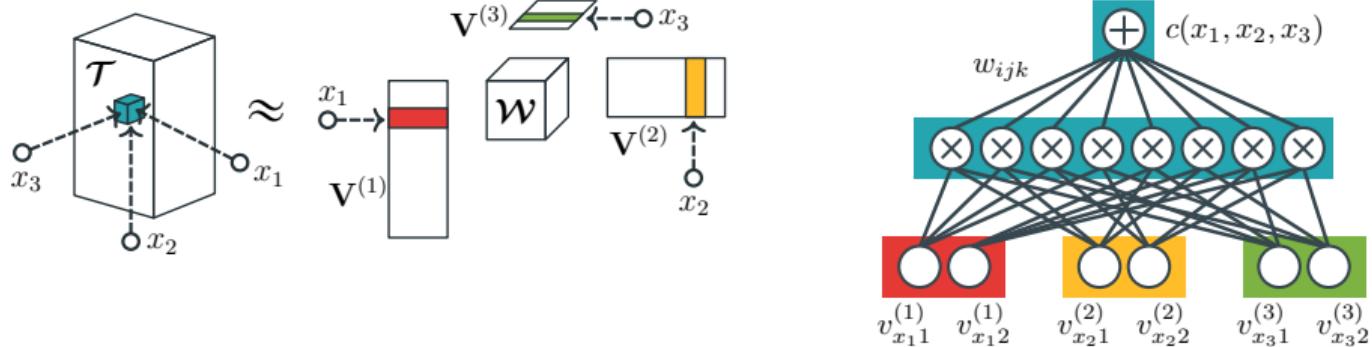
a tensorized definition

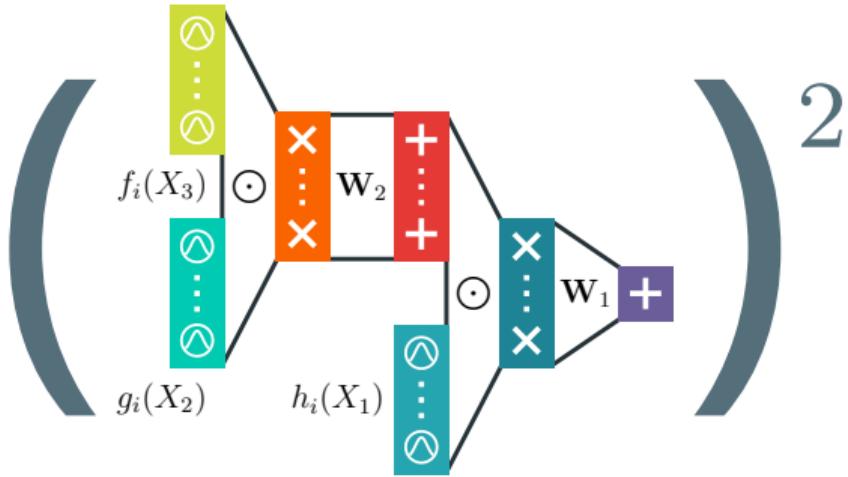


- I. A set of tractable functions is a circuit layer
 - II. A linear projection of a layer is a circuit layer
 - III. The product of two layers is a circuit layer
- stack layers to build a deep circuit!**

circuits layers

as tensor factorizations

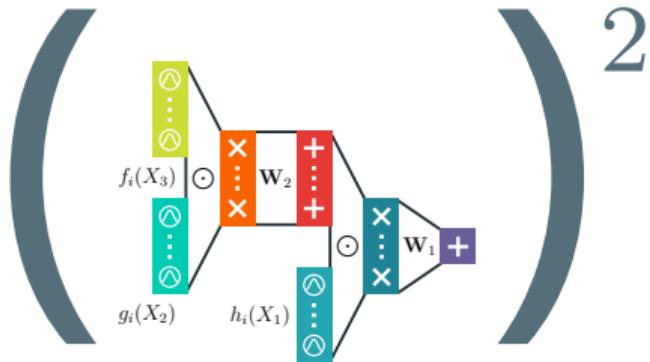




how to efficiently square (and *renormalize*) a deep PC?

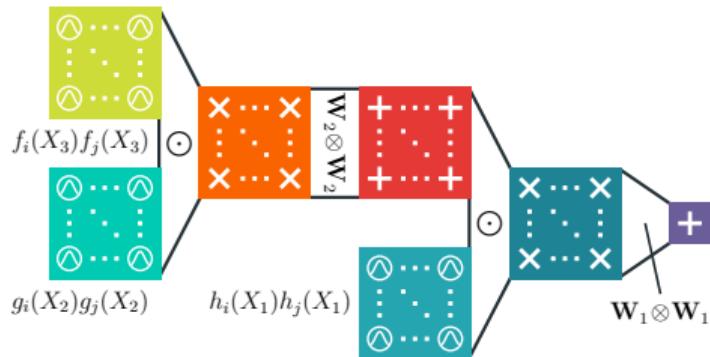
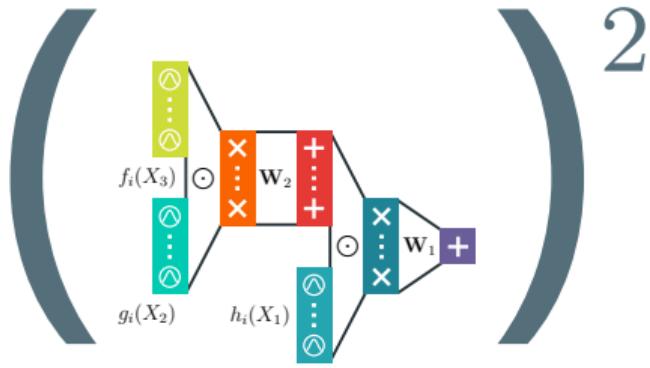
squaring deep PCs

the tensorized way



squaring deep PCs

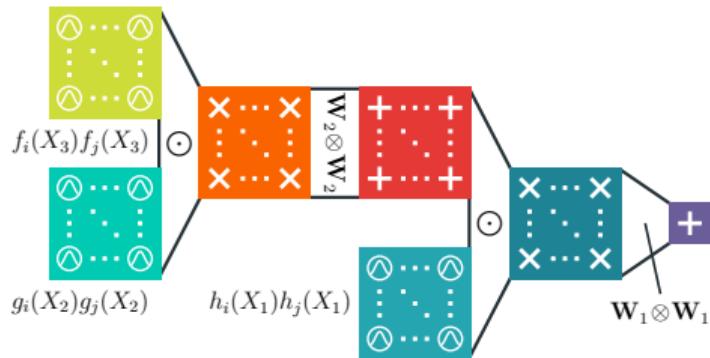
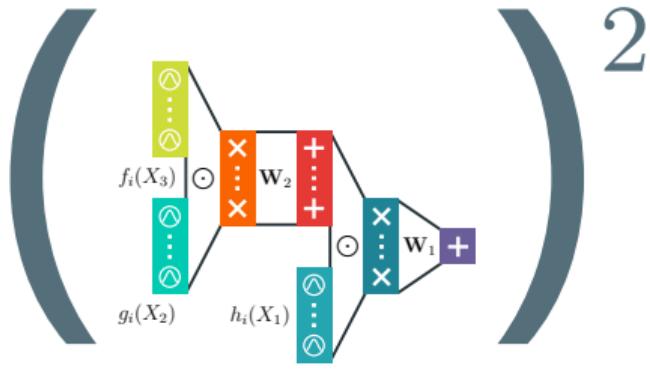
the tensorized way



squaring a circuit = squaring layers

squaring deep PCs

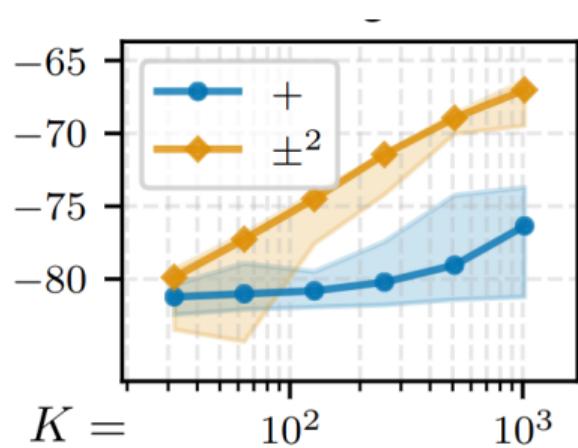
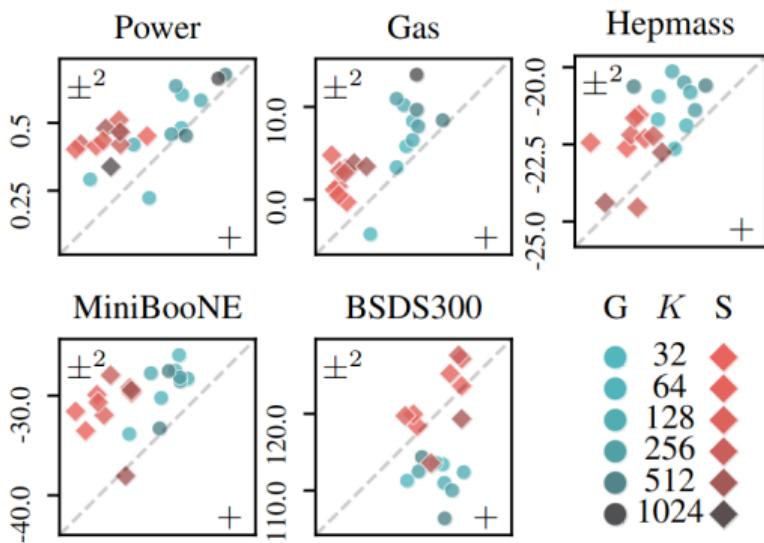
the tensorized way



exactly compute $\int c(x)c(x)dX$ **in time** $O(LK^2)$

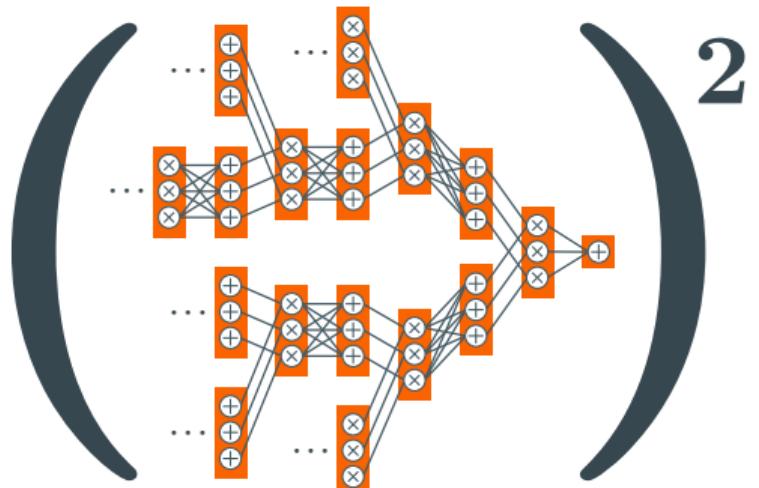
how more expressive?

for the ML crowd



theorem

$\exists p$ requiring exponentially large
squared non-mono circuits...

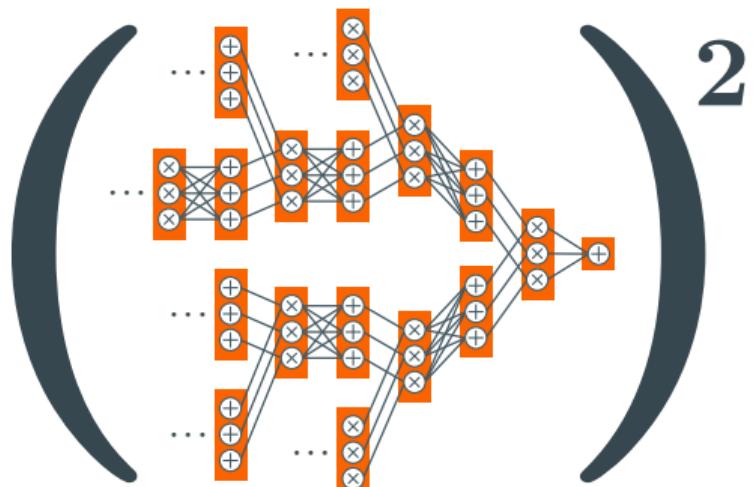


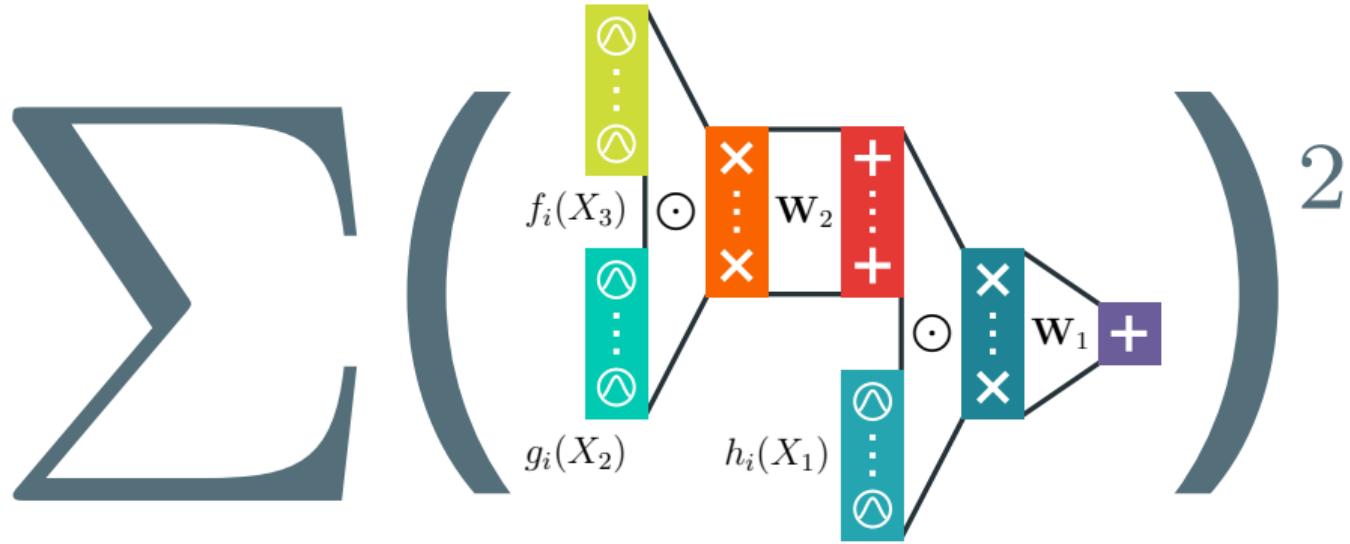
theorem



...but compact

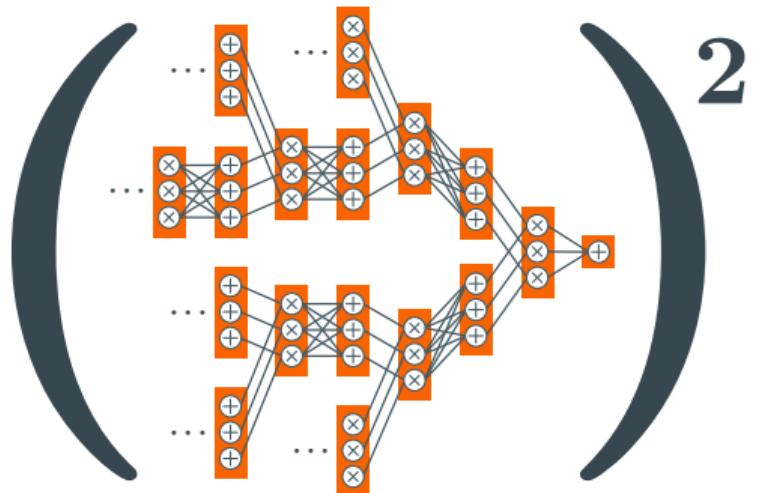
monotonic circuits...!





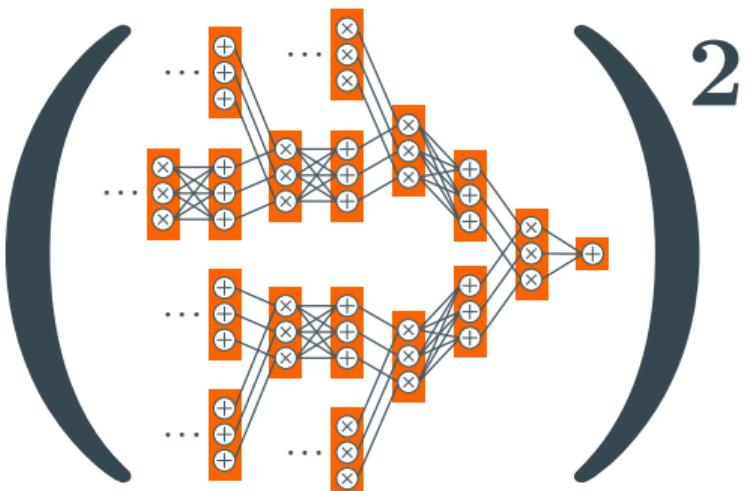
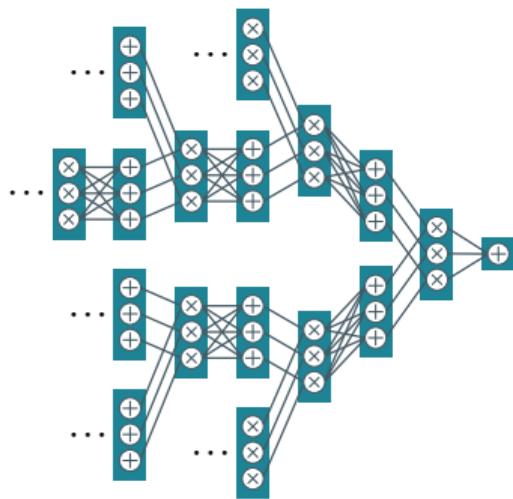
what if we use more than one square?

theorem



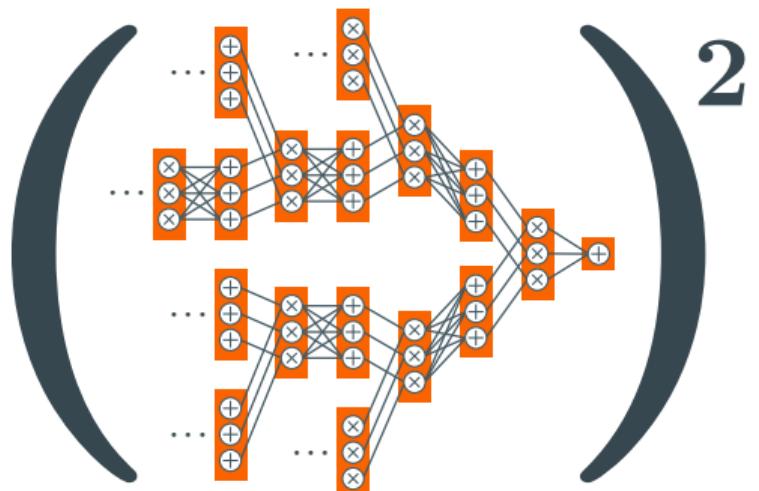
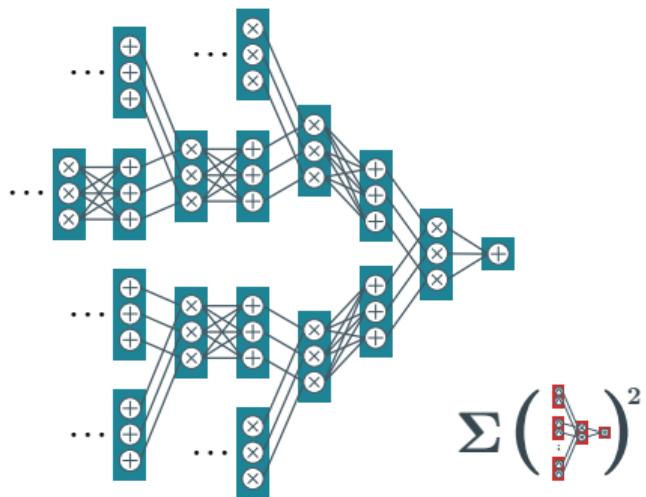
$\exists p$ requiring exponentially large **squared non-mono circuits...**

theorem



...exponentially large monotonic circuits...

theorem



...but compact **SOS circuits...**!

$$\pm_{\text{sd}} = \Delta \Sigma_{\text{cmp}}^2$$

(Theorem 5)

$$\Sigma_{\text{cmp}}^2 = \text{psd}$$

(Proposition 2)

$$+_{\text{sd}}$$

•
Open Question 1

•
Open Question 2

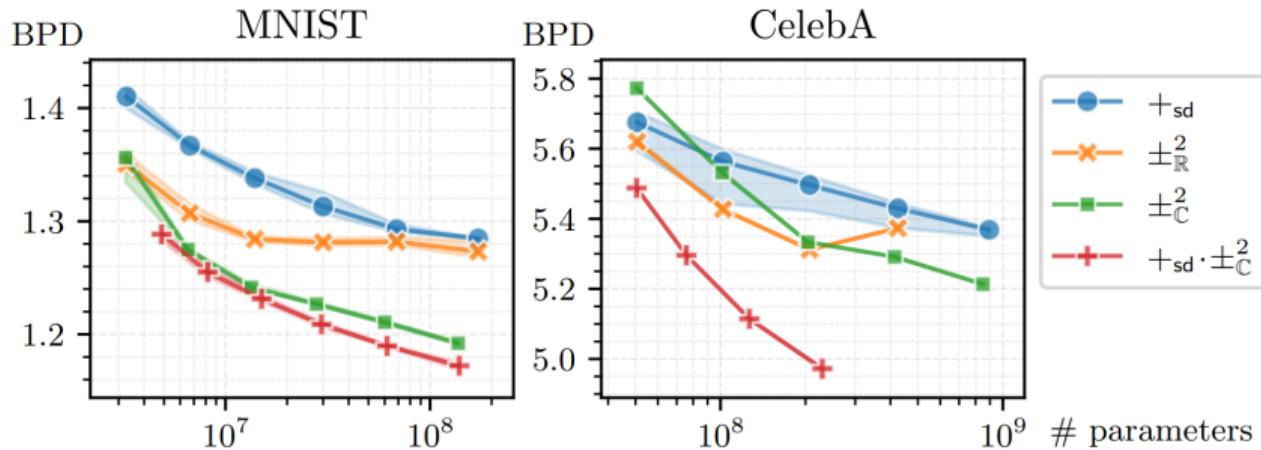
$$\pm_{\mathbb{R}}^2$$

• UDISJ
(Theorem 0)

• UPS
(Theorem 2)

• UTQ
(Theorem B.3)

a hierarchy of subtractive mixtures



complex circuits are SOS (and scale better!)

compositional inference I



```
1 from cirkit.symbolic.functional import integrate, multiply,
2     conjugate
3
4 # create a deep circuit with complex parameters
5 c = build_symbolic_complex_circuit('quad-tree-4')
6
7 # compute the partition function of c^2
8 def renormalize(c):
9     c1 = conjugate(c)
10    c2 = multiply(c, c1)
11    return integrate(c2)
```

EigenVI: score-based variational inference with orthogonal function expansions

Diana Cai
Flatiron Institute
dcai@flatironinstitute.org

Chirag Modi
Flatiron Institute
cmodi@flatironinstitute.org

Charles C. Margossian
Flatiron Institute
cmargossian@flatironinstitute.org

Robert M. Gower
Flatiron Institute
rgower@flatironinstitute.org

David M. Blei
Columbia University
david.blei@columbia.edu

Lawrence K. Saul
Flatiron Institute
lsaul@flatironinstitute.org

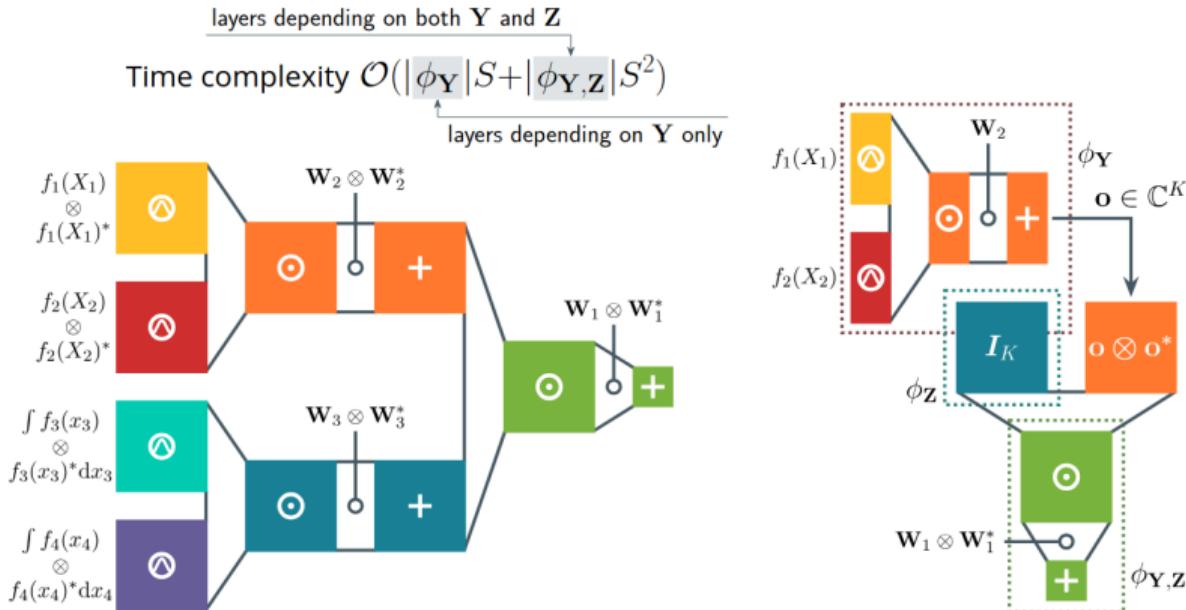
***what about deep orthonormal mixtures
and arbitrary marginals?***

On Faster Marginalization with Squared Circuits via Orthonormalization

Lorenzo Loconte¹ Antonio Vergari¹

¹ School of Informatics, University of Edinburgh, UK
l.loconte@sms.ed.ac.uk, avergari@ed.ac.uk

it suffices to orthonormalize each layer!



faster marginalization of arbitrary subsets of features

approximate inference

e.g., via sampling

Can we use a subtractive mixture model to approximate expectations?

$$\mathbb{E}_{\mathbf{x} \sim q(\mathbf{x})} [f(\mathbf{x})] \approx \frac{1}{S} \sum_{i=1}^S f(\mathbf{x}^{(i)}) \quad \text{with} \quad \mathbf{x}^{(i)} \sim q(\mathbf{x})$$

\Rightarrow but how to sample from q ?

approximate inference

e.g., via sampling

Can we use a subtractive mixture model to approximate expectations?

$$\mathbb{E}_{\mathbf{x} \sim q(\mathbf{x})} [f(\mathbf{x})] \approx \frac{1}{S} \sum_{i=1}^S f(\mathbf{x}^{(i)}) \quad \text{with} \quad \mathbf{x}^{(i)} \sim q(\mathbf{x})$$

\Rightarrow but how to sample from q ?

use **autoregressive inverse transform sampling**:

$$x_1 \sim q(x_1), \quad x_i \sim q(x_i | \mathbf{x}_{<i}) \quad \text{for } i \in \{2, \dots, d\}$$

\Rightarrow can be slow for large dimensions, requires **inverting the CDF**

approximate inference

difference of expectation estimator

Idea: represent q as a difference of two additive mixtures

$$q(\mathbf{x}) = Z_+ \cdot q_+(\mathbf{x}) - Z_- \cdot q_-(\mathbf{x})$$

⇒ *expectations will break down in two “parts”*

approximate inference

difference of expectation estimator

Idea: represent q as a difference of two additive mixtures

$$q(\mathbf{x}) = Z_+ \cdot q_+(\mathbf{x}) - Z_- \cdot q_-(\mathbf{x})$$

\Rightarrow *expectations will break down in two “parts”*

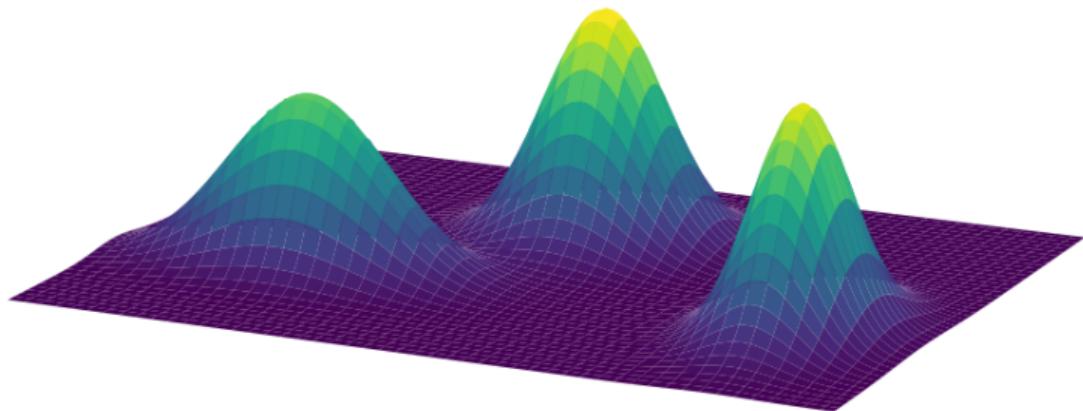
$$\frac{Z_+}{S_+} \sum_{s=1}^{S_+} f(\mathbf{x}_+^{(s)}) - \frac{Z_-}{S_-} \sum_{s=1}^{S_-} f(\mathbf{x}_-^{(s)}), \text{ where } \begin{aligned} \mathbf{x}_+^{(s)} &\sim q_+(\mathbf{x}_+) \\ \mathbf{x}_-^{(s)} &\sim q_-(\mathbf{x}_-) \end{aligned}, \quad (1)$$

approximate inference

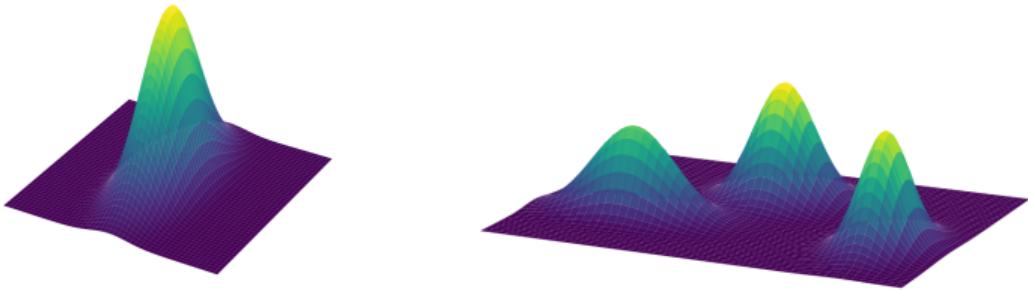
difference of expectation estimator

Method	d	Number of components (K)					
		2		4		6	
		$\log(\hat{I} - I)$	Time (s)	$\log(\hat{I} - I)$	Time (s)	$\log(\hat{I} - I)$	Time (s)
ΔExS	16	-19.507 ± 1.025	0.293 ± 0.004	-19.062 ± 0.823	1.049 ± 0.077	-19.497 ± 1.974	2.302 ± 0.159
ARITS	16	-19.111 ± 1.103	7.525 ± 0.038	-19.299 ± 1.611	7.52 ± 0.023	-18.739 ± 1.024	7.746 ± 0.032
ΔExS	32	-48.411 ± 1.265	0.325 ± 0.012	-48.046 ± 0.972	1.027 ± 0.107	-48.34 ± 0.814	2.213 ± 0.177
ARITS	32	-47.897 ± 1.165	15.196 ± 0.059	-47.349 ± 0.839	15.535 ± 0.059	-47.3 ± 0.978	17.371 ± 0.06
ΔExS	64	-108.095 ± 1.094	0.38 ± 0.034	-107.56 ± 0.616	0.9 ± 0.14	-107.653 ± 0.945	1.512 ± 0.383
ARITS	64	-107.898 ± 1.129	30.459 ± 0.098	-107.33 ± 0.929	33.892 ± 0.119	-107.374 ± 1.138	52.02 ± 0.127

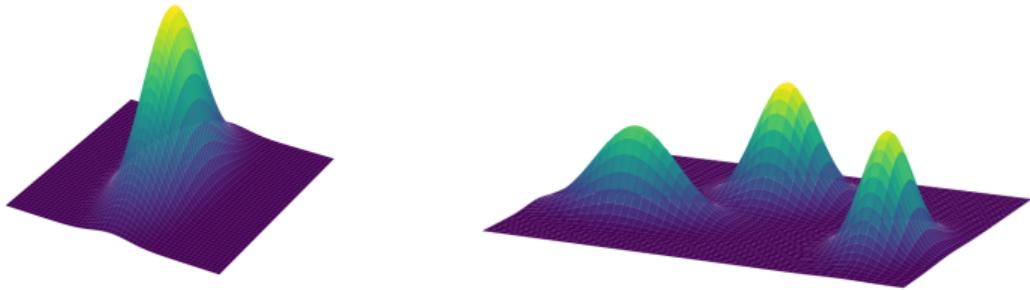
faster than autoregressive sampling



oh mixtures, you're so fine you blow my mind!



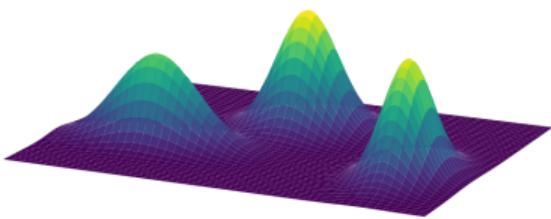
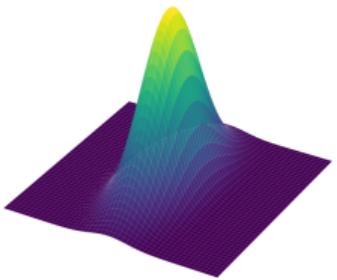
$$p(\mathbf{X}) \quad \xrightarrow{\text{orange arrow}} \quad \sum_{i=1}^K w_i p_i(\mathbf{X}) \quad w_i > 0$$



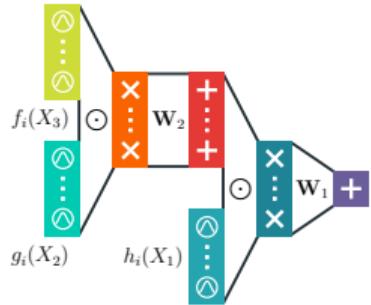
$$p(\mathbf{X}) \longrightarrow \sum_{i=1}^K w_i p_i(\mathbf{X}) \quad w_i > 0$$

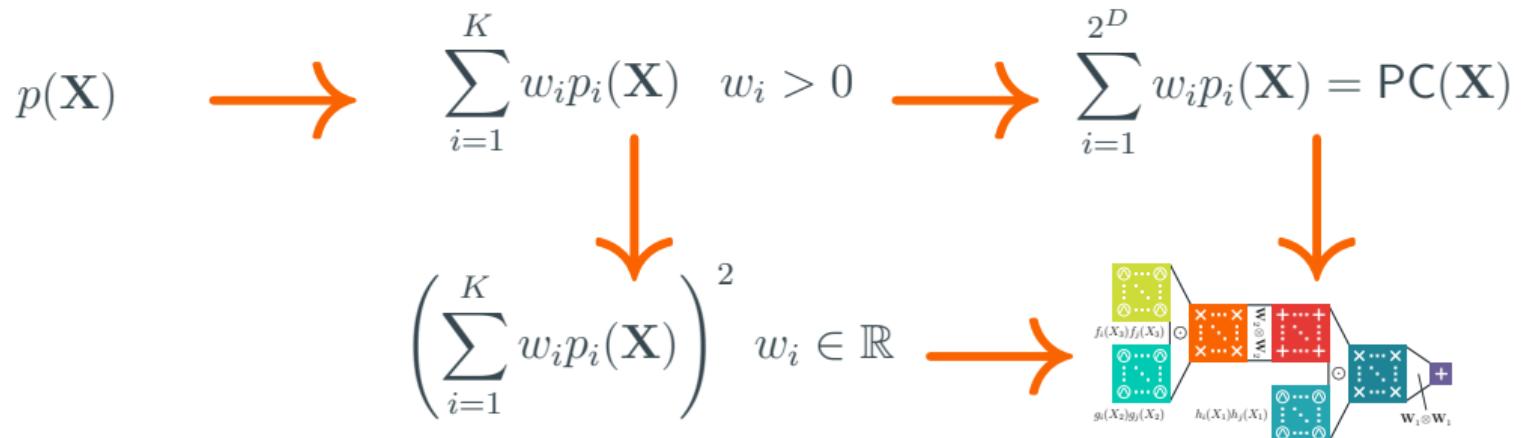
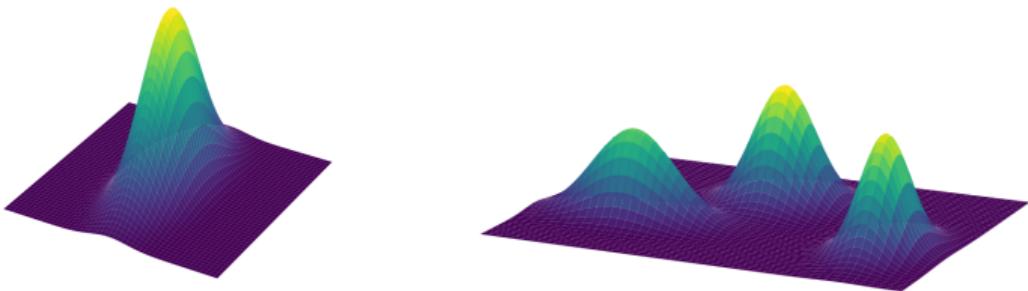
*"if someone publishes a paper on **model A**, there will be a paper about
mixtures of A soon, with high probability"*

A. Vergari



$$p(\mathbf{X}) \rightarrow \sum_{i=1}^K w_i p_i(\mathbf{X}) \quad w_i > 0 \rightarrow \sum_{i=1}^{2^D} w_i p_i(\mathbf{X}) = \text{PC}(\mathbf{X})$$

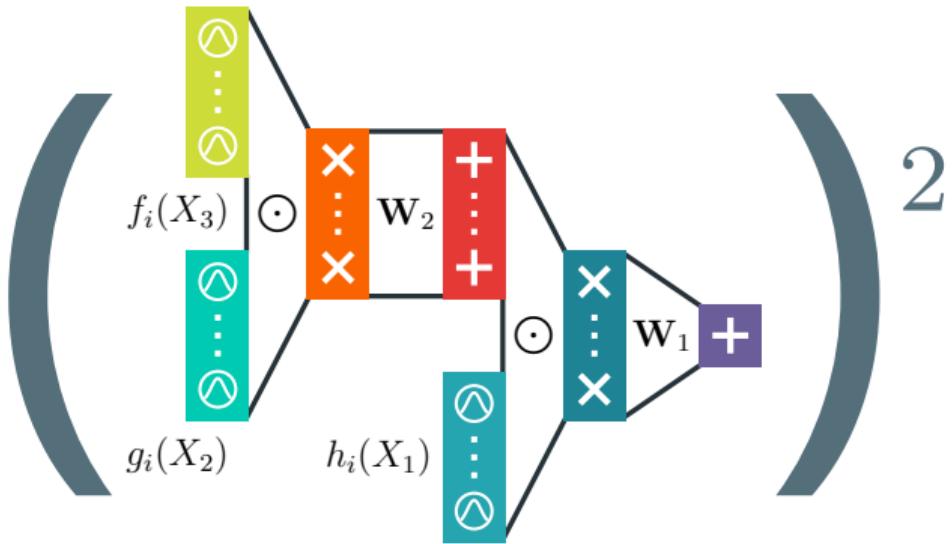






learning & reasoning with circuits in pytorch

github.com/april-tools/cirkit



questions?