



# *subtractive mixture models*

## *representation, learning & inference*

**antonio vergari** (he/him)

 @tetraduzione

4th Apr 2025 - 3rd GeMSS Sophia Antipolis

*april*

april-tools.github.io

# *april*

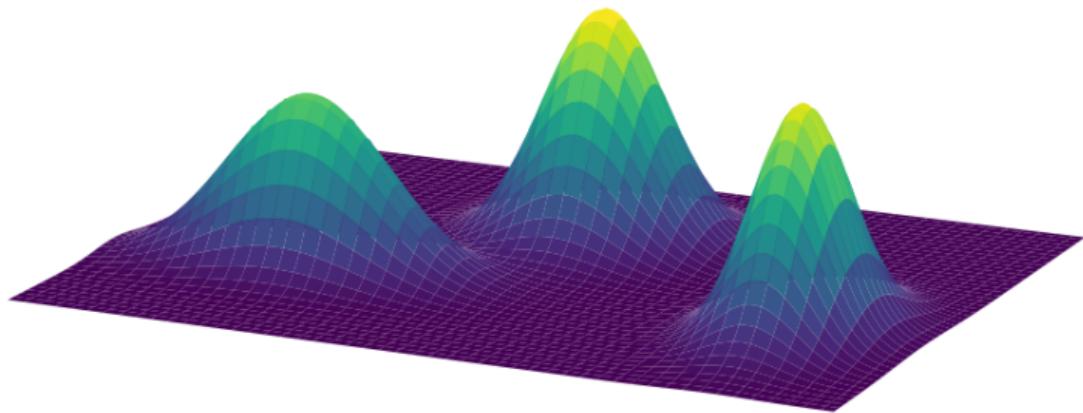
*about  
probabilities  
integrals &  
logic*

# *april*

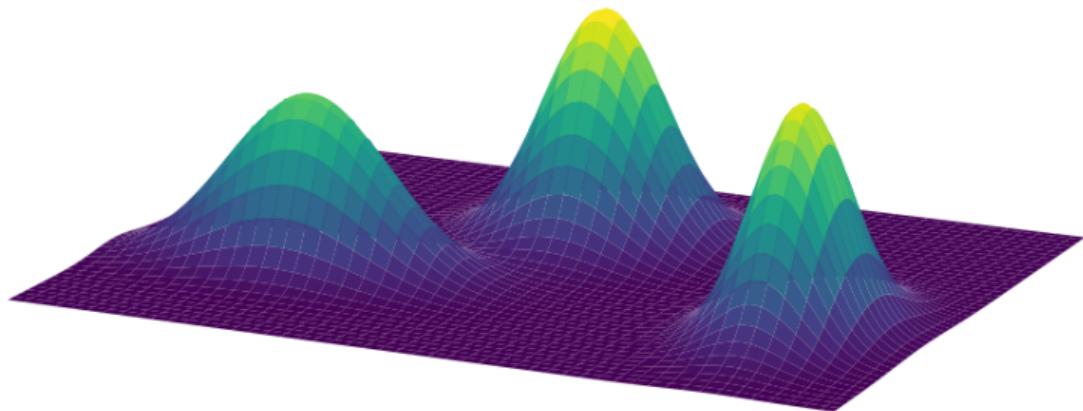
*autonomous &  
provably  
reliable  
intelligent  
learners*

# *april*

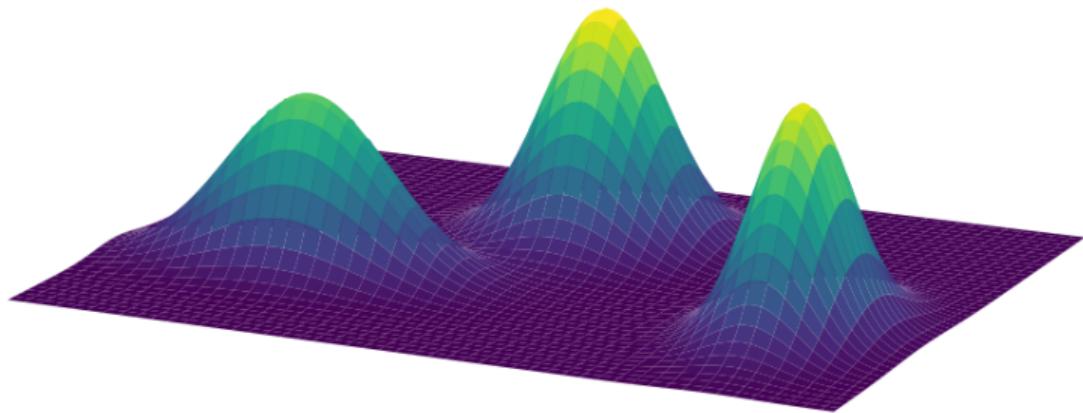
*april is  
probably a  
recursive  
identifier of a  
lab*



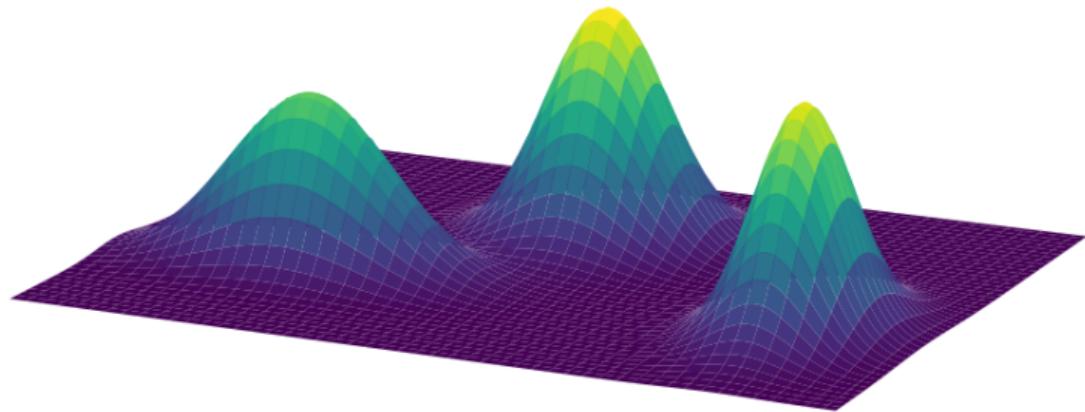
***who knows mixture models?***



**who *loves* mixture models?**



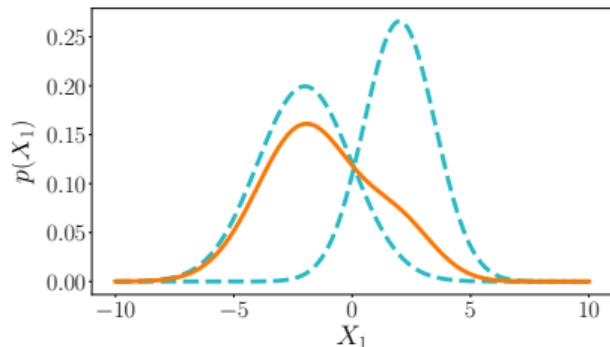
*a brief recap...*



$$c(\mathbf{X}) = \sum_{i=1}^K w_i c_i(\mathbf{X}), \quad \text{with} \quad w_i \geq 0, \quad \sum_{i=1}^K w_i = 1$$

# GMMS

as computational graphs

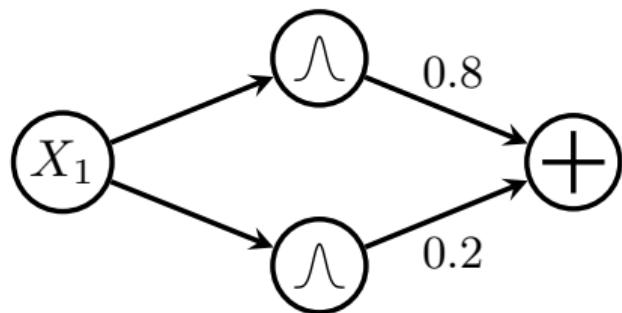


$$p(X) = w_1 \cdot p_1(X_1) + w_2 \cdot p_2(X_1)$$

⇒ translating inference to data structures...

# GMMs

as computational graphs

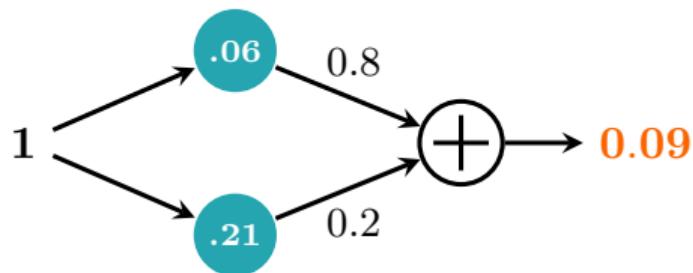


$$p(X_1) = 0.2 \cdot p_1(X_1) + 0.8 \cdot p_2(X_1)$$

⇒ ...e.g., as a weighted sum unit over Gaussian input distributions

# GMMS

as computational graphs

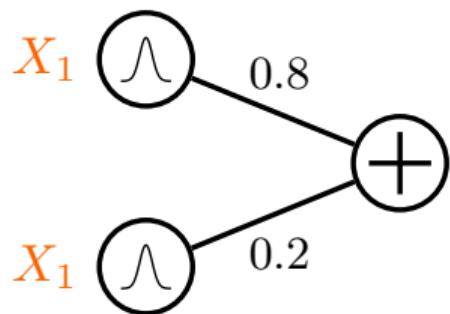


$$p(X = 1) = 0.2 \cdot p_1(X_1 = 1) + 0.8 \cdot p_2(X_1 = 1)$$

⇒ inference = feedforward evaluation

# GMMs

as computational graphs

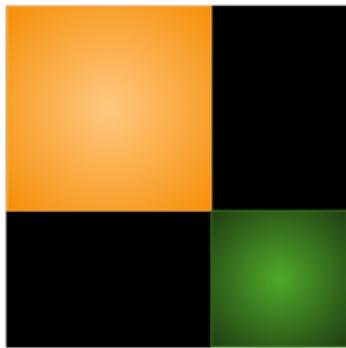
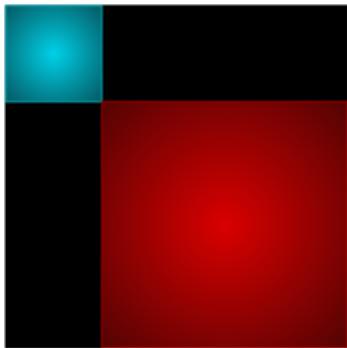


A simplified notation:

- ⇒ **scopes** attached to inputs
- ⇒ edge directions omitted

# GMMS

as computational graphs

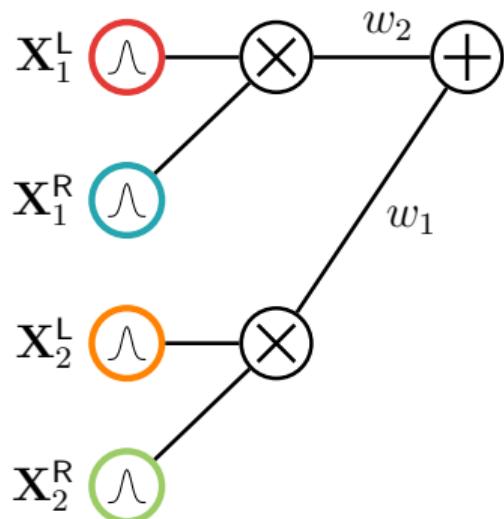


$$p(\mathbf{X}) = w_1 \cdot p_1(\mathbf{X}_1^L) \cdot p_1(\mathbf{X}_1^R) + \\ w_2 \cdot p_2(\mathbf{X}_2^L) \cdot p_2(\mathbf{X}_2^R)$$

⇒ local factorizations...

# GMMs

as computational graphs



$$p(\mathbf{X}) = w_1 \cdot p_1(\mathbf{X}_1^L) \cdot p_1(\mathbf{X}_1^R) + \\ w_2 \cdot p_2(\mathbf{X}_2^L) \cdot p_2(\mathbf{X}_2^R)$$

⇒ ...are product units

# **probabilistic circuits (PCs)**

*a grammar for tractable computational graphs*

I. A simple tractable function is a circuit

⇒ e.g., a multivariate Gaussian or  
orthonormal polynomial



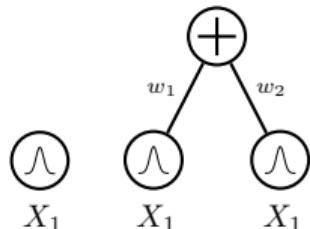
$X_1$

# **probabilistic circuits (PCs)**

*a grammar for tractable computational graphs*

I. A simple tractable function is a circuit

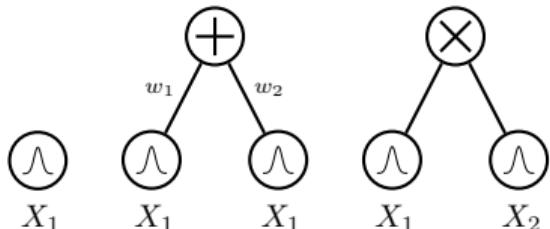
II. A weighted combination of circuits is a circuit



# **probabilistic circuits (PCs)**

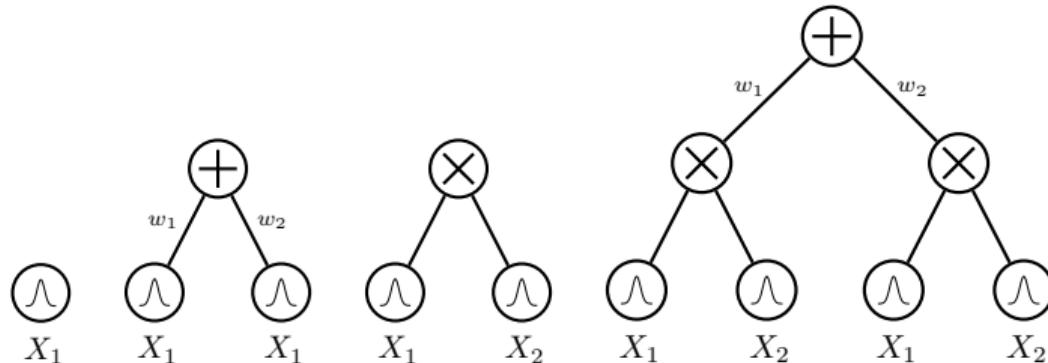
*a grammar for tractable computational graphs*

- I. A simple tractable function is a circuit
- II. A weighted combination of circuits is a circuit
- III. A product of circuits is a circuit



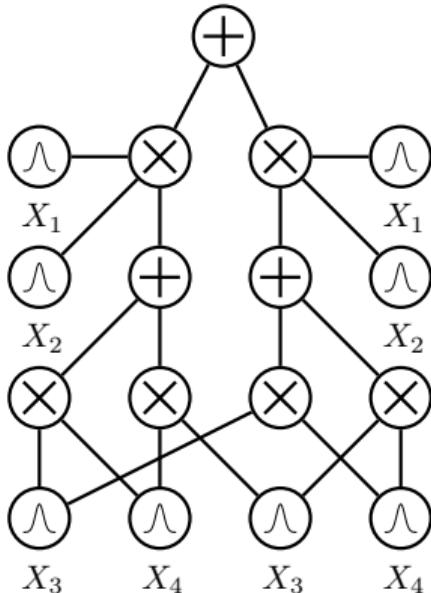
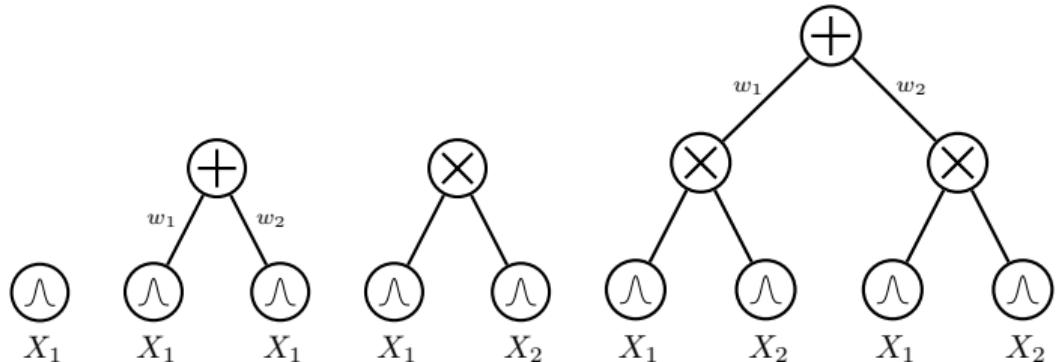
# **probabilistic circuits (PCs)**

*a grammar for tractable computational graphs*



# **probabilistic circuits (PCs)**

*a grammar for tractable computational graphs*



# **probabilistic circuits (PCs)**

*a tensorized definition*

- I. A set of tractable functions is a circuit layer



# **probabilistic circuits (PCs)**

*a tensorized definition*

- I. A set of tractable functions is a circuit layer
- II. A linear projection of a layer is a circuit layer

$$c(\mathbf{x}) = \mathbf{W}l(\mathbf{x})$$



# **probabilistic circuits (PCs)**

*a tensorized definition*

- I. A set of tractable functions is a circuit layer
- II. A linear projection of a layer is a circuit layer

$$c(\mathbf{x}) = \mathbf{W}l(\mathbf{x})$$



# **probabilistic circuits (PCs)**

*a tensorized definition*

- I. A set of tractable functions is a circuit layer
- II. A linear projection of a layer is a circuit layer
- III. The product of two layers is a circuit layer

$$c(\mathbf{x}) = \mathbf{l}(\mathbf{x}) \odot \mathbf{r}(\mathbf{x}) \quad // \text{ Hadamard}$$



# **probabilistic circuits (PCs)**

*a tensorized definition*

- I. A set of tractable functions is a circuit layer
- II. A linear projection of a layer is a circuit layer
- III. The product of two layers is a circuit layer

$$c(\mathbf{x}) = \mathbf{l}(\mathbf{x}) \odot \mathbf{r}(\mathbf{x}) \quad // \text{ Hadamard}$$

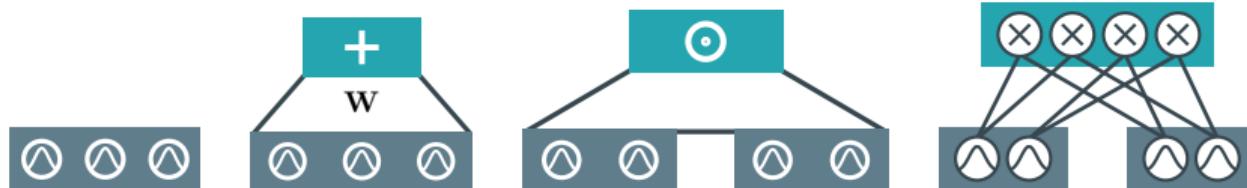


# **probabilistic circuits (PCs)**

*a tensorized definition*

- I. A set of tractable functions is a circuit layer
- II. A linear projection of a layer is a circuit layer
- III. The product of two layers is a circuit layer

$$c(\mathbf{x}) = \text{vec}(\mathbf{l}(\mathbf{x})\mathbf{r}(\mathbf{x})^\top) \quad // \text{ Kronecker}$$

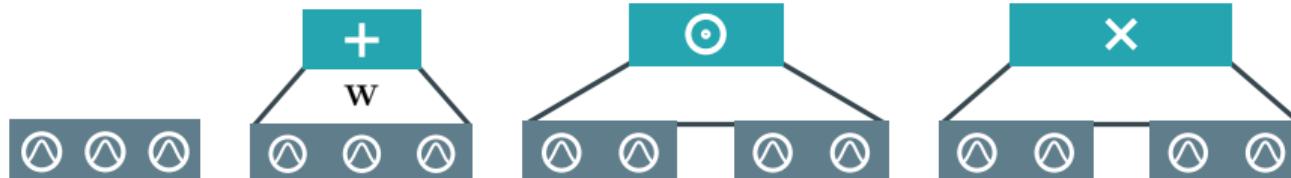


# **probabilistic circuits (PCs)**

*a tensorized definition*

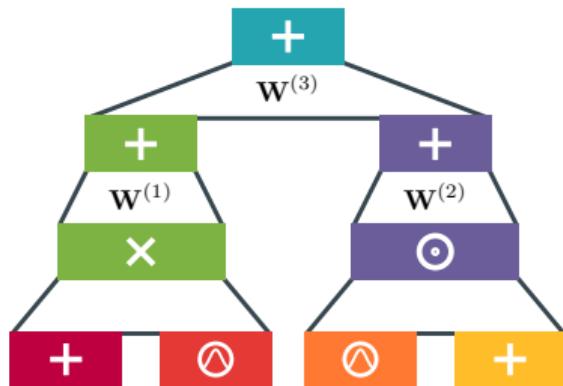
- I. A set of tractable functions is a circuit layer
- II. A linear projection of a layer is a circuit layer
- III. The product of two layers is a circuit layer

$$c(\mathbf{x}) = \text{vec}(\mathbf{l}(\mathbf{x})\mathbf{r}(\mathbf{x})^\top) \quad // \text{ Kronecker}$$



# **probabilistic circuits (PCs)**

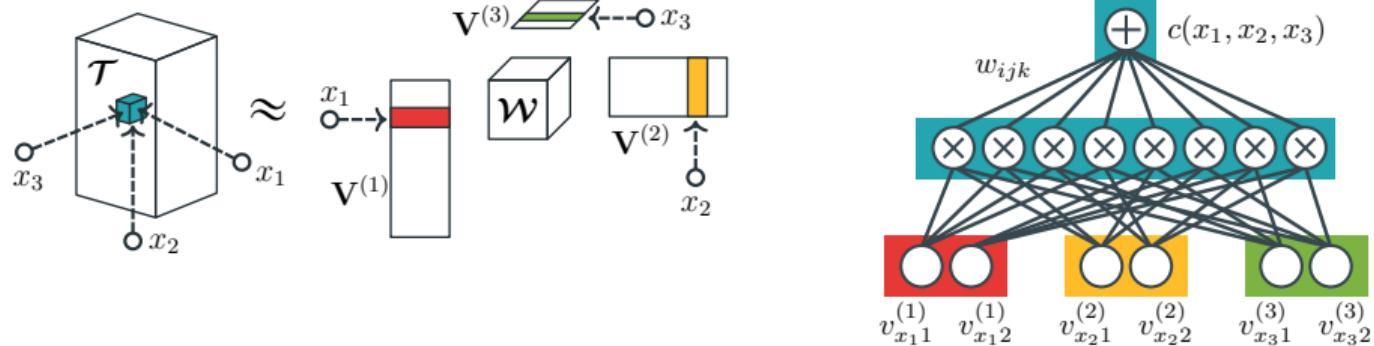
*a tensorized definition*



- I. A set of tractable functions is a circuit layer
  - II. A linear projection of a layer is a circuit layer
  - III. The product of two layers is a circuit layer
- stack layers to build a deep circuit!**

# *tensor factorizations*

*as circuits*



---

Loconte et al., "What is the Relationship between Tensor Factorizations and Circuits (and How Can We Exploit it)?", TMLR, 2025



***learning & reasoning with circuits in pytorch***

[github.com/april-tools/cirkit](https://github.com/april-tools/cirkit)

```
1 from cirkit.templates import circuit_templates  
2  
3 symbolic_circuit = circuit_templates.image_data(  
4     (1, 28, 28),                      # The shape of MNIST  
5     region_graph='quad-graph',  
6     input_layer='categorical',          # input distributions  
7     sum_product_layer='cp',            # CP, Tucker, CP-T  
8     num_input_units=64,                # overparameterizing  
9     num_sum_units=64,  
10    sum_weight_param=circuit_templates.Parameterization(  
11        activation='softmax',  
12        initialization='normal'  
13    )  
14 )
```

```
1 from cirkit.pipeline import compile
2 circuit = compile(symbolic_circuit)
3
4 with torch.no_grad():
5     test_lls = 0.0
6     for batch, _ in test_dataloader:
7         batch = batch.to(device).unsqueeze(dim=1)
8         log_likelihoods = circuit(batch)
9         test_lls += log_likelihoods.sum().item()
10 average_ll = test_lls / len(data_test)
11 bpd = -average_ll / (28 * 28 * np.log(2.0))
12 print(f"Average LL: {average_ll:.3f}") # Average LL:
13     → -682.916
14 print(f"Bits per dim: {bpds:.3f}") # Bits per dim: 1.257
```

The screenshot shows a Jupyter Notebook interface. At the top, there's a header bar with a file icon, a dropdown menu labeled "main", and the path "cirkit / notebooks / region-graphs-and-parametrisation.ipynb". To the right of the path is a search bar with the placeholder "Go to file" and a refresh icon. Below the header, a commit message from "loreloc" is displayed: "updated notebooks with respect to API changes" with a timestamp "e3e7e80 · 2 days ago" and a clock icon. Underneath the commit message, there are buttons for "Preview", "Code", and "Blame", followed by the statistics "1082 lines (1082 loc) · 793 KB". On the far right of the header, there are buttons for "Raw", "Copy", and "Download".

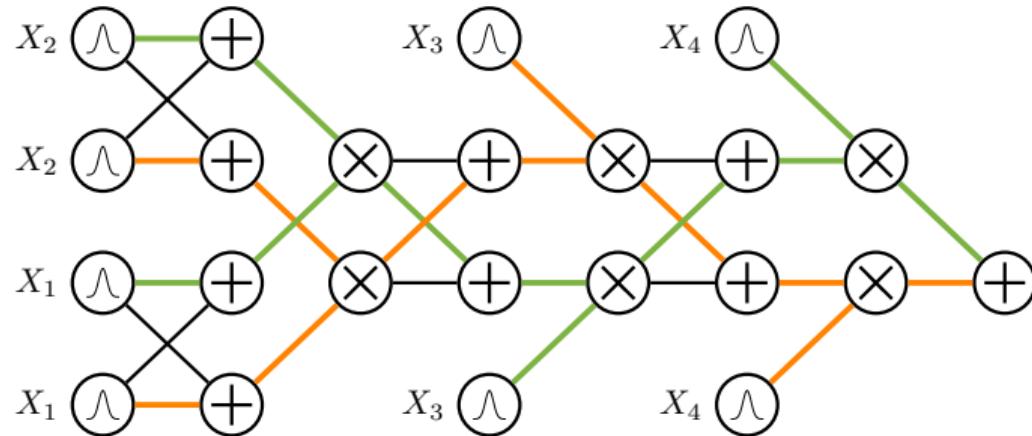
# Notebook on Region Graphs and Sum Product Layers

## Goals

By the end of this tutorial you will:

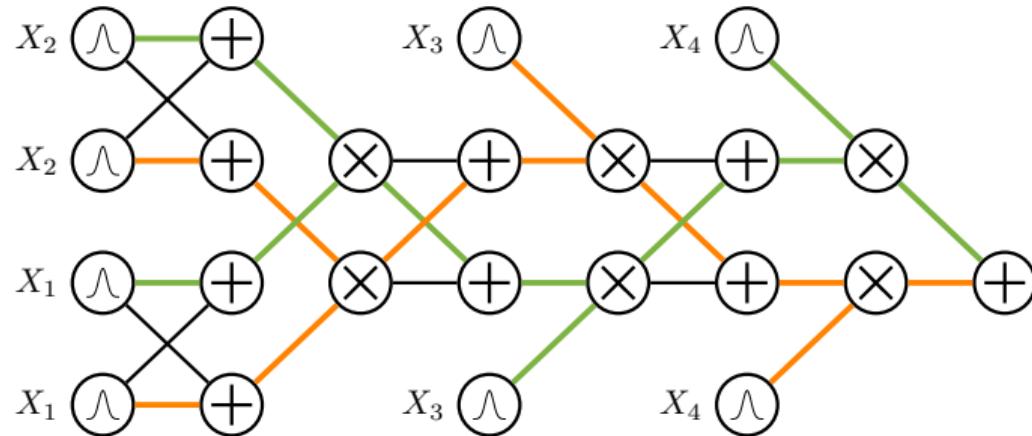
- know what a region graph is
- know how to choose between region graphs for your circuit
- understand how to parametrize a circuit by choosing a sum product layer
- build circuits to tractably estimate a probability distribution over images<sup>1</sup>

## deep mixtures



$$p(\mathbf{x}) = \sum_{\mathcal{T}} \left( \prod_{w_j \in \mathbf{w}_{\mathcal{T}}} w_j \right) \prod_{l \in \text{leaves}(\mathcal{T})} p_l(\mathbf{x})$$

## *deep mixtures*



*an exponential number of mixture components!*

# **...why PCs?**

## **1. A grammar for tractable models**

One formalism to represent many probabilistic models

⇒ #HMMs #Trees #XGBoost, Tensor Networks, ...

# **...why PCs?**

## **1. A grammar for tractable models**

One formalism to represent many probabilistic models

⇒ #HMMs #Trees #XGBoost, Tensor Networks, ...

## **2. Tractability == structural properties!!!**

Exact computations of reasoning tasks are certified by guaranteeing certain structural properties. #marginals #expectations #MAP, #product ...

# *structural properties*

*smoothness*

*decomposability*

*compatibility*

# *structural properties*

*property A*

*property B*

*property C*

# *structural properties*

**smoothness**

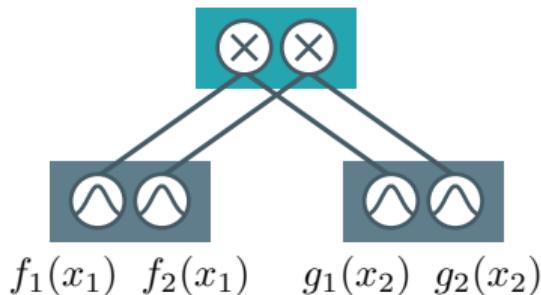
**decomposability**

**property C**

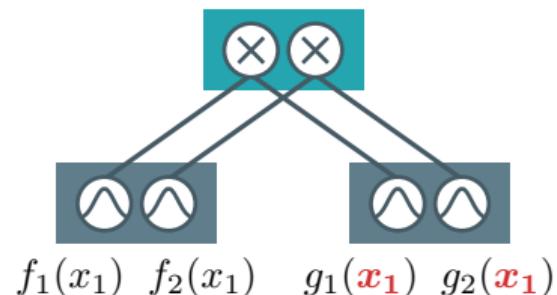
smoothness  $\wedge$  decomposability  
 $\implies$  multilinearity

# Multilinearity in circuits

the inputs of product units are defined over disjoint sets of variables



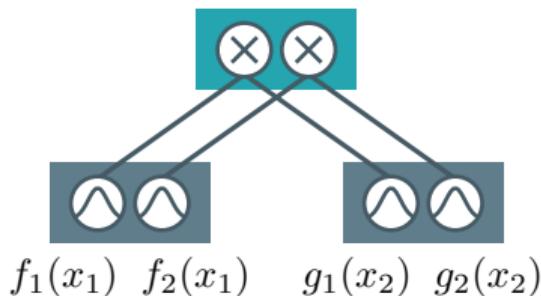
✓ **multilinear**



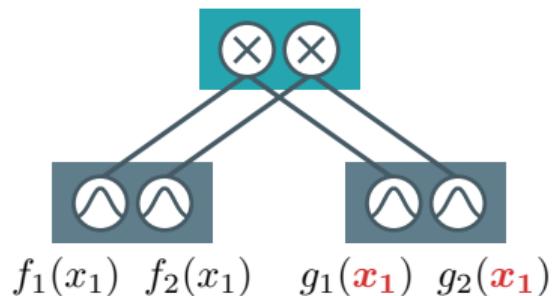
✗ **not multilinear**

# Multilinearity in circuits

the inputs of product units are defined over disjoint sets of variables



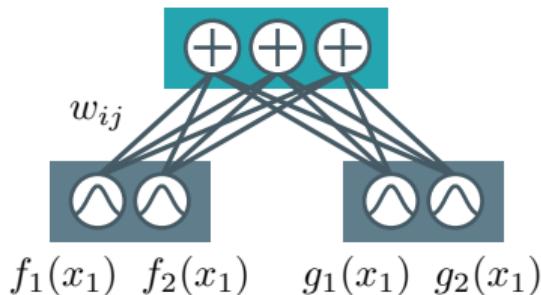
**decomposable circuit**



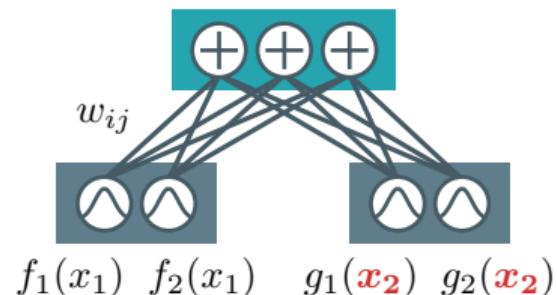
**non-decomposable circuit**

# Multilinearity in circuits

the inputs of sum units are defined over the same variables



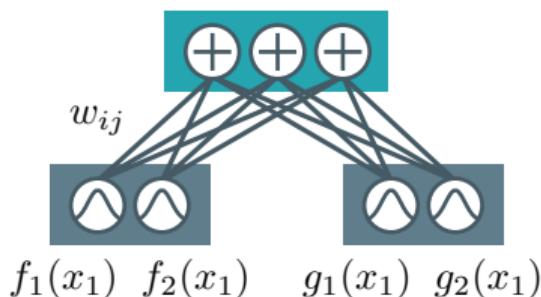
✓ **multilinear**



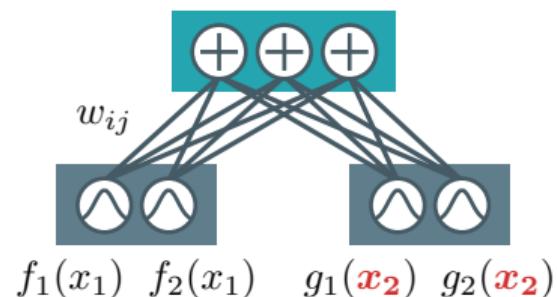
✗ **not multilinear**

# Multilinearity in circuits

the inputs of sum units are defined over the same variables



**smooth circuit**



**non-smooth circuit**

# *structural properties*

**smoothness**

**decomposability**

**property C**

smoothness  $\wedge$  decomposability  
 $\implies$  multilinearity

# *structural properties*

**smoothness**

tractable computation of arbitrary integrals  
in probabilistic circuits

**decomposability**

**property C**

$$p(\mathbf{y}) = \int p(\mathbf{y}, \mathbf{z}) d\mathbf{z}, \quad \forall \mathbf{Y} \subseteq \mathbf{X}, \quad \mathbf{Z} = \mathbf{X} \setminus \mathbf{Y}$$

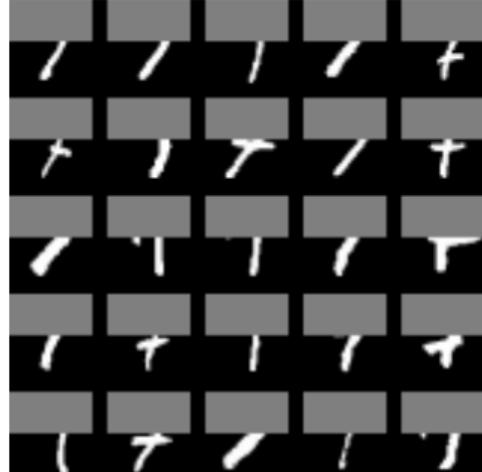
⇒ tractable partition function  
⇒ also any conditional is tractable

## *tractable marginals on PCs*

Original

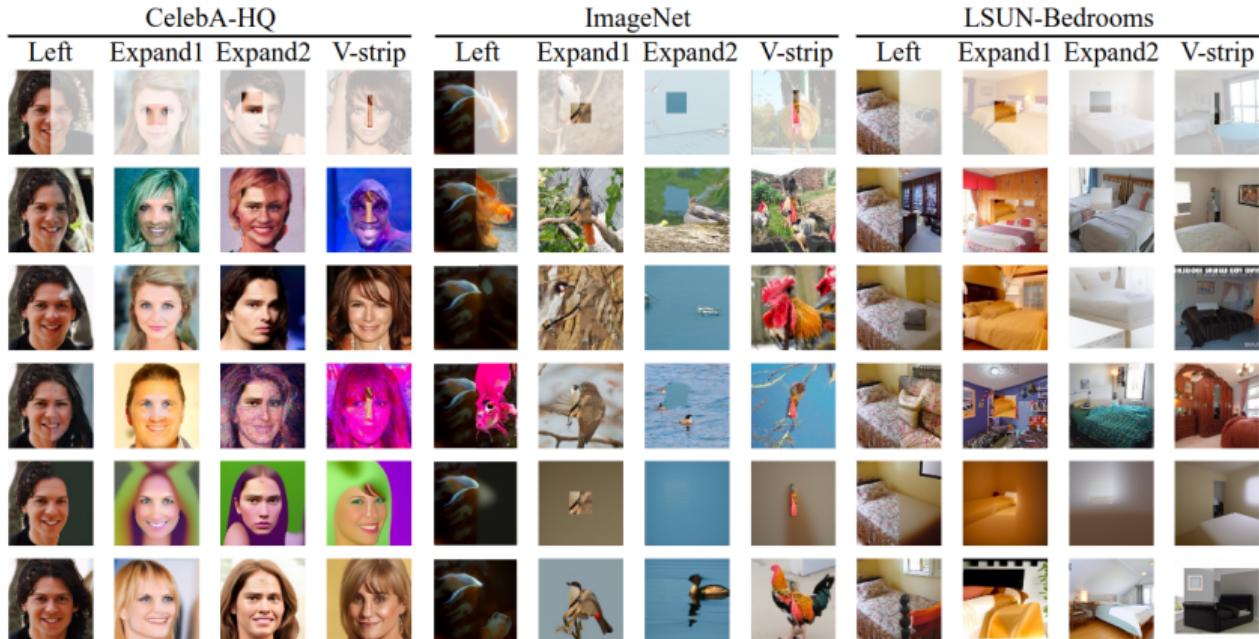


Missing



Conditional sample





# *structural properties*

***smoothness***

Integrals involving two or more functions:  
e.g., expectations

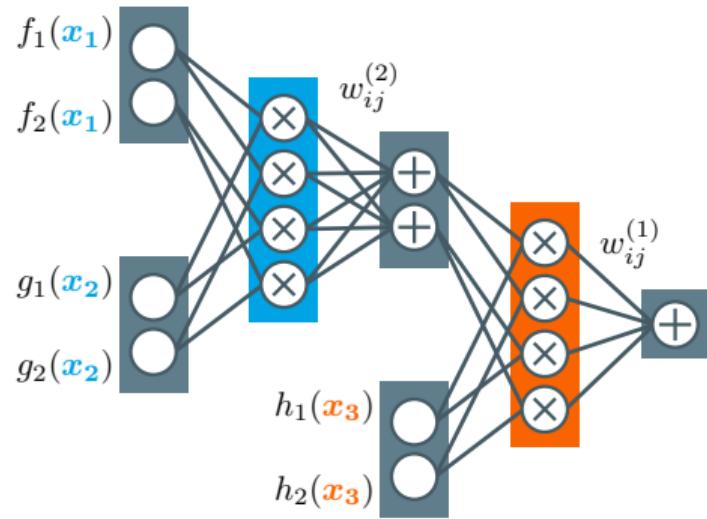
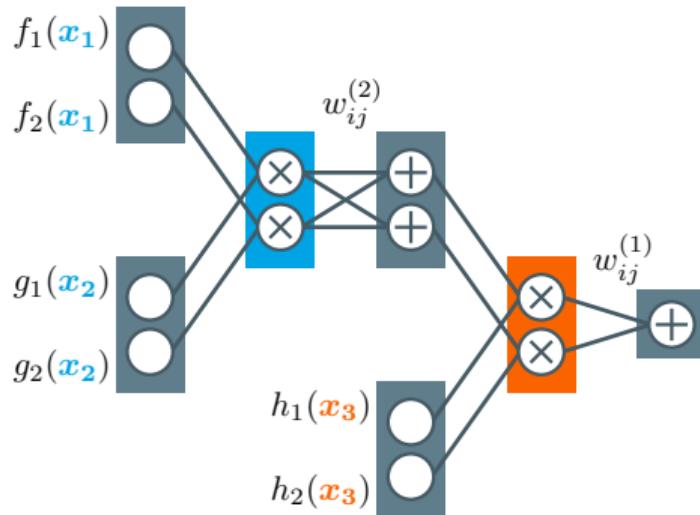
***decomposability***

$$\mathbb{E}_{\mathbf{x} \sim p} [f(\mathbf{x})] = \int p(\mathbf{x}) f(\mathbf{x}) d\mathbf{x}$$

***compatibility***

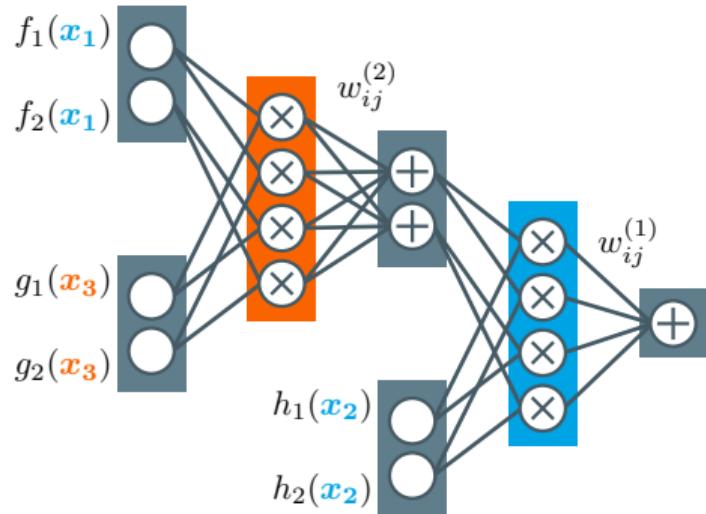
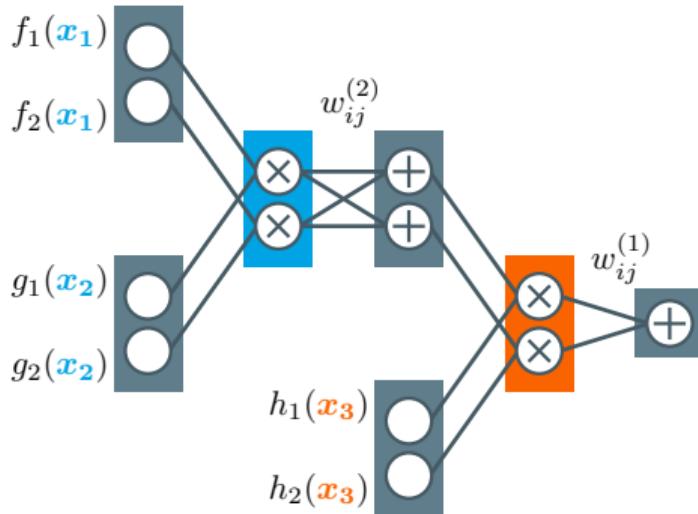
when both  $p(\mathbf{x})$  and  $f(\mathbf{x})$  are circuits

# *compatibility*



**compatible circuits**

# *compatibility*



**non-compatible circuits**

# *structural properties*

**smoothness**

**decomposability**

**compatibility**

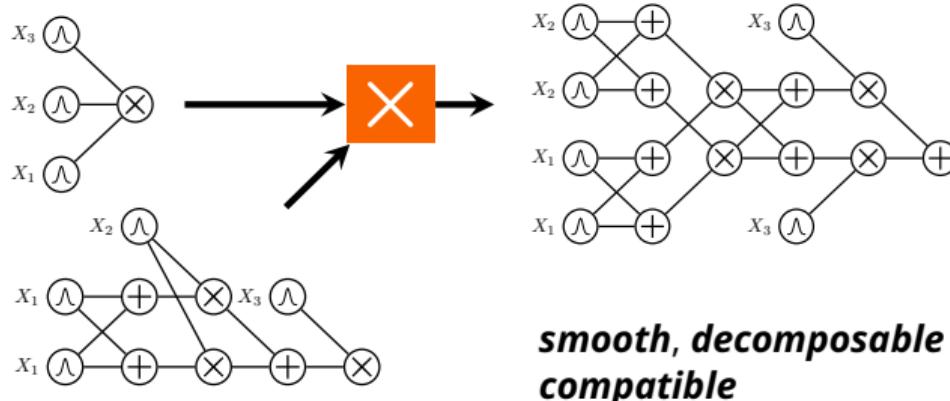
**compatibility**



**smoothness  $\wedge$  decomposability**

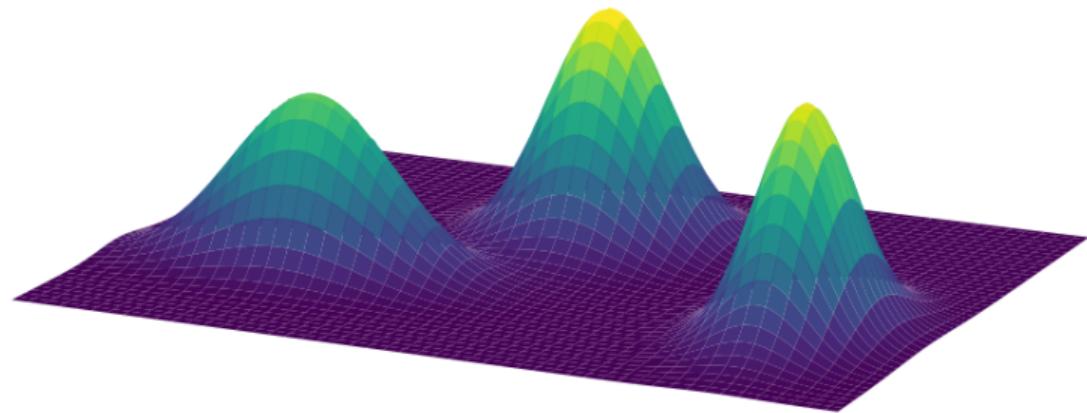
**compatibility**  $\Rightarrow$  tractable expectations

# Tractable products



**compute**  $\mathbb{E}_{\mathbf{x} \sim p} f(\mathbf{x}) = \int p(\mathbf{x}) f(\mathbf{x}) \mathrm{d}\mathbf{x}$  in  $O(|p| |f|)$

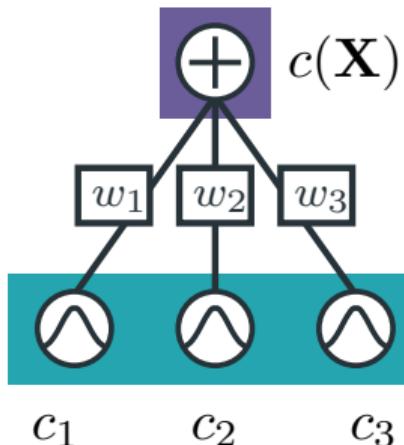
```
1 from cirkit.symbolic.circuit import Circuit
2 from cirkit.symbolic.functional import (
3     integrate, multiply)
4
5 # Circuits expectation  $\int [p(x) f(x)] dx$ 
6 def expectation(p: Circuit, f: Circuit) -> Circuit:
7     i = multiply(p, f)
8     return integrate(i)
9
10 # Squared loss  $\int [p(x)-q(x)]^2 dx = E_p[p] + E_q[q] - 2E_p[q]$ 
11 #           =  $\int p^2(x) dx + \int q^2(x) dx - 2\int p(x)q(x) dx$ 
12 def squared_loss(p: Circuit, q: Circuit) -> Circuit:
13     p2 = multiply(p, p)
14     q2 = multiply(q, q)
15     pq = multiply(p, q)
16     return integrate(p2) + integrate(q2) - 2 * integrate(pq)
```



$$c(\mathbf{X}) = \sum_{i=1}^K w_i c_i(\mathbf{X}), \quad \text{with } w_i \geq 0, \quad \sum_{i=1}^K w_i = 1$$

# **additive MMs**

*are so cool!*



easily represented as shallow PCs

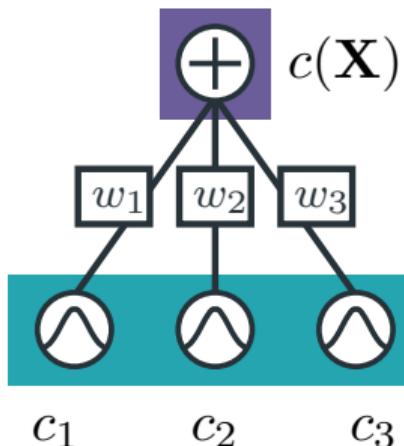
these are **monotonic** PCs

if marginals/conditionals are tractable for the components, they are tractable for the MM

universal approximators...

# **additive MMs**

*are so cool!*



easily represented as shallow PCs

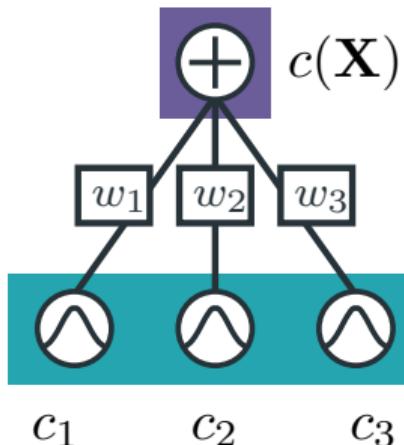
these are **monotonic** PCs

if marginals/conditionals are tractable for the components, they are tractable for the MM

universal approximators...

## **additive MMs**

*are so cool!*



easily represented as shallow PCs

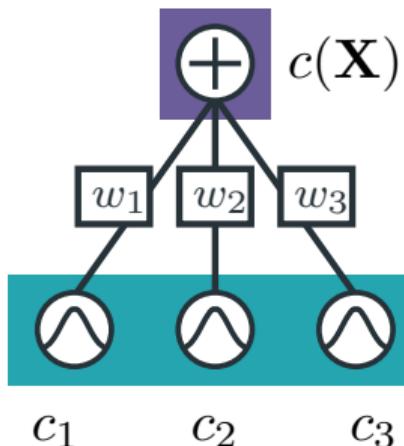
these are **monotonic** PCs

if marginals/conditionals are tractable for the components, they are tractable for the MM

universal approximators...

## **additive MMs**

are so cool!



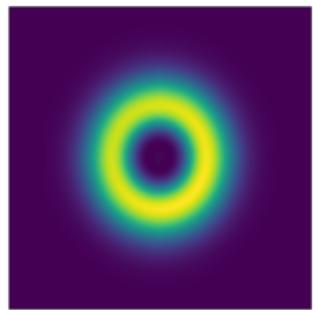
easily represented as shallow PCs

these are **monotonic** PCs

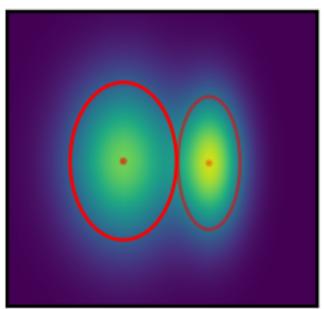
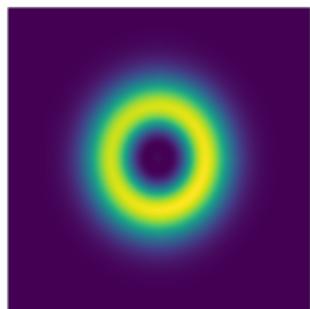
if marginals/conditionals are tractable for the components, they are tractable for the MM

universal approximators...

*however...*

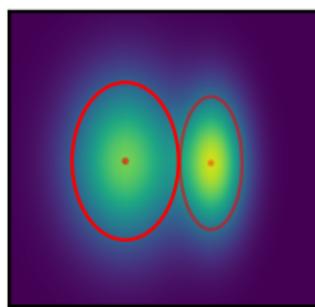
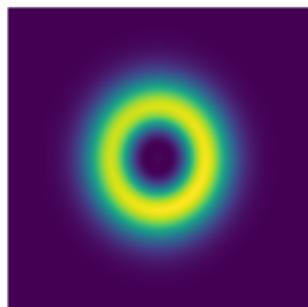


*however...*

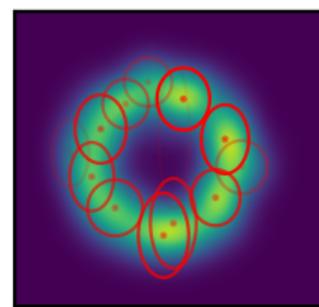


GMM ( $K = 2$ )

*however...*

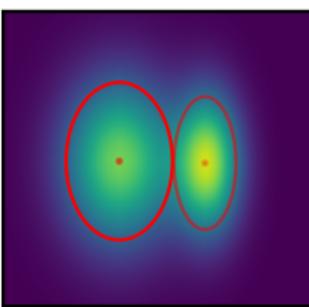
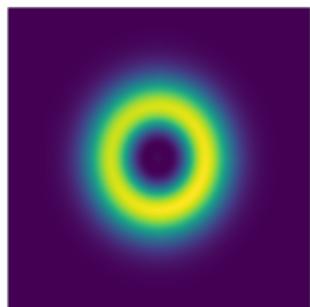


GMM ( $K = 2$ )

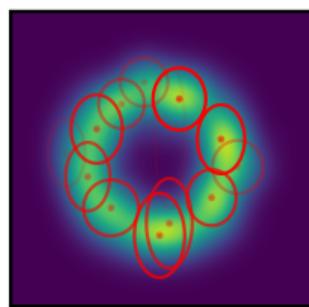


GMM ( $K = 16$ )

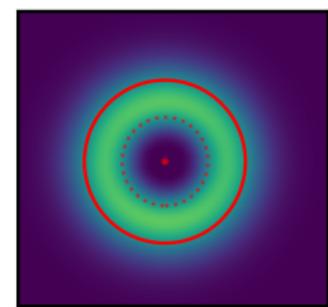
**however...**



GMM ( $K = 2$ )



GMM ( $K = 16$ )

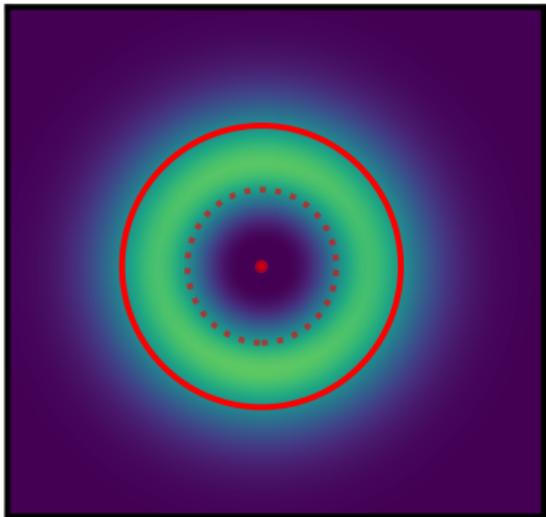


nGMM<sup>2</sup> ( $K = 2$ )

*spoiler*

**shallow mixtures  
with negative parameters  
can be *exponentially more compact* than  
deep ones with positive parameters.**

# **subtractive MMs**



also called negative/signed/**subtractive** MMs

⇒ or **non-monotonic** circuits,...

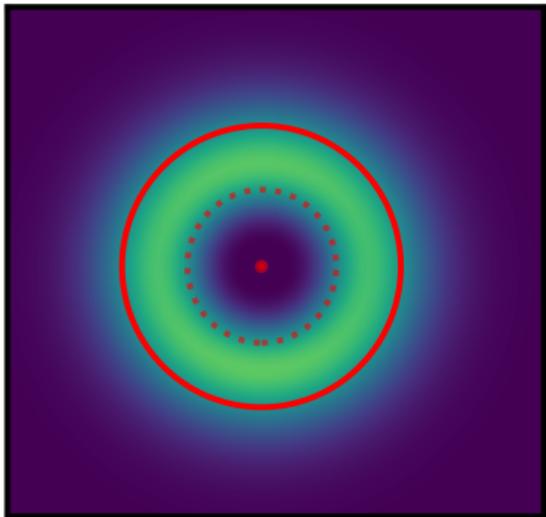
**issue:** how to preserve non-negative outputs?

well understood for simple parametric forms

e.g., Weibulls, Gaussians

⇒ constraints on variance, mean

# **subtractive MMs**



also called negative/signed/**subtractive** MMs

⇒ *or non-monotonic circuits,...*

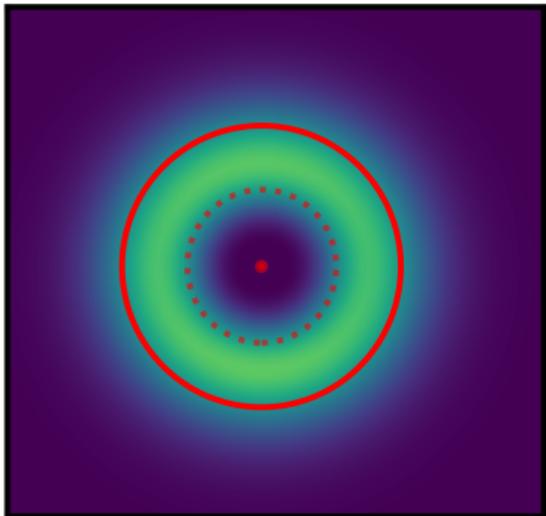
**issue:** how to preserve non-negative outputs?

well understood for simple parametric forms

e.g., Weibulls, Gaussians

⇒ *constraints on variance, mean*

# **subtractive MMs**



also called negative/signed/**subtractive** MMs

⇒ or **non-monotonic** circuits,...

**issue:** how to preserve non-negative outputs?

well understood for simple parametric forms

e.g., Weibulls, Gaussians

⇒ constraints on variance, mean

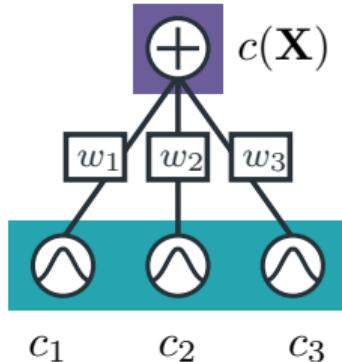
**tl;dr**

***“Understand when and how  
we can use negative parameters  
in deep subtractive mixture models”***

**tl;dr**

***"Understand when and how  
we can use negative parameters  
in deep **non-monotonic circuits**"***

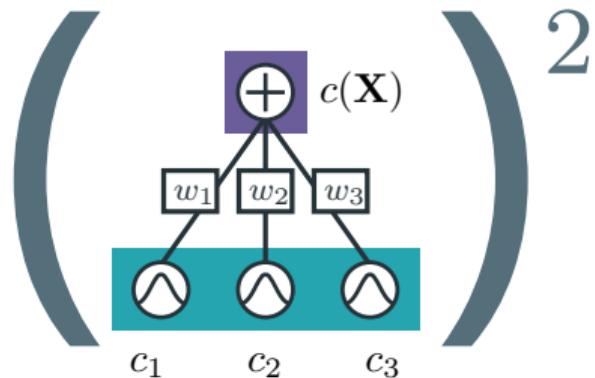
## *subtractive MMs as circuits*



a **non-monotonic** smooth and (structured)  
decomposable circuit  
⇒ possibly with negative outputs

$$c(\mathbf{X}) = \sum_{i=1}^K w_i c_i(\mathbf{X}), \quad w_i \in \mathbb{R},$$

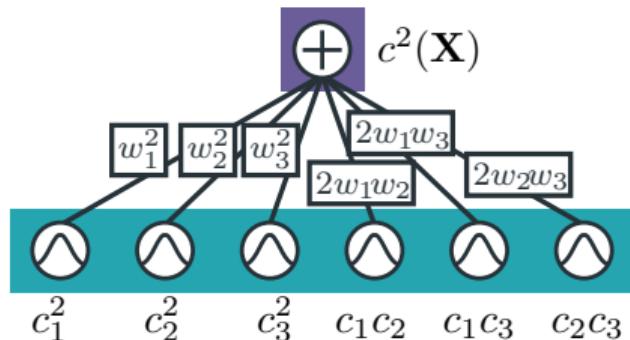
## *squaring shallow MMs*



$$c^2(\mathbf{X}) = \left( \sum_{i=1}^K w_i c_i(\mathbf{X}) \right)^2$$

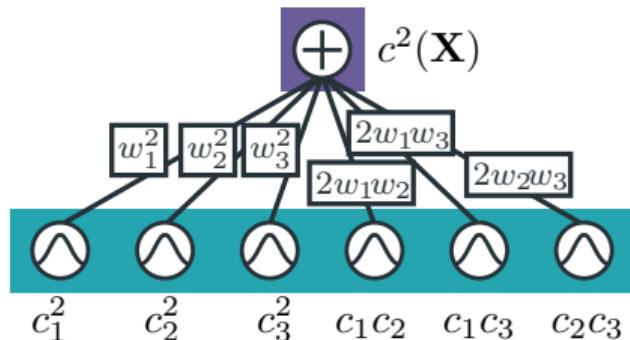
$\Rightarrow$  ensure non-negative output

# *squaring shallow MMs*



$$\begin{aligned} c^2(\mathbf{X}) &= \left( \sum_{i=1}^K w_i c_i(\mathbf{X}) \right)^2 \\ &= \sum_{i=1}^K \sum_{j=1}^K w_i w_j c_i(\mathbf{X}) c_j(\mathbf{X}) \end{aligned}$$

## *squaring shallow MMs*

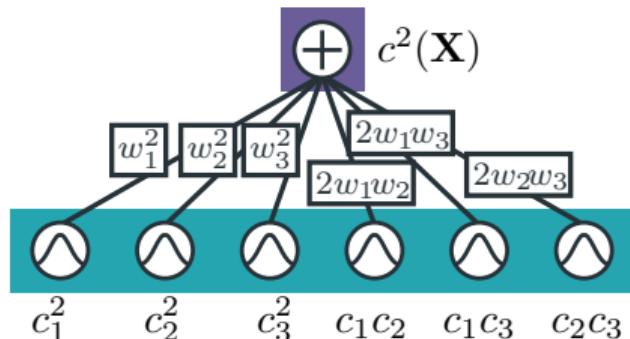


$$\begin{aligned} c^2(\mathbf{X}) &= \left( \sum_{i=1}^K w_i c_i(\mathbf{X}) \right)^2 \\ &= \sum_{i=1}^K \sum_{j=1}^K w_i w_j c_i(\mathbf{X}) c_j(\mathbf{X}) \end{aligned}$$

still a smooth and (str) decomposable PC with  $\mathcal{O}(K^2)$  components!

$\Rightarrow$  but still  $\mathcal{O}(K)$  parameters

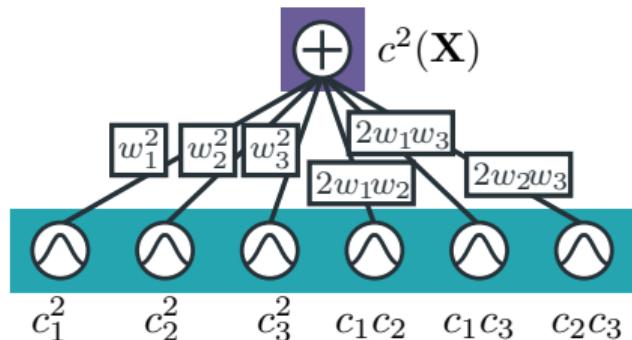
# *squaring shallow MMs*



$$\begin{aligned} c^2(\mathbf{X}) &= \left( \sum_{i=1}^K w_i c_i(\mathbf{X}) \right)^2 \\ &= \sum_{i=1}^K \sum_{j=1}^K w_i w_j c_i(\mathbf{X}) c_j(\mathbf{X}) \end{aligned}$$

to **renormalize**, we have to compute  $\sum_i \sum_j w_i w_j \int c_i(\mathbf{x}) c_j(\mathbf{x}) d\mathbf{x}$

# *squaring shallow MMs*



$$\begin{aligned} c^2(\mathbf{X}) &= \left( \sum_{i=1}^K w_i c_i(\mathbf{X}) \right)^2 \\ &= \sum_{i=1}^K \sum_{j=1}^K w_i w_j c_i(\mathbf{X}) c_j(\mathbf{X}) \end{aligned}$$

to **renormalize**, we have to compute  $\sum_i \sum_j w_i w_j \int c_i(\mathbf{x}) c_j(\mathbf{x}) d\mathbf{x}$   
⇒ or we pick  $c_i, c_j$  to be **orthonormal**...!

---

## EigenVI: score-based variational inference with orthogonal function expansions

---

**Diana Cai**  
Flatiron Institute  
[dcai@flatironinstitute.org](mailto:dcai@flatironinstitute.org)

**Chirag Modi**  
Flatiron Institute  
[cmodi@flatironinstitute.org](mailto:cmodi@flatironinstitute.org)

**Charles C. Margossian**  
Flatiron Institute  
[cmargossian@flatironinstitute.org](mailto:cmargossian@flatironinstitute.org)

**Robert M. Gower**  
Flatiron Institute  
[rgower@flatironinstitute.org](mailto:rgower@flatironinstitute.org)

**David M. Blei**  
Columbia University  
[david.blei@columbia.edu](mailto:david.blei@columbia.edu)

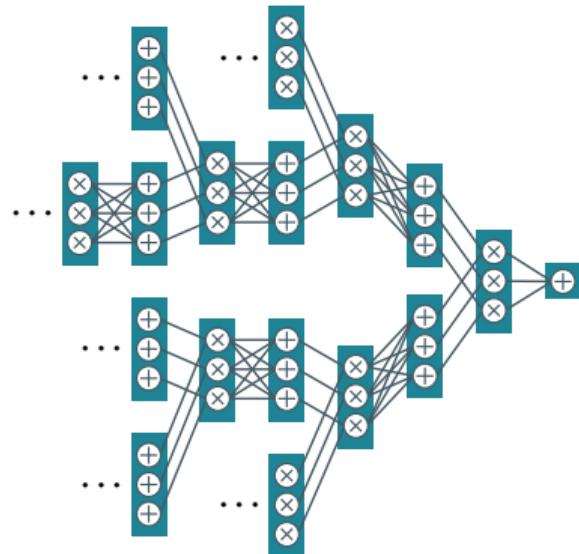
**Lawrence K. Saul**  
Flatiron Institute  
[lsaul@flatironinstitute.org](mailto:lsaul@flatironinstitute.org)

*orthonormal squared mixtures for VI*

*wait...*

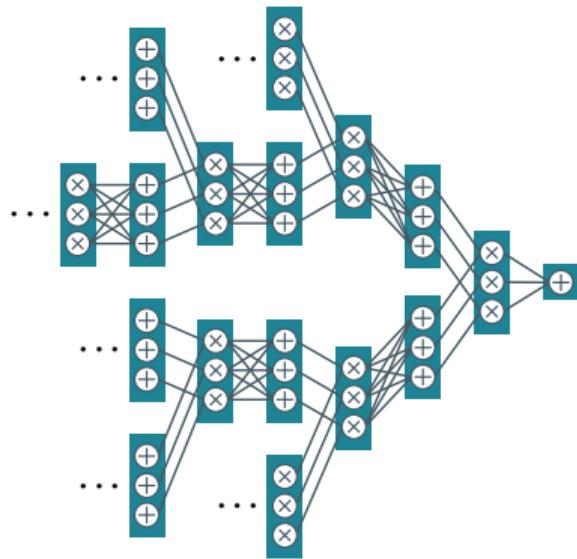
*“do negative parameters  
really boost expressiveness?  
and...always?”*

# theorem



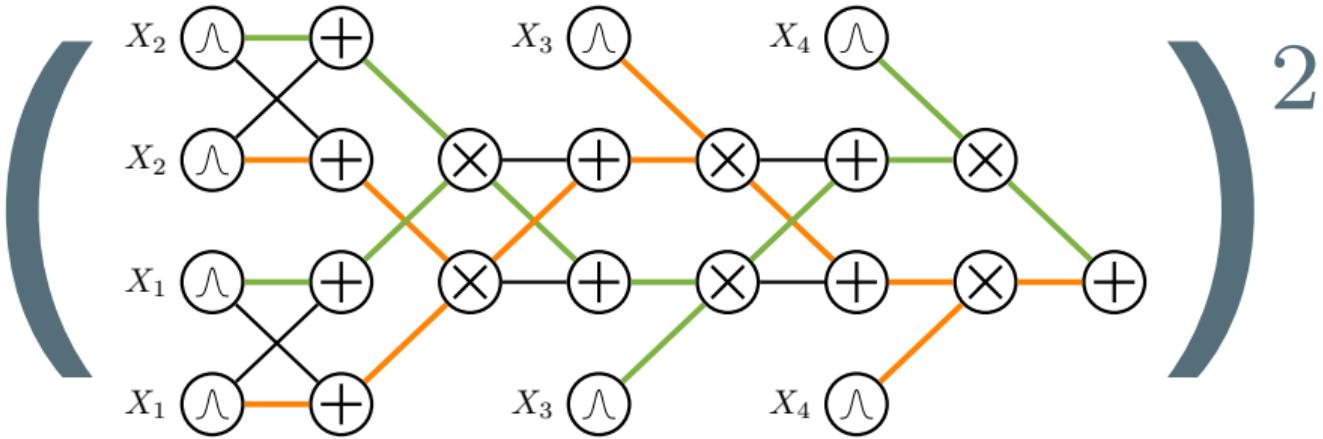
$\exists p$  requiring exponentially large  
monotonic circuits...

# theorem



$$\left( \begin{array}{c} \text{...} \\ \text{...} \\ \text{...} \end{array} \right)^2$$

...but compact  
squared non-monotonic circuits

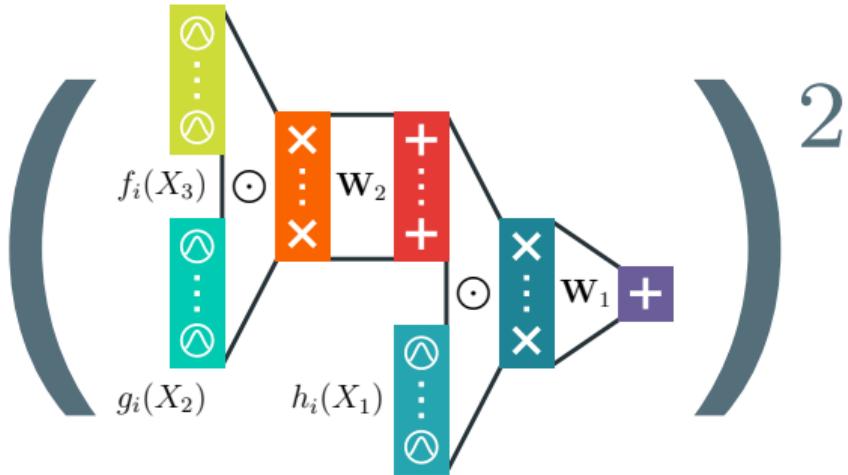


**how to efficiently square (and *renormalize*) a deep PC?**

# *compositional inference I*



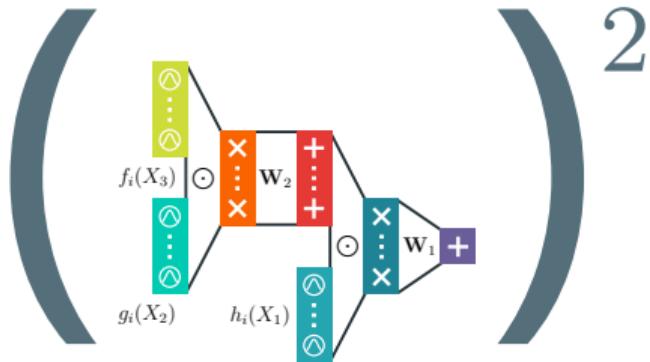
```
1 from cirkit.symbolic.functional import integrate, multiply  
2  
3 #  
4 # create a deep circuit  
5 c = build_symbolic_circuit('quad-tree-4')  
6  
7 #  
8 # compute the partition function of c^2  
9 def renormalize(c):  
10    c2 = multiply(c, c)  
11    return integrate(c2)
```



**how to efficiently square (and *renormalize*) a deep PC?**

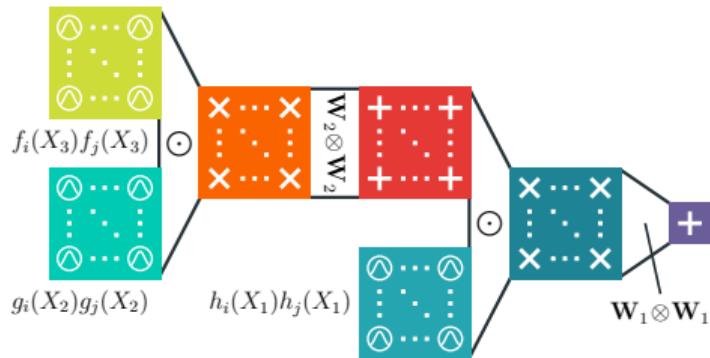
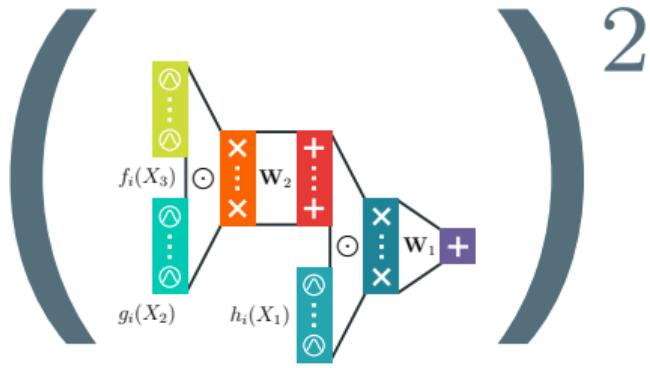
# *squaring deep PCs*

*the tensorized way*



# *squaring deep PCs*

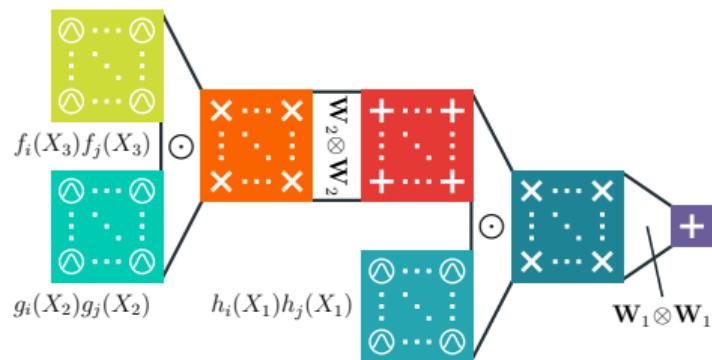
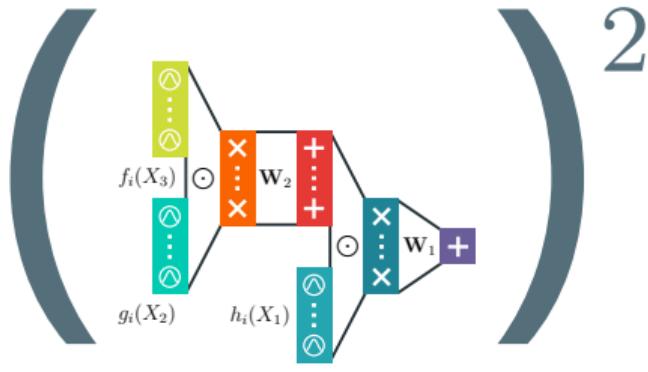
*the tensorized way*



**squaring a circuit = squaring layers**

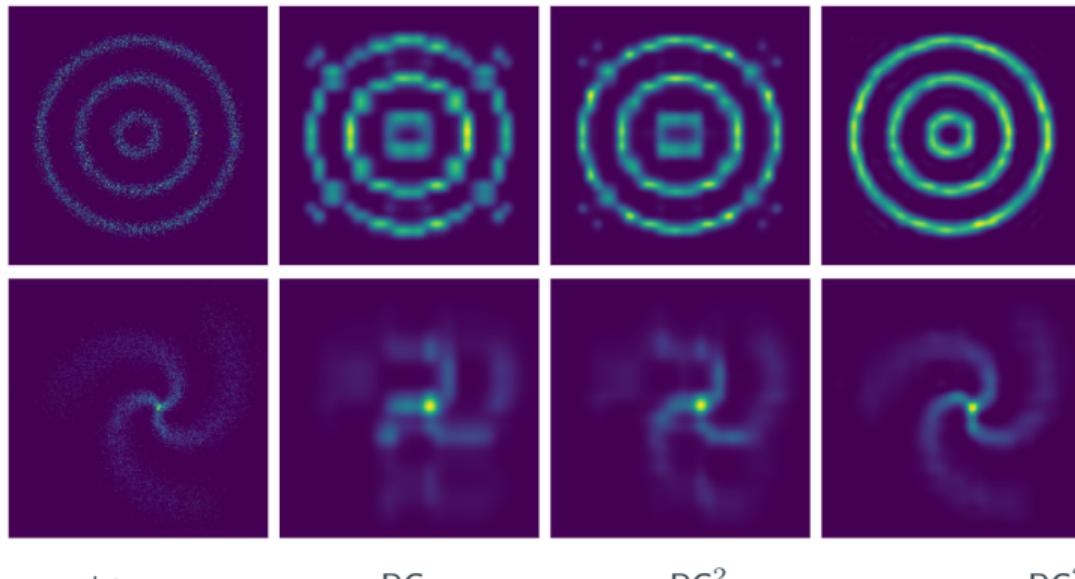
# *squaring deep PCs*

*the tensorized way*

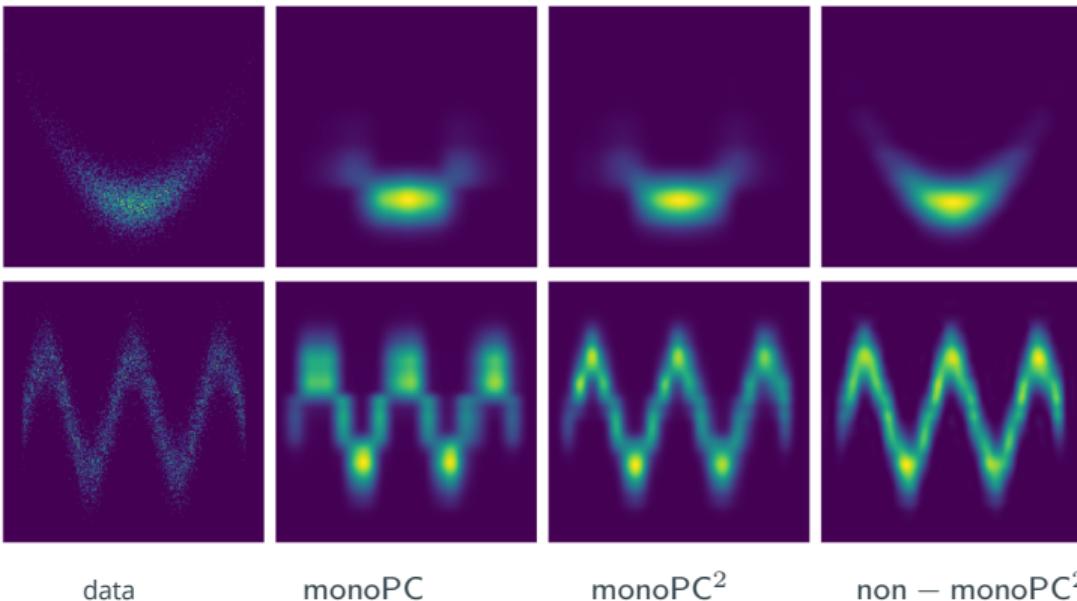


**exactly compute**  $\int c(x)c(x)dX$  **in time**  $O(LK^2)$

# *more expressive?*

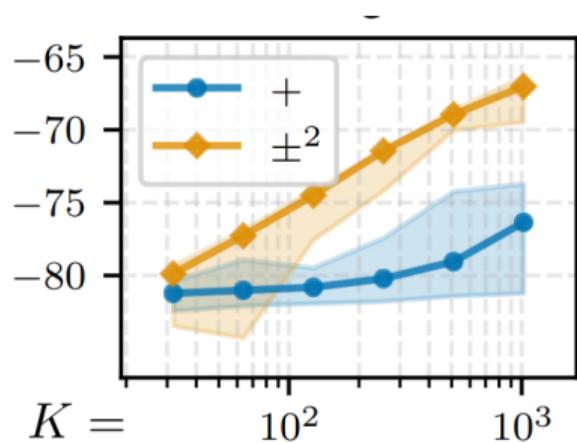
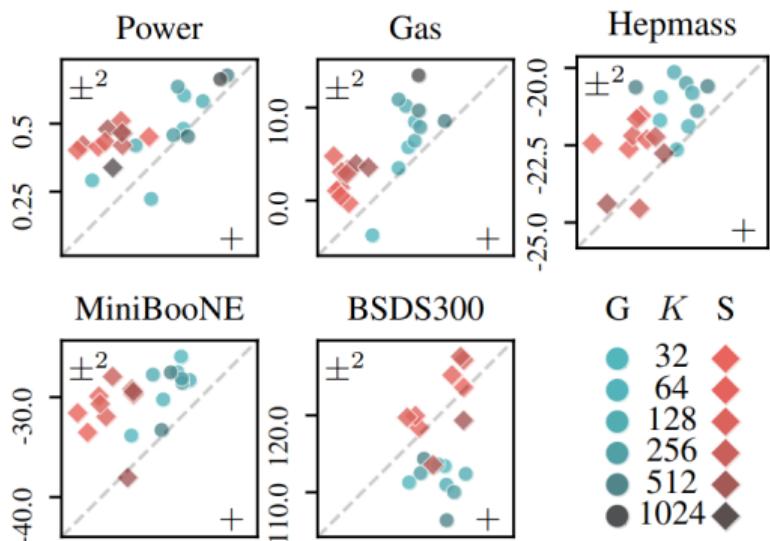


# *more expressive?*



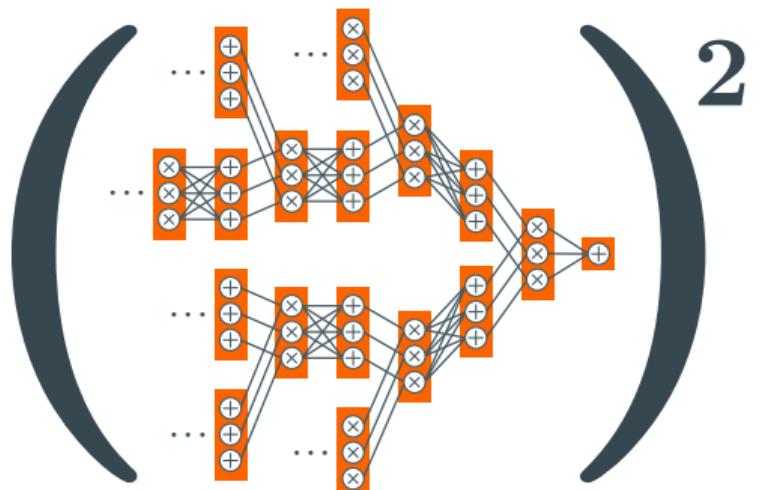
# how more expressive?

real-world data



# *theorem*

$\exists p$  requiring exponentially large  
squared non-mono circuits...

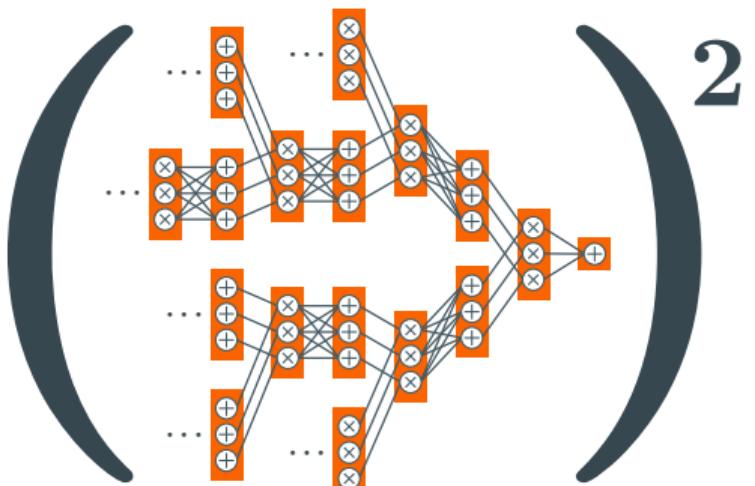


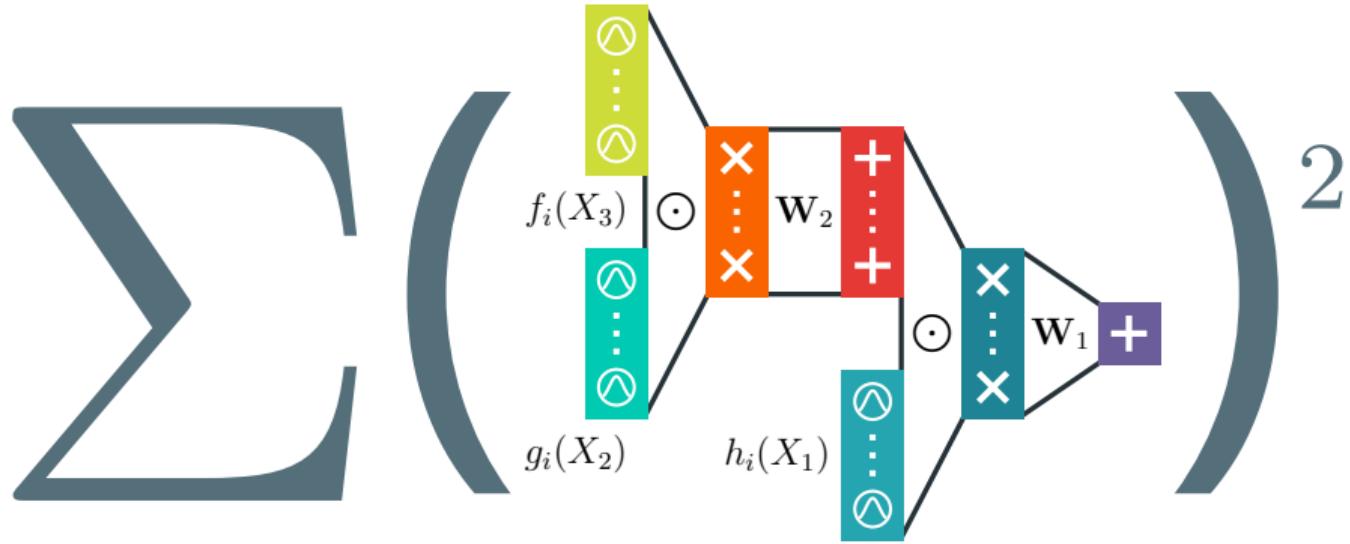
# *theorem*



**...but compact**

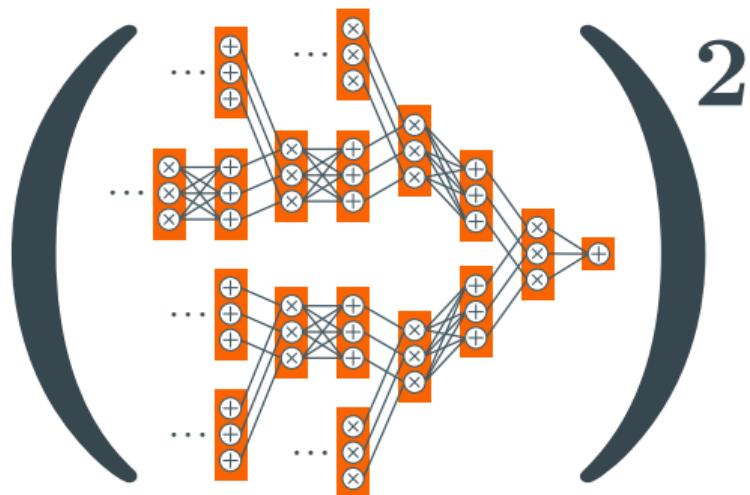
**monotonic circuits...!**





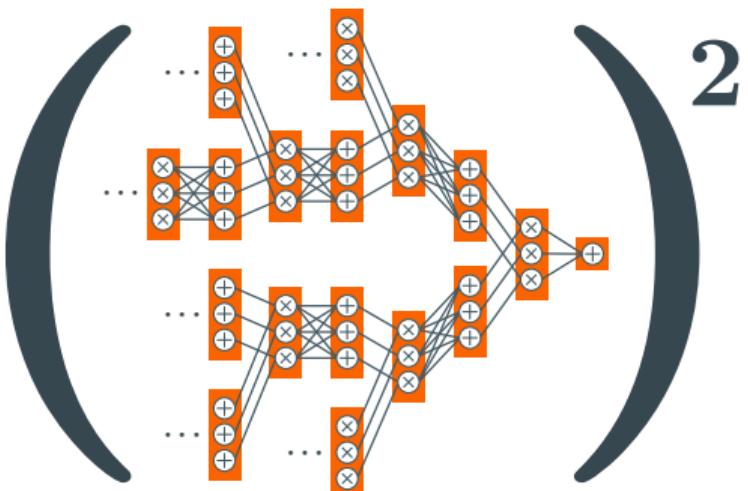
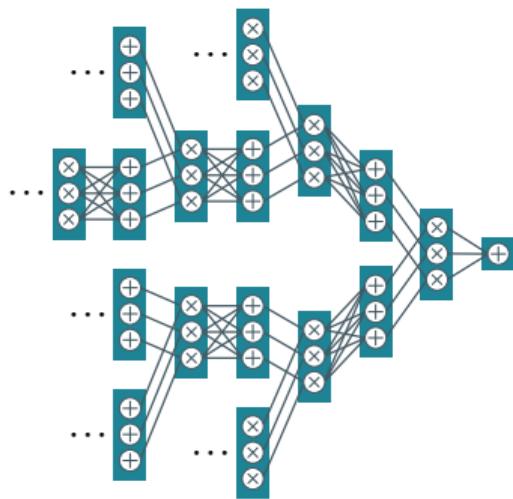
***what if we use more than one square?***

## **theorem**



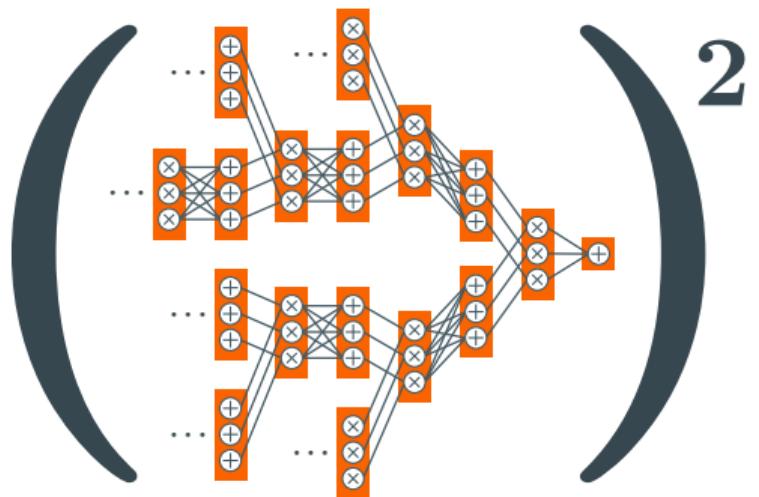
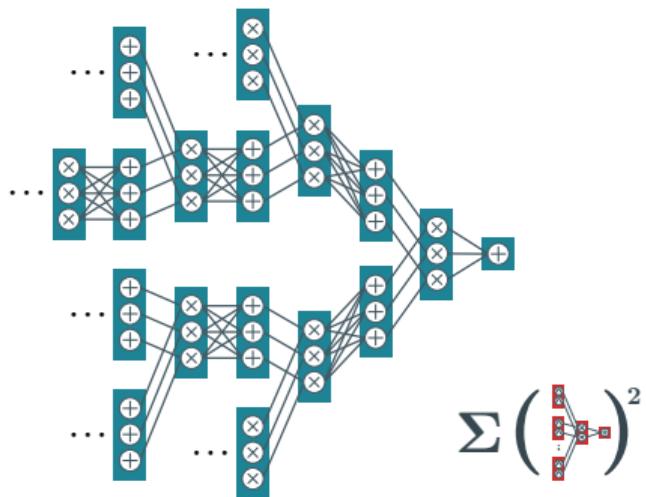
$\exists p$  requiring exponentially large **squared non-mono circuits**...

# *theorem*



**...exponentially large monotonic circuits...**

# theorem



...but compact **SOS circuits...**!

$$\pm_{\text{sd}} = \Delta \Sigma_{\text{cmp}}^2$$

(Theorem 5)

$$\Sigma_{\text{cmp}}^2 = \text{psd}$$

(Proposition 2)

$$+_{\text{sd}}$$

•  
Open Question 1

•  
Open Question 2

$$\pm_{\mathbb{R}}^2$$

• UDISJ  
(Theorem 0)

• UPS  
(Theorem 2)

• UTQ  
(Theorem B.3)

## *a hierarchy of subtractive mixtures*

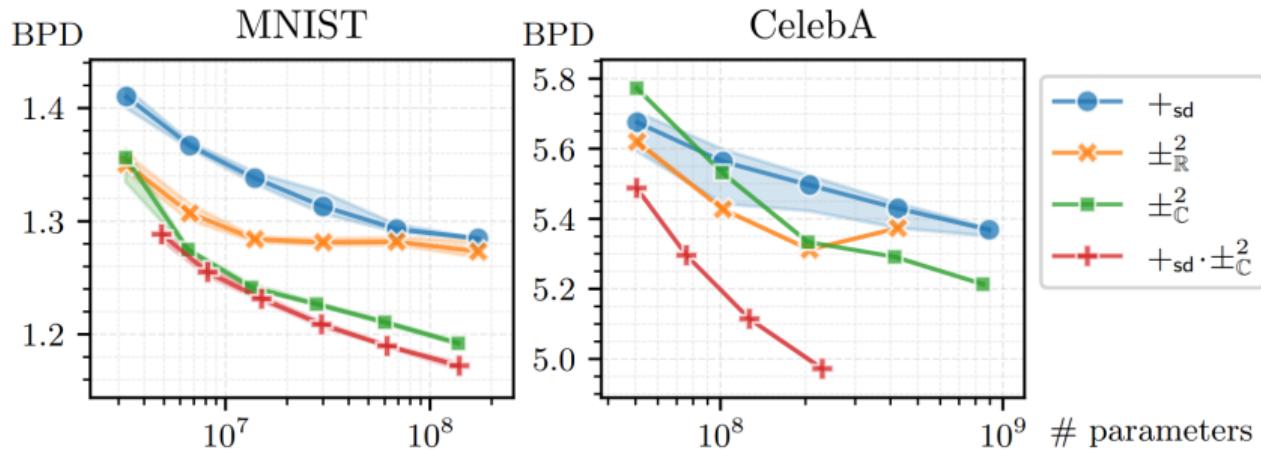
we can define circuits (and hence mixtures) over the Complex:

$$c^2(\mathbf{x}) = c(\mathbf{x})^\dagger c(\mathbf{x}), \quad c(\mathbf{x}) \in \mathbb{C}$$

and then we can note that they can be written as a SOS form

$$c^2(\mathbf{x}) = r(\mathbf{x})^2 + i(\mathbf{x})^2, \quad r(\mathbf{x}), i(\mathbf{x}) \in \mathbb{R}$$

***complex circuits are SOS (and scale better!)***



***complex circuits are SOS (and scale better!)***

*takeaway*

*“use squared mixtures  
over complex numbers  
and you get a SOS for free”*

## *takeaway*

*“use squared mixtures  
over complex numbers  
and you get a SOS for free”*

⇒    *but how to implement them?*

# *compositional inference I*



```
1 from cirkit.symbolic.functional import integrate, multiply,
2     conjugate
3
4 # create a deep circuit with complex parameters
5 c = build_symbolic_complex_circuit('quad-tree-4')
6
7 # compute the partition function of c^2
8 def renormalize(c):
9     c1 = conjugate(c)
10    c2 = multiply(c, c1)
11    return integrate(c2)
```

# On Faster Marginalization with Squared Circuits via Orthonormalization

Lorenzo Loconte<sup>1</sup>      Antonio Vergari<sup>1</sup>

<sup>1</sup> School of Informatics, University of Edinburgh, UK

l.loconte@sms.ed.ac.uk, avergari@ed.ac.uk

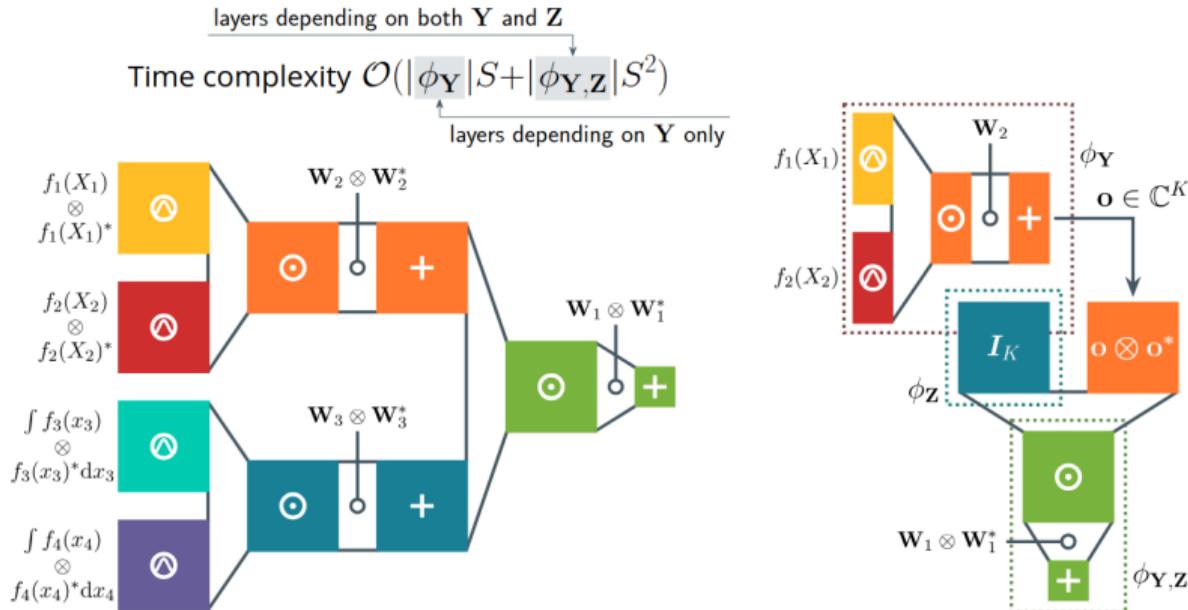
***what about deep orthonormal mixtures  
and arbitrary marginals?***

# On Faster Marginalization with Squared Circuits via Orthonormalization

Lorenzo Loconte<sup>1</sup>      Antonio Vergari<sup>1</sup>

<sup>1</sup> School of Informatics, University of Edinburgh, UK  
l.loconte@sms.ed.ac.uk, avergari@ed.ac.uk

*it suffices to orthonormalize each layer!*



***faster marginalization of arbitrary subsets of features***

# **approximate inference**

e.g., via sampling

Can we use a subtractive mixture model to approximate expectations?

$$\mathbb{E}_{\mathbf{x} \sim q(\mathbf{x})} [f(\mathbf{x})] \approx \frac{1}{S} \sum_{i=1}^S f(\mathbf{x}^{(i)}) \quad \text{with} \quad \mathbf{x}^{(i)} \sim q(\mathbf{x})$$

$\Rightarrow$  but how to sample from  $q$ ?

# approximate inference

e.g., via sampling

Can we use a subtractive mixture model to approximate expectations?

$$\mathbb{E}_{\mathbf{x} \sim q(\mathbf{x})} [f(\mathbf{x})] \approx \frac{1}{S} \sum_{i=1}^S f(\mathbf{x}^{(i)}) \quad \text{with} \quad \mathbf{x}^{(i)} \sim q(\mathbf{x})$$

$\Rightarrow$  but how to sample from  $q$ ?

use **autoregressive inverse transform sampling**:

$$x_1 \sim q(x_1), \quad x_i \sim q(x_i | \mathbf{x}_{<i}) \quad \text{for } i \in \{2, \dots, d\}$$

$\Rightarrow$  can be slow for large dimensions, requires **inverting the CDF**

# **approximate inference**

*difference of expectation estimator*

**Idea:** represent  $q$  as a difference of two additive mixtures

$$q(\mathbf{x}) = Z_+ \cdot q_+(\mathbf{x}) - Z_- \cdot q_-(\mathbf{x})$$

⇒ *expectations will break down in two “parts”*

# **approximate inference**

*difference of expectation estimator*

**Idea:** represent  $q$  as a difference of two additive mixtures

$$q(\mathbf{x}) = Z_+ \cdot q_+(\mathbf{x}) - Z_- \cdot q_-(\mathbf{x})$$

$\Rightarrow$  *expectations will break down in two “parts”*

$$\frac{Z_+}{S_+} \sum_{s=1}^{S_+} f(\mathbf{x}_+^{(s)}) - \frac{Z_-}{S_-} \sum_{s=1}^{S_-} f(\mathbf{x}_-^{(s)}), \text{ where } \begin{aligned} \mathbf{x}_+^{(s)} &\sim q_+(\mathbf{x}_+) \\ \mathbf{x}_-^{(s)} &\sim q_-(\mathbf{x}_-) \end{aligned}, \quad (1)$$

# *approximate inference*

*difference of expectation estimator*

Method	$d$	Number of components ( $K$ )					
		2		4		6	
		$\log( \hat{I} - I )$	Time (s)	$\log( \hat{I} - I )$	Time (s)	$\log( \hat{I} - I )$	Time (s)
$\Delta$ ExS	16	-19.507 $\pm$ 1.025	0.293 $\pm$ 0.004	-19.062 $\pm$ 0.823	1.049 $\pm$ 0.077	-19.497 $\pm$ 1.974	2.302 $\pm$ 0.159
ARITS	16	-19.111 $\pm$ 1.103	7.525 $\pm$ 0.038	-19.299 $\pm$ 1.611	7.52 $\pm$ 0.023	-18.739 $\pm$ 1.024	7.746 $\pm$ 0.032
$\Delta$ ExS	32	-48.411 $\pm$ 1.265	0.325 $\pm$ 0.012	-48.046 $\pm$ 0.972	1.027 $\pm$ 0.107	-48.34 $\pm$ 0.814	2.213 $\pm$ 0.177
ARITS	32	-47.897 $\pm$ 1.165	15.196 $\pm$ 0.059	-47.349 $\pm$ 0.839	15.535 $\pm$ 0.059	-47.3 $\pm$ 0.978	17.371 $\pm$ 0.06
$\Delta$ ExS	64	-108.095 $\pm$ 1.094	0.38 $\pm$ 0.034	-107.56 $\pm$ 0.616	0.9 $\pm$ 0.14	-107.653 $\pm$ 0.945	1.512 $\pm$ 0.383
ARITS	64	-107.898 $\pm$ 1.129	30.459 $\pm$ 0.098	-107.33 $\pm$ 0.929	33.892 $\pm$ 0.119	-107.374 $\pm$ 1.138	52.02 $\pm$ 0.127

*faster than autoregressive sampling*

# **...why PCs?**

## **1. A grammar for tractable models**

One formalism to represent many probabilistic models

⇒ #HMMs #Trees #XGBoost, Tensor Networks, ...

## **2. Tractability == structural properties!!!**

Exact computations of reasoning tasks are certified by guaranteeing certain structural properties. #marginals #expectations #MAP, #product ...

# **...why PCs?**

## **1. A grammar for tractable models**

One formalism to represent many probabilistic models

→ #HMMs #Trees #XGBoost, Tensor Networks, ...

## **2. Tractability == structural properties!!!**

Exact computations of reasoning tasks are certified by guaranteeing certain structural properties. #marginals #expectations #MAP, #product ...

## **3. Reliable neuro-symbolic AI**

logical constraints as circuits, multiplied to probabilistic circuits

---

# Semantic Probabilistic Layers for Neuro-Symbolic Learning

---

**Kareem Ahmed**

CS Department  
UCLA

ahmedk@cs.ucla.edu

**Stefano Teso**

CIMeC and DISI  
University of Trento

stefano.teso@unitn.it

**Kai-Wei Chang**

CS Department  
UCLA

kwchang@cs.ucla.edu

**Guy Van den Broeck**

CS Department  
UCLA

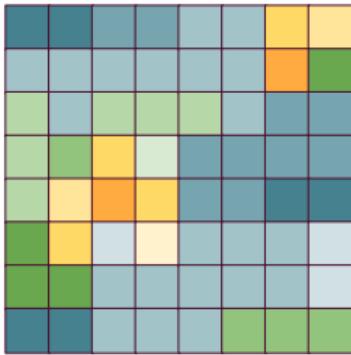
guyvdb@cs.ucla.edu

**Antonio Vergari**

School of Informatics  
University of Edinburgh

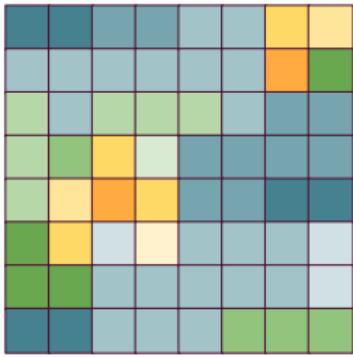
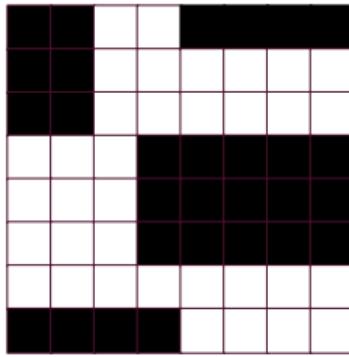
avergari@ed.ac.uk

***enforce constraints in neural networks at NeurIPS 2022***

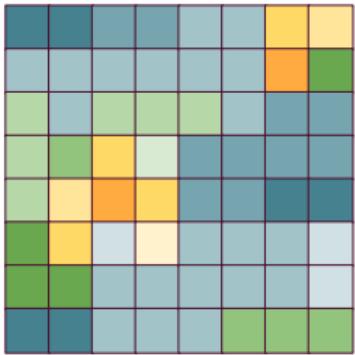
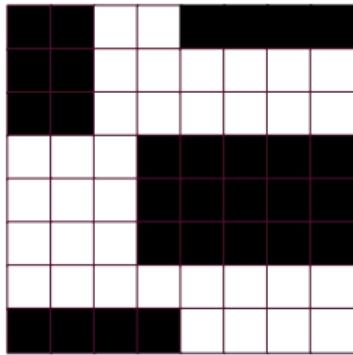
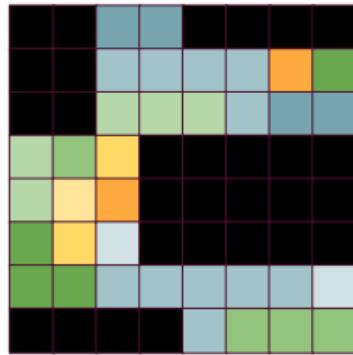


$$q(\mathbf{x})$$

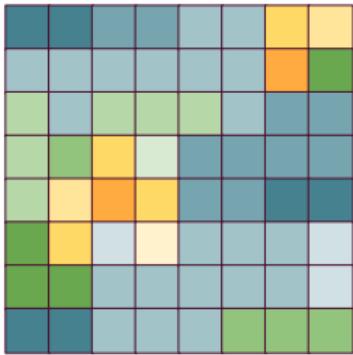
**start from a distribution  $q(\mathbf{x})$ ...**

 $q(\mathbf{x})$  $c(\mathbf{x})$ 

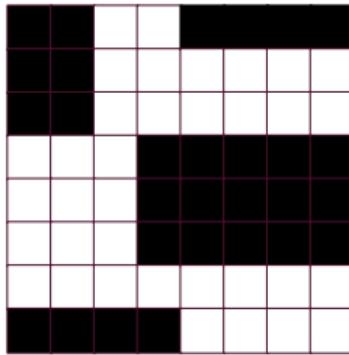
**...and cut its support by a constraint  $c(\mathbf{x})$**

 $q(\mathbf{x})$  $c(\mathbf{x})$  $q(\mathbf{x}) \cdot c(\mathbf{x})$ 

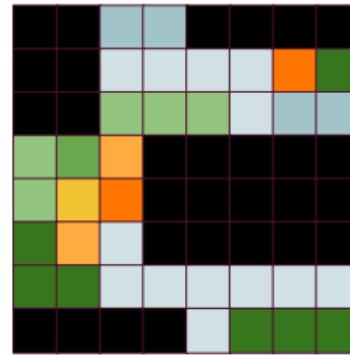
**by multiplying them**  $q(\mathbf{x})c(\mathbf{x})\dots$



$$q(\mathbf{x})$$

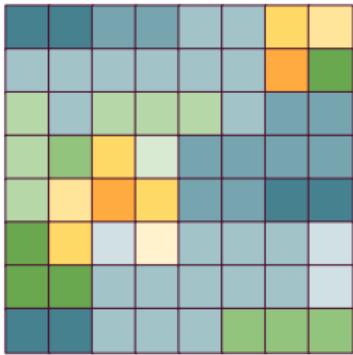


$$c(\mathbf{x})$$

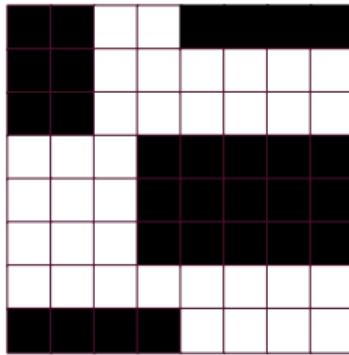


$$\frac{q(\mathbf{x}) \cdot c(\mathbf{x})}{\sum_{\mathbf{x}} q(\mathbf{x}) \cdot c(\mathbf{x})}$$

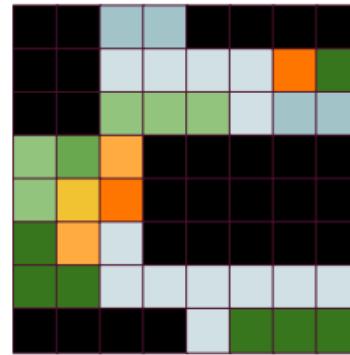
**and then renormalizing them!**



$q(\mathbf{x})$

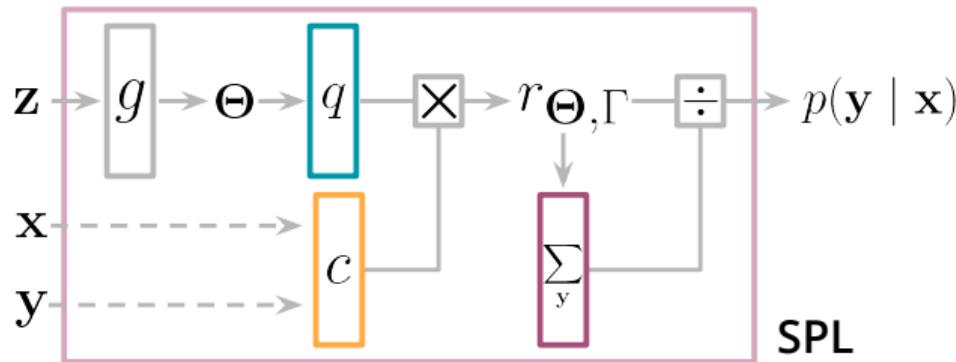


$c(\mathbf{x})$



$$\frac{q(\mathbf{x}) \cdot c(\mathbf{x})}{\sum_{\mathbf{x}} q(\mathbf{x}) \cdot c(\mathbf{x})}$$

**states with zero probability will never be predicted  
(nor sampled)**



$$p(\mathbf{y} \mid \mathbf{x}) = \mathbf{q}_{\Theta}(\mathbf{y} \mid g(\mathbf{z})) \cdot \mathbf{c}_K(\mathbf{x}, \mathbf{y}) / \mathcal{Z}(\mathbf{x})$$

$$\mathcal{Z}(\mathbf{x}) = \sum_{\mathbf{y}} q_{\Theta}(\mathbf{y} \mid \mathbf{x}) \cdot c_K(\mathbf{x}, \mathbf{y})$$



Ground Truth



ResNet-18



Semantic Loss



circuits

***predictions guarantee a logical constraint 100% of the time!***

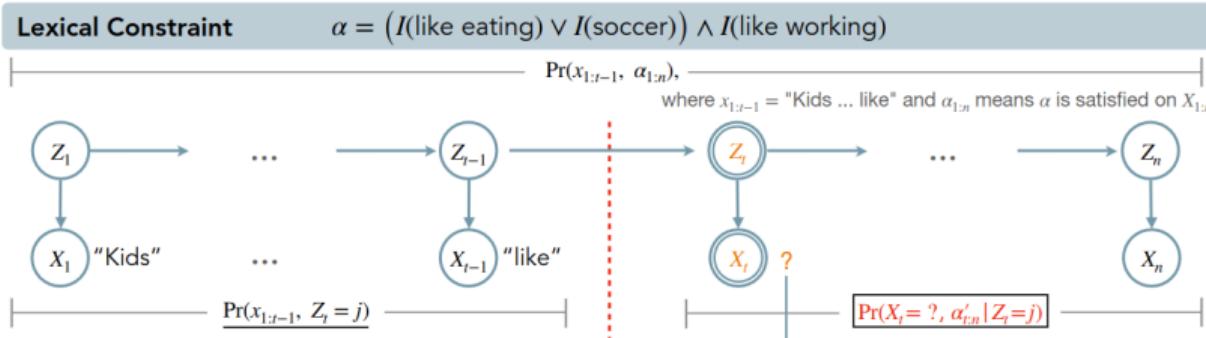
**SPL**  
**(and variants)**  
*everywhere*

---

## Tractable Control for Autoregressive Language Generation

---

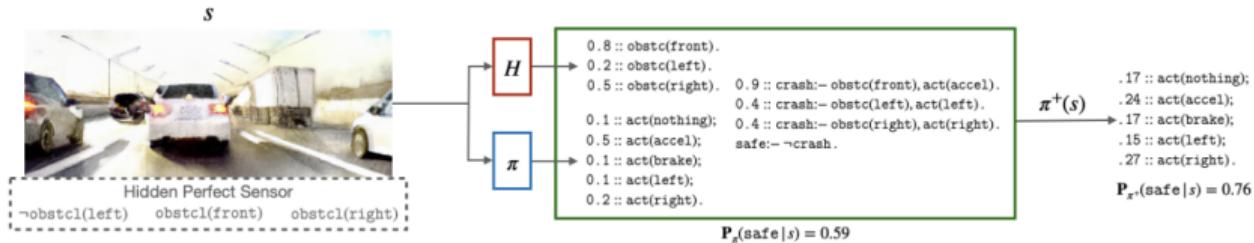
Honghua Zhang<sup>\* 1</sup> Meihua Dang<sup>\* 1</sup> Nanyun Peng<sup>1</sup> Guy Van den Broeck<sup>1</sup>



***constrained text generation with LLMs (ICML 2023)***

# Safe Reinforcement Learning via Probabilistic Logic Shields

Wen-Chi Yang<sup>1</sup>, Giuseppe Marra<sup>1</sup>, Gavin Rens and Luc De Raedt<sup>1,2</sup>



***reliable reinforcement learning (AAAI 23)***

---

# How to Turn Your Knowledge Graph Embeddings into Generative Models

---

**Lorenzo Loconte**  
University of Edinburgh, UK  
l.loconte@sms.ed.ac.uk

**Nicola Di Mauro**  
University of Bari, Italy  
nicola.dimauro@uniba.it

**Robert Peharz**  
TU Graz, Austria  
robert.peharz@tugraz.at

**Antonio Vergari**  
University of Edinburgh, UK  
avergari@ed.ac.uk

*enforce constraints in knowledge graph embeddings  
oral at NeurIPS 2023*

---

# Logically Consistent Language Models via Neuro-Symbolic Integration

---

The figure displays three panels illustrating logical reasoning:

- Forward Implication**:  $A \rightarrow \neg B$ . Given facts A: (albatross, isA, bird) and B: (albatross, isA, fish).
  - Question: Is an albatross a bird? Response: Yes. Logical: ✕ Factual: ✕
  - Question: Is an albatross a fish? Response: No. Logical: ✓ Factual: ✓
- Reverse Implication**:  $\neg B \rightarrow \neg A$ . Given facts B: (albatross, isNotA, organism) and C: (albatross, isNotA, living thing).
  - Question: Is it true that an albatross is not an organism? Response: No. Logical: ✕ Factual: ✕
  - Question: Is it true that an albatross is not a living thing? Response: Yes. Logical: ✕ Factual: ✕
- Negation**:  $A \bullet \neg A$ . Given facts A: (computer, isA, airplane) and C: (computer, isNotA, airplane).
  - Question: Is a computer a airplane? Response: No. Logical: ✕ Factual: ✕
  - Question: Is it true that a computer is not a airplane? Response: Yes. Logical: ✓ Factual: ✓

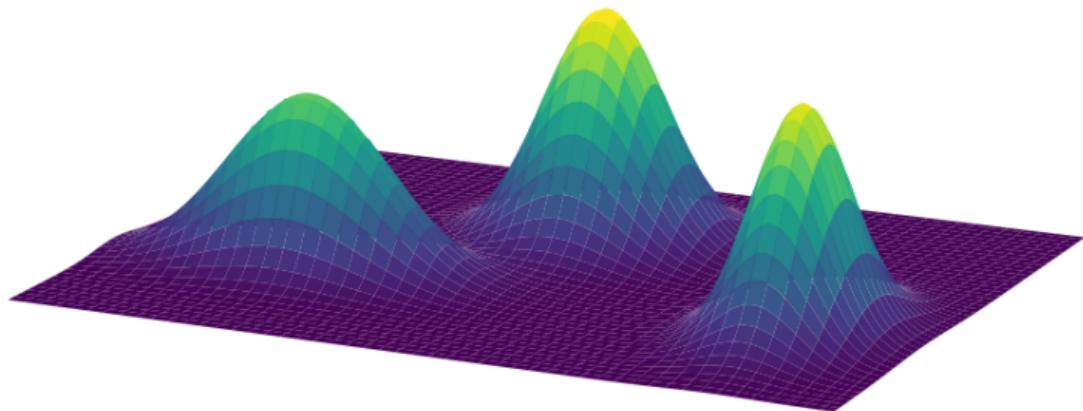
Legend:  
LLM2 = LoCo-LLMs 2  
LoCo-LLMs 2 = LLaMA 2

*improving logical (self-)consistency in LLMs at ICLR 2025*

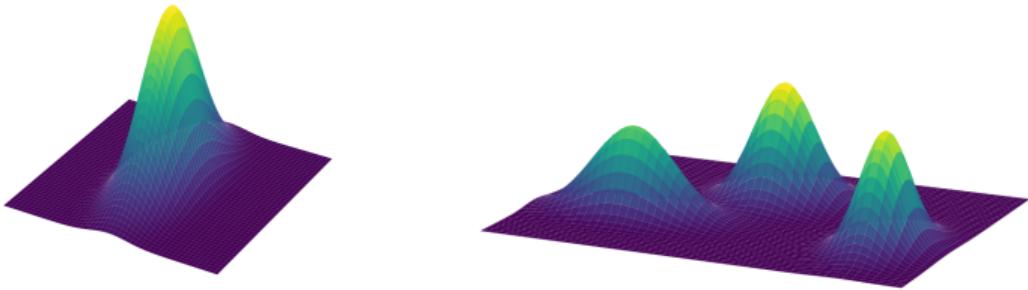


***learning & reasoning with circuits in pytorch***

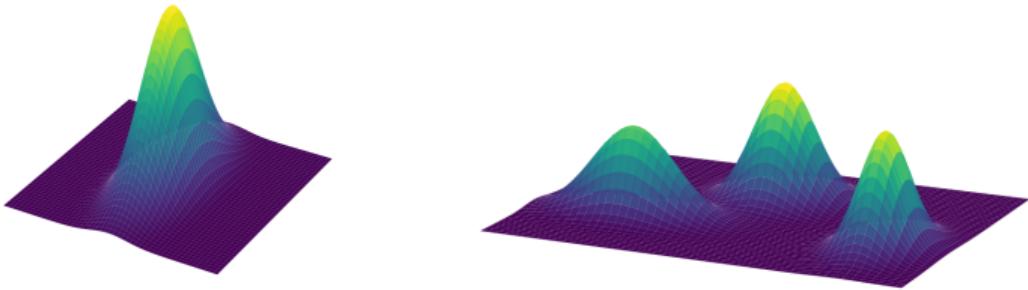
[github.com/april-tools/cirkit](https://github.com/april-tools/cirkit)



*oh mixtures, you're so fine you blow my mind!*



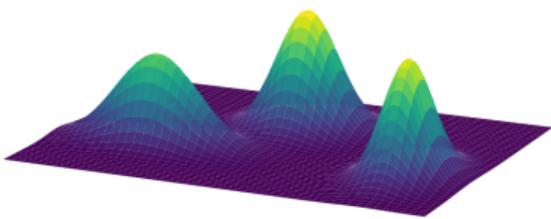
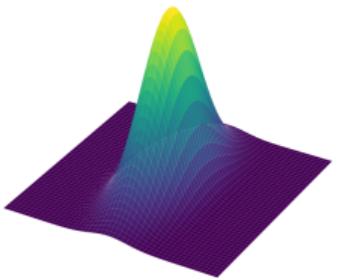
$$p(\mathbf{X}) \quad \xrightarrow{\text{orange arrow}} \quad \sum_{i=1}^K w_i p_i(\mathbf{X}) \quad w_i > 0$$



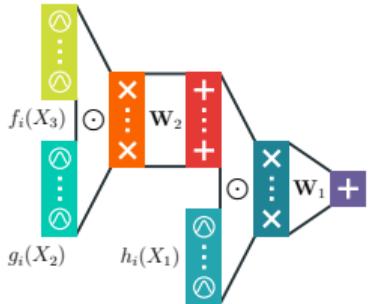
$$p(\mathbf{X}) \longrightarrow \sum_{i=1}^K w_i p_i(\mathbf{X}) \quad w_i > 0$$

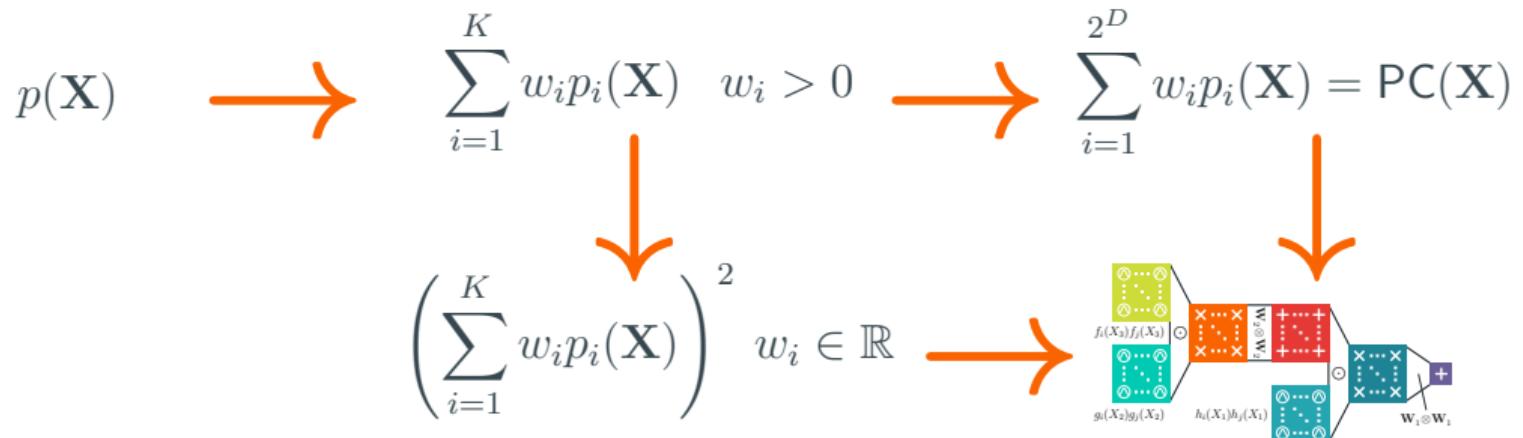
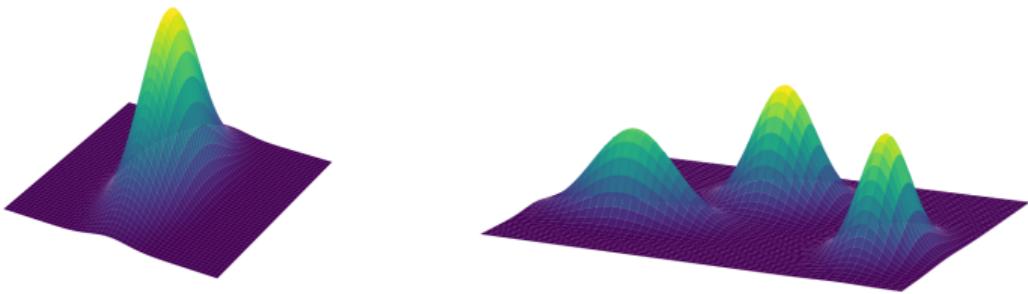
*"if someone publishes a paper on **model A**, there will be a paper about  
**mixtures of A** soon, with high probability"*

A. Vergari



$$p(\mathbf{X}) \rightarrow \sum_{i=1}^K w_i p_i(\mathbf{X}) \quad w_i > 0 \rightarrow \sum_{i=1}^{2^D} w_i p_i(\mathbf{X}) = \text{PC}(\mathbf{X})$$

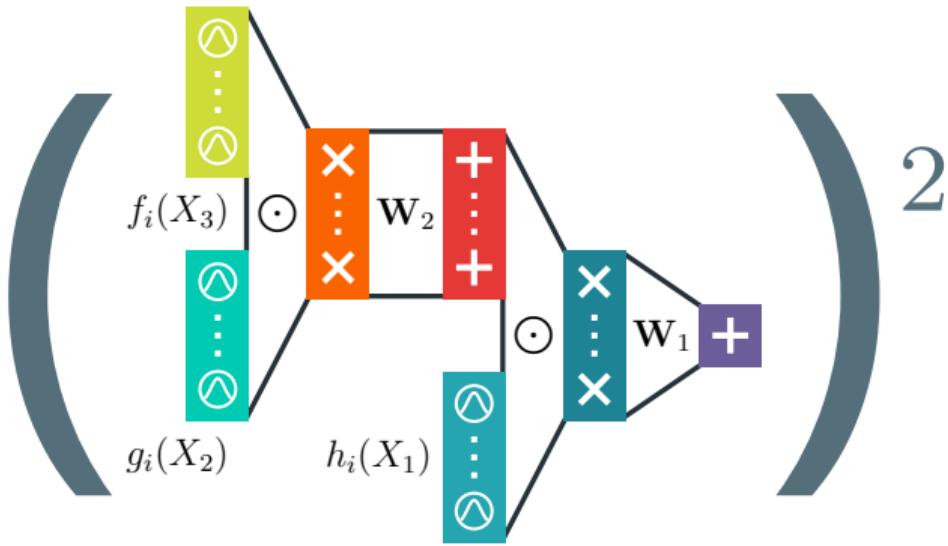






***learning & reasoning with circuits in pytorch***

[github.com/april-tools/cirkit](https://github.com/april-tools/cirkit)



**questions?**